

Homework-3

Problem 1 - *SSD, ONNX model, ORT Inferencing* (20 points)

In this problem, we will be inferencing SSD ONNX model using ONNX Runtime Server. You will follow the github repo and ONNX tutorials (links provided below). You will start with a pretrained Pytorch SSD model and retrain it for your target categories. Then you will convert this Pytorch model to ONNX and deploy it on ONNX runtime server for inferencing.

1. Download pretrained pytorch MobilenetV1 SSD and test it locally

Download pretrained pytorch MobilenetV1 SSD and test it locally using [Pascal VOC 2007 dataset](#). Select any two related categories from [Google Open Images dataset](#) and finetune the pretrained SSD model. Examples include, Aircraft and Aeroplane, Handgun and Shotgun. You can use the

`open_images_downloader.py` script provided at the github to download the data. For finetuning you can use the same parameters as in the tutorial below. Compute the accuracy of the test data for these categories before and after finetuning.

(3 points)

```
In [18]: import torch
import torchvision
from torchvision.models.detection import ssdlite320_mobilenet_v3_large
import torch.nn as nn
from torchvision import transforms
import onnx
import onnxruntime as ort
import numpy as np
from PIL import Image, ImageDraw
import matplotlib.pyplot as plt
import os
from torch.utils.data import Dataset, DataLoader

# Define paths
OPEN_IMAGES_ROOT = '/scratch/poh2005/data/open_images_temp'
VOC_ROOT = '/scratch/poh2005/data/VOC'

def test_pretrained_ssd():
    """Test pretrained SSDLite320 MobileNetV3 model on VOC test dataset"""

```

```

print("Loading pretrained SSDLite320 MobileNetV3 model...")
model = ssdlite320_mobilenet_v3_large(pretrained=True)
model.eval()

# Load VOC test dataset
print("\nLoading Pascal VOC 2007 test dataset...")
test_dataset = torchvision.datasets.VOCDetection(
    root=VOC_ROOT,
    year='2007',
    image_set='test',
    download=True
)

transform = transforms.Compose([
    transforms.Resize((320, 320)),
    transforms.ToTensor(),
    transforms.Normalize(mean=[0.485, 0.456, 0.406],
                        std=[0.229, 0.224, 0.225])
])

device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
model = model.to(device)

# Test results storage
correct = 0
total = 0
results = []

print("\nTesting pretrained model on VOC dataset...")
for idx in range(len(test_dataset)):
    img, annotation = test_dataset[idx]
    objects = annotation['annotation']['object']
    if not isinstance(objects, list):
        objects = [objects]

    has_target_class = False
    target_boxes = []
    target_labels = []

# Get ground truth boxes and labels
for obj in objects:
    if obj['name'] in ['aeroplane', 'boat']:
        has_target_class = True
        bbox = obj['bndbox']
        target_boxes.append([
            float(bbox['xmin']), float(bbox['ymin']),
            float(bbox['xmax']), float(bbox['ymax'])
        ])
        target_labels.append(1 if obj['name'] == 'aeroplane' else 2)

if has_target_class:
    # Transform image
    img_tensor = transform(img).unsqueeze(0).to(device)

# Get predictions
    with torch.no_grad():

```



```

    ])

fig, axes = plt.subplots(1, num_images, figsize=(20, 4))
if num_images == 1:
    axes = [axes]

count = 0
for idx in range(len(test_dataset)):
    if count >= num_images:
        break

    img, annotation = test_dataset[idx]
    objects = annotation['annotation']['object']
    if not isinstance(objects, list):
        objects = [objects]

    has_target = False
    for obj in objects:
        if obj['name'] in ['aeroplane', 'boat']:
            has_target = True
            break

    if has_target:
        # Get predictions
        img_tensor = transform(img).unsqueeze(0).to(device)

        with torch.no_grad():
            predictions = model(img_tensor)

        # Draw original image with predictions
        img_draw = img.copy()
        draw = ImageDraw.Draw(img_draw)

        # Draw ground truth boxes
        for obj in objects:
            if obj['name'] in ['aeroplane', 'boat']:
                bbox = obj['bbox']
                xmin, ymin = float(bbox['xmin']), float(bbox['ymin'])
                xmax, ymax = float(bbox['xmax']), float(bbox['ymax'])
                draw.rectangle([xmin, ymin, xmax, ymax], outline='green')
                draw.text((xmin, ymin-10), obj['name'], fill='green')

        # Draw predictions
        pred_boxes = predictions[0]['boxes'].cpu()
        pred_scores = predictions[0]['scores'].cpu()
        pred_labels = predictions[0]['labels'].cpu()

        # Filter predictions
        mask = (pred_labels == 1) | (pred_labels == 2)
        pred_boxes = pred_boxes[mask]
        pred_scores = pred_scores[mask]
        pred_labels = pred_labels[mask]

        for box, score, label in zip(pred_boxes, pred_scores, pred_labels):
            if score > 0.5: # Confidence threshold
                box = box.numpy()

```

```

        draw.rectangle(box.tolist(), outline='red', width=2)
        class_name = 'airplane' if label == 1 else 'boat'
        draw.text((box[0], box[1]-20),
                  f'{class_name} {score:.2f}', fill='red')

    # Display
    axes[count].imshow(img_draw)
    axes[count].axis('off')
    axes[count].set_title(f'Image {count+1}')
    count += 1

plt.tight_layout()
plt.show()

def main():
    # Test pretrained model
    print("Testing pretrained SSD model...")
    model, pretrained_accuracy, pretrained_results = test_pretrained_ssd()

    # Visualize some detections
    print("\nVisualizing detections...")
    test_dataset = torchvision.datasets.VOCDetection(
        root=VOC_ROOT,
        year='2007',
        image_set='test',
        download=True
    )
    visualize_detections(model, test_dataset)

    return model, pretrained_results

if __name__ == '__main__':
    # Create output directories
    os.makedirs(VOC_ROOT, exist_ok=True)
    os.makedirs(OPEN_IMAGES_ROOT, exist_ok=True)

    # Run main process
    model, results = main()

```

```
Testing pretrained SSD model...
Loading pretrained SSDLite320 MobileNetV3 model...

Loading Pascal VOC 2007 test dataset...
Using downloaded and verified file: /scratch/poh2005/data/VOC/VOCtest_06-Nov-2007.tar
Extracting /scratch/poh2005/data/VOC/VOCtest_06-Nov-2007.tar to /scratch/poh2005/data/VOC

Testing pretrained model on VOC dataset...

Test Result 1:
Image: 000067.jpg
True labels: ['airplane']
Predicted: airplane (Confidence: 9.83%)

Test Result 2:
Image: 000069.jpg
True labels: ['boat']
Predicted: airplane (Confidence: 6.10%)

Test Result 3:
Image: 000080.jpg
True labels: ['boat', 'boat', 'boat', 'boat', 'boat', 'boat', 'boat', 'boat', 'boat', 'boat']
Predicted: airplane (Confidence: 2.42%)

Test Result 4:
Image: 000105.jpg
True labels: ['boat']
Predicted: airplane (Confidence: 12.68%)

Test Result 5:
Image: 000128.jpg
True labels: ['airplane']
Predicted: airplane (Confidence: 16.29%)

Test Result 6:
Image: 000179.jpg
True labels: ['boat', 'boat', 'boat', 'boat', 'boat']
Predicted: airplane (Confidence: 3.68%)

Test Result 7:
Image: 000216.jpg
True labels: ['airplane']
Predicted: airplane (Confidence: 3.68%)

Test Result 8:
Image: 000240.jpg
True labels: ['boat', 'boat', 'boat']
Predicted: airplane (Confidence: 5.97%)

Test Result 9:
Image: 000243.jpg
True labels: ['airplane']
Predicted: airplane (Confidence: 1.46%)
```

Test Result 10:
Image: 000260.jpg
True labels: ['airplane', 'airplane']
Predicted: airplane (Confidence: 0.54%)

Test Result 11:
Image: 000295.jpg
True labels: ['boat', 'boat']
Predicted: airplane (Confidence: 8.33%)

Test Result 12:
Image: 000316.jpg
True labels: ['airplane']
Predicted: airplane (Confidence: 3.53%)

Test Result 13:
Image: 000350.jpg
True labels: ['boat', 'boat', 'boat']
Predicted: airplane (Confidence: 6.25%)

Test Result 14:
Image: 000371.jpg
True labels: ['boat']
Predicted: airplane (Confidence: 5.26%)

Test Result 15:
Image: 000375.jpg
True labels: ['boat']
Predicted: airplane (Confidence: 10.19%)

Test Result 16:
Image: 000418.jpg
True labels: ['airplane']
Predicted: airplane (Confidence: 11.87%)

Test Result 17:
Image: 000444.jpg
True labels: ['boat']
Predicted: airplane (Confidence: 14.38%)

Test Result 18:
Image: 000449.jpg
True labels: ['boat']
Predicted: airplane (Confidence: 17.14%)

Test Result 19:
Image: 000473.jpg
True labels: ['airplane']
Predicted: airplane (Confidence: 3.02%)

Test Result 20:
Image: 000481.jpg
True labels: ['boat', 'boat', 'boat']
Predicted: airplane (Confidence: 16.13%)

Test Result 21:
Image: 000521.jpg
True labels: ['airplane']
Predicted: airplane (Confidence: 1.23%)

Test Result 22:
Image: 000529.jpg
True labels: ['boat']
Predicted: airplane (Confidence: 0.87%)

Test Result 23:
Image: 000538.jpg
True labels: ['boat']
Predicted: airplane (Confidence: 1.04%)

Test Result 24:
Image: 000560.jpg
True labels: ['airplane']
Predicted: airplane (Confidence: 2.28%)

Test Result 25:
Image: 000576.jpg
True labels: ['boat']
Predicted: airplane (Confidence: 3.47%)

Test Result 26:
Image: 000600.jpg
True labels: ['airplane']
Predicted: airplane (Confidence: 33.38%)

Test Result 27:
Image: 000618.jpg
True labels: ['airplane', 'airplane']
Predicted: airplane (Confidence: 14.34%)

Test Result 28:
Image: 000631.jpg
True labels: ['boat', 'boat']
Predicted: airplane (Confidence: 3.22%)

Test Result 29:
Image: 000665.jpg
True labels: ['airplane']
Predicted: airplane (Confidence: 0.70%)

Test Result 30:
Image: 000668.jpg
True labels: ['airplane']
Predicted: airplane (Confidence: 3.75%)

Test Result 31:
Image: 000696.jpg
True labels: ['airplane']
Predicted: airplane (Confidence: 5.70%)

Test Result 32:

Image: 000706.jpg
True labels: ['airplane', 'airplane', 'airplane', 'airplane', 'airplane']
Predicted: airplane (Confidence: 8.61%)

Test Result 33:
Image: 000792.jpg
True labels: ['boat', 'boat']
Predicted: airplane (Confidence: 2.77%)

Test Result 34:
Image: 000811.jpg
True labels: ['boat']
Predicted: airplane (Confidence: 2.20%)

Test Result 35:
Image: 000817.jpg
True labels: ['airplane']
Predicted: airplane (Confidence: 7.92%)

Test Result 36:
Image: 000837.jpg
True labels: ['boat']
Predicted: airplane (Confidence: 3.02%)

Test Result 37:
Image: 000846.jpg
True labels: ['airplane', 'airplane', 'airplane', 'airplane', 'airplane', 'airplane', 'airplane', 'airplane']
Predicted: airplane (Confidence: 2.00%)

Test Result 38:
Image: 000907.jpg
True labels: ['boat']
Predicted: airplane (Confidence: 8.68%)

Test Result 39:
Image: 000914.jpg
True labels: ['boat']
Predicted: airplane (Confidence: 10.68%)

Test Result 40:
Image: 000968.jpg
True labels: ['airplane']
Predicted: airplane (Confidence: 4.72%)

Test Result 41:
Image: 000976.jpg
True labels: ['airplane']
Predicted: airplane (Confidence: 4.17%)

Test Result 42:
Image: 000995.jpg
True labels: ['boat', 'boat']
Predicted: airplane (Confidence: 0.62%)

Test Result 43:

Image: 001049.jpg
True labels: ['boat']
Predicted: airplane (Confidence: 0.89%)

Test Result 44:
Image: 001059.jpg
True labels: ['boat', 'boat']
Predicted: airplane (Confidence: 6.74%)

Test Result 45:
Image: 001076.jpg
True labels: ['boat', 'boat']
Predicted: airplane (Confidence: 3.88%)

Test Result 46:
Image: 001080.jpg
True labels: ['airplane']
Predicted: airplane (Confidence: 0.34%)

Test Result 47:
Image: 001099.jpg
True labels: ['airplane']
Predicted: airplane (Confidence: 2.62%)

Test Result 48:
Image: 001141.jpg
True labels: ['boat', 'boat']
Predicted: airplane (Confidence: 10.86%)

Test Result 49:
Image: 001153.jpg
True labels: ['airplane']
Predicted: airplane (Confidence: 3.90%)

Test Result 50:
Image: 001155.jpg
True labels: ['boat']
Predicted: airplane (Confidence: 9.24%)

Test Result 51:
Image: 001188.jpg
True labels: ['boat', 'boat', 'boat']
Predicted: airplane (Confidence: 11.00%)

Test Result 52:
Image: 001223.jpg
True labels: ['airplane']
Predicted: airplane (Confidence: 4.40%)

Test Result 53:
Image: 001227.jpg
True labels: ['boat']
Predicted: airplane (Confidence: 21.00%)

Test Result 54:

Image: 001305.jpg
True labels: ['airplane']
Predicted: airplane (Confidence: 8.01%)

Test Result 55:
Image: 001342.jpg
True labels: ['boat', 'boat', 'boat', 'boat', 'boat', 'boat', 'boat']
Predicted: airplane (Confidence: 1.04%)

Test Result 56:
Image: 001355.jpg
True labels: ['boat']
Predicted: airplane (Confidence: 18.19%)

Test Result 57:
Image: 001373.jpg
True labels: ['airplane']
Predicted: airplane (Confidence: 4.71%)

Test Result 58:
Image: 001377.jpg
True labels: ['airplane']
Predicted: airplane (Confidence: 3.32%)

Test Result 59:
Image: 001394.jpg
True labels: ['boat']
Predicted: airplane (Confidence: 1.57%)

Test Result 60:
Image: 001410.jpg
True labels: ['boat', 'boat', 'boat']
Predicted: airplane (Confidence: 7.86%)

Test Result 61:
Image: 001433.jpg
True labels: ['airplane']
Predicted: airplane (Confidence: 2.15%)

Test Result 62:
Image: 001474.jpg
True labels: ['boat']
Predicted: airplane (Confidence: 4.66%)

Test Result 63:
Image: 001487.jpg
True labels: ['boat']
Predicted: airplane (Confidence: 18.97%)

Test Result 64:
Image: 001505.jpg
True labels: ['airplane']
Predicted: airplane (Confidence: 0.70%)

Test Result 65:

Image: 001527.jpg
True labels: ['boat', 'boat', 'boat']
Predicted: airplane (Confidence: 0.41%)

Test Result 66:
Image: 001547.jpg
True labels: ['airplane']
Predicted: airplane (Confidence: 3.11%)

Test Result 67:
Image: 001568.jpg
True labels: ['airplane']
Predicted: airplane (Confidence: 29.65%)

Test Result 68:
Image: 001591.jpg
True labels: ['boat']
Predicted: airplane (Confidence: 1.36%)

Test Result 69:
Image: 001621.jpg
True labels: ['airplane']
Predicted: airplane (Confidence: 2.05%)

Test Result 70:
Image: 001634.jpg
True labels: ['boat']
Predicted: airplane (Confidence: 4.20%)

Test Result 71:
Image: 001646.jpg
True labels: ['boat', 'boat']
Predicted: airplane (Confidence: 4.69%)

Test Result 72:
Image: 001660.jpg
True labels: ['airplane']
Predicted: airplane (Confidence: 1.01%)

Test Result 73:
Image: 001698.jpg
True labels: ['boat', 'boat']
Predicted: airplane (Confidence: 5.97%)

Test Result 74:
Image: 001705.jpg
True labels: ['boat']
Predicted: airplane (Confidence: 15.62%)

Test Result 75:
Image: 001770.jpg
True labels: ['boat']
Predicted: airplane (Confidence: 7.08%)

Test Result 76:
Image: 001783.jpg

True labels: ['airplane', 'airplane']
Predicted: airplane (Confidence: 5.18%)

Test Result 77:
Image: 001815.jpg
True labels: ['boat']
Predicted: airplane (Confidence: 8.03%)

Test Result 78:
Image: 001822.jpg
True labels: ['boat']
Predicted: airplane (Confidence: 2.18%)

Test Result 79:
Image: 001848.jpg
True labels: ['airplane']
Predicted: airplane (Confidence: 8.26%)

Test Result 80:
Image: 001850.jpg
True labels: ['airplane']
Predicted: airplane (Confidence: 2.86%)

Test Result 81:
Image: 001884.jpg
True labels: ['airplane']
Predicted: airplane (Confidence: 8.88%)

Test Result 82:
Image: 001885.jpg
True labels: ['airplane', 'airplane']
Predicted: airplane (Confidence: 5.89%)

Test Result 83:
Image: 001895.jpg
True labels: ['boat', 'boat', 'boat', 'boat']
Predicted: airplane (Confidence: 3.98%)

Test Result 84:
Image: 001912.jpg
True labels: ['airplane']
Predicted: airplane (Confidence: 16.22%)

Test Result 85:
Image: 001974.jpg
True labels: ['boat', 'boat']
Predicted: airplane (Confidence: 73.34%)

Test Result 86:
Image: 001994.jpg
True labels: ['airplane']
Predicted: airplane (Confidence: 4.24%)

Test Result 87:
Image: 001996.jpg
True labels: ['airplane', 'airplane']

Predicted: airplane (Confidence: 3.53%)

Test Result 88:

Image: 002007.jpg

True labels: ['airplane']

Predicted: airplane (Confidence: 12.57%)

Test Result 89:

Image: 002014.jpg

True labels: ['airplane']

Predicted: airplane (Confidence: 0.89%)

Test Result 90:

Image: 002029.jpg

True labels: ['boat']

Predicted: airplane (Confidence: 3.93%)

Test Result 91:

Image: 002052.jpg

True labels: ['airplane']

Predicted: airplane (Confidence: 5.80%)

Test Result 92:

Image: 002085.jpg

True labels: ['boat']

Predicted: airplane (Confidence: 9.67%)

Test Result 93:

Image: 002089.jpg

True labels: ['airplane']

Predicted: airplane (Confidence: 8.56%)

Test Result 94:

Image: 002107.jpg

True labels: ['airplane']

Predicted: airplane (Confidence: 0.45%)

Test Result 95:

Image: 002110.jpg

True labels: ['airplane']

Predicted: airplane (Confidence: 7.37%)

Test Result 96:

Image: 002157.jpg

True labels: ['airplane']

Predicted: airplane (Confidence: 3.91%)

Test Result 97:

Image: 002175.jpg

True labels: ['boat']

Predicted: airplane (Confidence: 1.03%)

Test Result 98:

Image: 002198.jpg

True labels: ['airplane']

Predicted: airplane (Confidence: 4.36%)

Test Result 99:
Image: 002217.jpg
True labels: ['airplane', 'airplane', 'airplane', 'airplane', 'airplane', 'airplane']
Predicted: airplane (Confidence: 3.08%)

Test Result 100:
Image: 002225.jpg
True labels: ['boat']
Predicted: airplane (Confidence: 7.49%)

Test Result 101:
Image: 002246.jpg
True labels: ['airplane']
Predicted: airplane (Confidence: 1.09%)

Test Result 102:
Image: 002274.jpg
True labels: ['boat', 'boat']
Predicted: airplane (Confidence: 2.55%)

Test Result 103:
Image: 002390.jpg
True labels: ['boat', 'boat']
Predicted: airplane (Confidence: 9.16%)

Test Result 104:
Image: 002395.jpg
True labels: ['boat']
Predicted: airplane (Confidence: 16.84%)

Test Result 105:
Image: 002449.jpg
True labels: ['boat', 'boat']
Predicted: airplane (Confidence: 4.86%)

Test Result 106:
Image: 002474.jpg
True labels: ['boat', 'boat', 'boat']
Predicted: airplane (Confidence: 3.35%)

Test Result 107:
Image: 002580.jpg
True labels: ['boat', 'boat', 'boat', 'boat', 'boat']
Predicted: airplane (Confidence: 1.87%)

Test Result 108:
Image: 002583.jpg
True labels: ['airplane']
Predicted: airplane (Confidence: 1.30%)

Test Result 109:
Image: 002619.jpg
True labels: ['airplane']
Predicted: airplane (Confidence: 3.34%)

Test Result 110:
Image: 002629.jpg
True labels: ['airplane']
Predicted: airplane (Confidence: 1.78%)

Test Result 111:
Image: 002661.jpg
True labels: ['airplane']
Predicted: airplane (Confidence: 0.90%)

Test Result 112:
Image: 002665.jpg
True labels: ['airplane']
Predicted: airplane (Confidence: 3.48%)

Test Result 113:
Image: 002703.jpg
True labels: ['airplane']
Predicted: airplane (Confidence: 4.18%)

Test Result 114:
Image: 002707.jpg
True labels: ['airplane']
Predicted: airplane (Confidence: 11.75%)

Test Result 115:
Image: 002719.jpg
True labels: ['boat']
Predicted: airplane (Confidence: 1.95%)

Test Result 116:
Image: 002754.jpg
True labels: ['airplane']
Predicted: airplane (Confidence: 5.97%)

Test Result 117:
Image: 002764.jpg
True labels: ['boat']
Predicted: airplane (Confidence: 0.46%)

Test Result 118:
Image: 002769.jpg
True labels: ['airplane']
Predicted: airplane (Confidence: 0.58%)

Test Result 119:
Image: 002811.jpg
True labels: ['boat', 'boat']
Predicted: airplane (Confidence: 10.68%)

Test Result 120:
Image: 002821.jpg
True labels: ['boat']
Predicted: airplane (Confidence: 1.59%)

Test Result 121:
Image: 002822.jpg
True labels: ['airplane']
Predicted: airplane (Confidence: 7.00%)

Test Result 122:
Image: 002843.jpg
True labels: ['airplane', 'airplane']
Predicted: airplane (Confidence: 1.42%)

Test Result 123:
Image: 002851.jpg
True labels: ['airplane']
Predicted: airplane (Confidence: 5.85%)

Test Result 124:
Image: 002894.jpg
True labels: ['airplane', 'airplane']
Predicted: airplane (Confidence: 12.90%)

Test Result 125:
Image: 002908.jpg
True labels: ['airplane']
Predicted: airplane (Confidence: 3.16%)

Test Result 126:
Image: 002948.jpg
True labels: ['boat']
Predicted: airplane (Confidence: 0.70%)

Test Result 127:
Image: 002949.jpg
True labels: ['airplane']
Predicted: airplane (Confidence: 1.28%)

Test Result 128:
Image: 002971.jpg
True labels: ['airplane']
Predicted: airplane (Confidence: 8.91%)

Test Result 129:
Image: 002983.jpg
True labels: ['boat']
Predicted: airplane (Confidence: 3.36%)

Test Result 130:
Image: 003012.jpg
True labels: ['boat']
Predicted: airplane (Confidence: 37.82%)

Test Result 131:
Image: 003037.jpg
True labels: ['airplane', 'airplane']
Predicted: airplane (Confidence: 31.90%)

Test Result 132:

Image: 003073.jpg
True labels: ['boat']
Predicted: airplane (Confidence: 3.99%)

Test Result 133:
Image: 003131.jpg
True labels: ['boat']
Predicted: airplane (Confidence: 22.55%)

Test Result 134:
Image: 003144.jpg
True labels: ['airplane']
Predicted: airplane (Confidence: 1.59%)

Test Result 135:
Image: 003209.jpg
True labels: ['airplane', 'airplane']
Predicted: airplane (Confidence: 0.28%)

Test Result 136:
Image: 003230.jpg
True labels: ['airplane']
Predicted: airplane (Confidence: 1.72%)

Test Result 137:
Image: 003268.jpg
True labels: ['airplane', 'airplane']
Predicted: airplane (Confidence: 4.05%)

Test Result 138:
Image: 003347.jpg
True labels: ['airplane']
Predicted: airplane (Confidence: 16.30%)

Test Result 139:
Image: 003372.jpg
True labels: ['airplane']
Predicted: airplane (Confidence: 2.30%)

Test Result 140:
Image: 003381.jpg
True labels: ['airplane']
Predicted: airplane (Confidence: 6.05%)

Test Result 141:
Image: 003409.jpg
True labels: ['boat', 'boat', 'boat']
Predicted: airplane (Confidence: 0.43%)

Test Result 142:
Image: 003445.jpg
True labels: ['airplane', 'airplane']
Predicted: airplane (Confidence: 4.63%)

Test Result 143:
Image: 003478.jpg

True labels: ['airplane']
Predicted: airplane (Confidence: 31.71%)

Test Result 144:
Image: 003494.jpg
True labels: ['airplane']
Predicted: airplane (Confidence: 5.08%)

Test Result 145:
Image: 003498.jpg
True labels: ['boat']
Predicted: airplane (Confidence: 5.90%)

Test Result 146:
Image: 003543.jpg
True labels: ['boat', 'boat', 'boat', 'boat']
Predicted: airplane (Confidence: 77.29%)

Test Result 147:
Image: 003571.jpg
True labels: ['boat']
Predicted: airplane (Confidence: 6.48%)

Test Result 148:
Image: 003574.jpg
True labels: ['airplane']
Predicted: airplane (Confidence: 11.89%)

Test Result 149:
Image: 003652.jpg
True labels: ['boat', 'boat', 'boat', 'boat', 'boat', 'boat']
Predicted: airplane (Confidence: 15.79%)

Test Result 150:
Image: 003799.jpg
True labels: ['boat', 'boat']
Predicted: airplane (Confidence: 2.89%)

Test Result 151:
Image: 003823.jpg
True labels: ['airplane']
Predicted: airplane (Confidence: 11.85%)

Test Result 152:
Image: 003836.jpg
True labels: ['boat']
Predicted: airplane (Confidence: 2.37%)

Test Result 153:
Image: 003892.jpg
True labels: ['boat']
Predicted: airplane (Confidence: 3.03%)

Test Result 154:
Image: 003900.jpg
True labels: ['boat', 'boat']

Predicted: airplane (Confidence: 3.75%)

Test Result 155:

Image: 003910.jpg

True labels: ['boat']

Predicted: airplane (Confidence: 49.29%)

Test Result 156:

Image: 003929.jpg

True labels: ['boat', 'boat']

Predicted: airplane (Confidence: 9.65%)

Test Result 157:

Image: 003931.jpg

True labels: ['airplane']

Predicted: airplane (Confidence: 2.44%)

Test Result 158:

Image: 003952.jpg

True labels: ['boat']

Predicted: airplane (Confidence: 0.89%)

Test Result 159:

Image: 004078.jpg

True labels: ['airplane']

Predicted: airplane (Confidence: 3.90%)

Test Result 160:

Image: 004123.jpg

True labels: ['boat', 'boat']

Predicted: airplane (Confidence: 2.66%)

Test Result 161:

Image: 004153.jpg

True labels: ['airplane', 'airplane']

Predicted: airplane (Confidence: 3.38%)

Test Result 162:

Image: 004160.jpg

True labels: ['boat', 'boat']

Predicted: airplane (Confidence: 5.82%)

Test Result 163:

Image: 004184.jpg

True labels: ['boat']

Predicted: airplane (Confidence: 2.85%)

Test Result 164:

Image: 004199.jpg

True labels: ['airplane']

Predicted: airplane (Confidence: 2.12%)

Test Result 165:

Image: 004225.jpg

True labels: ['boat', 'boat']

Predicted: airplane (Confidence: 0.77%)

Test Result 166:
Image: 004278.jpg
True labels: ['airplane']
Predicted: airplane (Confidence: 1.48%)

Test Result 167:
Image: 004299.jpg
True labels: ['boat', 'boat', 'boat']
Predicted: airplane (Confidence: 7.03%)

Test Result 168:
Image: 004314.jpg
True labels: ['airplane']
Predicted: airplane (Confidence: 4.77%)

Test Result 169:
Image: 004374.jpg
True labels: ['boat']
Predicted: airplane (Confidence: 4.15%)

Test Result 170:
Image: 004377.jpg
True labels: ['airplane', 'airplane', 'airplane', 'airplane', 'airplane', 'airplane']
Predicted: airplane (Confidence: 2.95%)

Test Result 171:
Image: 004382.jpg
True labels: ['boat']
Predicted: airplane (Confidence: 2.95%)

Test Result 172:
Image: 004398.jpg
True labels: ['airplane']
Predicted: airplane (Confidence: 0.81%)

Test Result 173:
Image: 004443.jpg
True labels: ['boat', 'boat', 'boat', 'boat']
Predicted: airplane (Confidence: 9.96%)

Test Result 174:
Image: 004497.jpg
True labels: ['airplane']
Predicted: airplane (Confidence: 4.53%)

Test Result 175:
Image: 004533.jpg
True labels: ['airplane']
Predicted: airplane (Confidence: 1.35%)

Test Result 176:
Image: 004556.jpg
True labels: ['airplane']
Predicted: airplane (Confidence: 4.49%)

Test Result 177:
Image: 004572.jpg
True labels: ['airplane']
Predicted: airplane (Confidence: 3.19%)

Test Result 178:
Image: 004613.jpg
True labels: ['boat']
Predicted: airplane (Confidence: 5.25%)

Test Result 179:
Image: 004633.jpg
True labels: ['airplane']
Predicted: airplane (Confidence: 3.51%)

Test Result 180:
Image: 004665.jpg
True labels: ['boat', 'boat']
Predicted: airplane (Confidence: 8.54%)

Test Result 181:
Image: 004680.jpg
True labels: ['airplane']
Predicted: airplane (Confidence: 0.97%)

Test Result 182:
Image: 004752.jpg
True labels: ['airplane']
Predicted: airplane (Confidence: 5.38%)

Test Result 183:
Image: 004784.jpg
True labels: ['airplane']
Predicted: airplane (Confidence: 1.31%)

Test Result 184:
Image: 004824.jpg
True labels: ['airplane', 'airplane']
Predicted: airplane (Confidence: 48.82%)

Test Result 185:
Image: 004843.jpg
True labels: ['boat', 'boat', 'boat']
Predicted: airplane (Confidence: 7.50%)

Test Result 186:
Image: 004899.jpg
True labels: ['boat', 'boat', 'boat']
Predicted: airplane (Confidence: 13.60%)

Test Result 187:
Image: 004906.jpg
True labels: ['airplane']
Predicted: airplane (Confidence: 1.44%)

Test Result 188:
Image: 004914.jpg
True labels: ['boat']
Predicted: airplane (Confidence: 4.49%)

Test Result 189:
Image: 004937.jpg
True labels: ['boat', 'boat']
Predicted: airplane (Confidence: 17.88%)

Test Result 190:
Image: 004988.jpg
True labels: ['airplane']
Predicted: airplane (Confidence: 5.05%)

Test Result 191:
Image: 005008.jpg
True labels: ['airplane']
Predicted: airplane (Confidence: 2.48%)

Test Result 192:
Image: 005022.jpg
True labels: ['boat']
Predicted: airplane (Confidence: 0.66%)

Test Result 193:
Image: 005034.jpg
True labels: ['airplane']
Predicted: airplane (Confidence: 0.48%)

Test Result 194:
Image: 005043.jpg
True labels: ['airplane']
Predicted: airplane (Confidence: 3.44%)

Test Result 195:
Image: 005074.jpg
True labels: ['boat']
Predicted: airplane (Confidence: 3.72%)

Test Result 196:
Image: 005076.jpg
True labels: ['airplane']
Predicted: airplane (Confidence: 66.84%)

Test Result 197:
Image: 005098.jpg
True labels: ['boat']
Predicted: airplane (Confidence: 3.01%)

Test Result 198:
Image: 005106.jpg
True labels: ['airplane', 'airplane']
Predicted: airplane (Confidence: 2.47%)

Test Result 199:

Image: 005112.jpg
True labels: ['boat', 'boat', 'boat', 'boat', 'boat']
Predicted: airplane (Confidence: 6.44%)

Test Result 200:
Image: 005213.jpg
True labels: ['boat']
Predicted: airplane (Confidence: 2.91%)

Test Result 201:
Image: 005234.jpg
True labels: ['airplane', 'airplane']
Predicted: airplane (Confidence: 4.66%)

Test Result 202:
Image: 005238.jpg
True labels: ['airplane']
Predicted: airplane (Confidence: 2.36%)

Test Result 203:
Image: 005243.jpg
True labels: ['boat']
Predicted: airplane (Confidence: 3.53%)

Test Result 204:
Image: 005249.jpg
True labels: ['airplane']
Predicted: airplane (Confidence: 59.14%)

Test Result 205:
Image: 005272.jpg
True labels: ['boat', 'boat']
Predicted: airplane (Confidence: 10.70%)

Test Result 206:
Image: 005329.jpg
True labels: ['boat', 'boat', 'boat', 'boat', 'boat', 'boat', 'boat', 'boat', 'boat']
Predicted: airplane (Confidence: 0.90%)

Test Result 207:
Image: 005376.jpg
True labels: ['airplane']
Predicted: airplane (Confidence: 0.52%)

Test Result 208:
Image: 005382.jpg
True labels: ['boat']
Predicted: airplane (Confidence: 2.21%)

Test Result 209:
Image: 005392.jpg
True labels: ['airplane']
Predicted: airplane (Confidence: 78.13%)

Test Result 210:

Image: 005402.jpg
True labels: ['airplane']
Predicted: airplane (Confidence: 2.09%)

Test Result 211:
Image: 005427.jpg
True labels: ['boat', 'boat']
Predicted: airplane (Confidence: 18.92%)

Test Result 212:
Image: 005458.jpg
True labels: ['airplane', 'airplane']
Predicted: airplane (Confidence: 8.09%)

Test Result 213:
Image: 005537.jpg
True labels: ['airplane']
Predicted: airplane (Confidence: 3.57%)

Test Result 214:
Image: 005575.jpg
True labels: ['airplane']
Predicted: airplane (Confidence: 7.18%)

Test Result 215:
Image: 005604.jpg
True labels: ['boat', 'boat', 'boat']
Predicted: airplane (Confidence: 2.97%)

Test Result 216:
Image: 005616.jpg
True labels: ['boat', 'boat']
Predicted: airplane (Confidence: 4.21%)

Test Result 217:
Image: 005633.jpg
True labels: ['airplane', 'airplane', 'airplane', 'airplane', 'airplane']
Predicted: airplane (Confidence: 6.27%)

Test Result 218:
Image: 005670.jpg
True labels: ['boat', 'boat', 'boat']
Predicted: airplane (Confidence: 13.06%)

Test Result 219:
Image: 005720.jpg
True labels: ['boat']
Predicted: airplane (Confidence: 5.80%)

Test Result 220:
Image: 005771.jpg
True labels: ['boat', 'boat', 'boat']
Predicted: airplane (Confidence: 0.51%)

Test Result 221:
Image: 005849.jpg
True labels: ['boat']
Predicted: airplane (Confidence: 21.76%)

Test Result 222:
Image: 005927.jpg
True labels: ['airplane', 'airplane']
Predicted: airplane (Confidence: 7.82%)

Test Result 223:
Image: 005950.jpg
True labels: ['boat', 'boat']
Predicted: airplane (Confidence: 11.30%)

Test Result 224:
Image: 005965.jpg
True labels: ['airplane']
Predicted: airplane (Confidence: 6.45%)

Test Result 225:
Image: 006014.jpg
True labels: ['boat']
Predicted: airplane (Confidence: 7.23%)

Test Result 226:
Image: 006024.jpg
True labels: ['airplane']
Predicted: airplane (Confidence: 6.26%)

Test Result 227:
Image: 006048.jpg
True labels: ['boat']
Predicted: airplane (Confidence: 10.75%)

Test Result 228:
Image: 006051.jpg
True labels: ['airplane', 'airplane', 'airplane', 'airplane', 'airplane', 'airplane']
Predicted: airplane (Confidence: 1.09%)

Test Result 229:
Image: 006110.jpg
True labels: ['boat', 'boat']
Predicted: airplane (Confidence: 3.05%)

Test Result 230:
Image: 006160.jpg
True labels: ['boat']
Predicted: airplane (Confidence: 3.89%)

Test Result 231:
Image: 006164.jpg
True labels: ['boat']
Predicted: airplane (Confidence: 29.18%)

Test Result 232:
Image: 006205.jpg
True labels: ['boat', 'boat', 'boat', 'boat']
Predicted: airplane (Confidence: 21.34%)

Test Result 233:
Image: 006232.jpg
True labels: ['airplane']
Predicted: airplane (Confidence: 48.81%)

Test Result 234:
Image: 006294.jpg
True labels: ['airplane']
Predicted: airplane (Confidence: 4.15%)

Test Result 235:
Image: 006302.jpg
True labels: ['boat', 'boat', 'boat', 'boat']
Predicted: airplane (Confidence: 4.62%)

Test Result 236:
Image: 006332.jpg
True labels: ['boat']
Predicted: airplane (Confidence: 2.50%)

Test Result 237:
Image: 006347.jpg
True labels: ['airplane', 'airplane', 'airplane']
Predicted: airplane (Confidence: 0.43%)

Test Result 238:
Image: 006406.jpg
True labels: ['airplane', 'airplane', 'airplane', 'airplane', 'airplane']
Predicted: airplane (Confidence: 2.36%)

Test Result 239:
Image: 006408.jpg
True labels: ['boat']
Predicted: airplane (Confidence: 5.97%)

Test Result 240:
Image: 006490.jpg
True labels: ['boat', 'boat', 'boat']
Predicted: airplane (Confidence: 50.40%)

Test Result 241:
Image: 006528.jpg
True labels: ['boat']
Predicted: airplane (Confidence: 0.79%)

Test Result 242:
Image: 006533.jpg
True labels: ['airplane']
Predicted: airplane (Confidence: 4.11%)

Test Result 243:
Image: 006577.jpg
True labels: ['airplane']
Predicted: airplane (Confidence: 2.02%)

Test Result 244:
Image: 006592.jpg
True labels: ['airplane']
Predicted: airplane (Confidence: 1.32%)

Test Result 245:
Image: 006604.jpg
True labels: ['boat', 'boat', 'boat', 'boat']
Predicted: airplane (Confidence: 3.36%)

Test Result 246:
Image: 006651.jpg
True labels: ['airplane']
Predicted: airplane (Confidence: 1.08%)

Test Result 247:
Image: 006659.jpg
True labels: ['boat']
Predicted: airplane (Confidence: 9.10%)

Test Result 248:
Image: 006713.jpg
True labels: ['airplane']
Predicted: airplane (Confidence: 2.20%)

Test Result 249:
Image: 006752.jpg
True labels: ['airplane']
Predicted: airplane (Confidence: 2.34%)

Test Result 250:
Image: 006758.jpg
True labels: ['airplane', 'airplane', 'airplane']
Predicted: airplane (Confidence: 20.55%)

Test Result 251:
Image: 006888.jpg
True labels: ['airplane', 'airplane']
Predicted: airplane (Confidence: 14.37%)

Test Result 252:
Image: 006946.jpg
True labels: ['airplane']
Predicted: airplane (Confidence: 2.90%)

Test Result 253:
Image: 006977.jpg
True labels: ['boat']
Predicted: airplane (Confidence: 1.37%)

Test Result 254:

Image: 006997.jpg
True labels: ['boat']
Predicted: airplane (Confidence: 7.64%)

Test Result 255:
Image: 007026.jpg
True labels: ['airplane', 'airplane', 'airplane']
Predicted: airplane (Confidence: 8.74%)

Test Result 256:
Image: 007066.jpg
True labels: ['boat']
Predicted: airplane (Confidence: 2.10%)

Test Result 257:
Image: 007091.jpg
True labels: ['boat']
Predicted: airplane (Confidence: 1.54%)

Test Result 258:
Image: 007096.jpg
True labels: ['airplane']
Predicted: airplane (Confidence: 2.67%)

Test Result 259:
Image: 007135.jpg
True labels: ['airplane']
Predicted: airplane (Confidence: 1.96%)

Test Result 260:
Image: 007155.jpg
True labels: ['airplane']
Predicted: airplane (Confidence: 6.80%)

Test Result 261:
Image: 007157.jpg
True labels: ['airplane', 'airplane']
Predicted: airplane (Confidence: 0.63%)

Test Result 262:
Image: 007164.jpg
True labels: ['airplane']
Predicted: airplane (Confidence: 4.89%)

Test Result 263:
Image: 007169.jpg
True labels: ['airplane']
Predicted: airplane (Confidence: 1.45%)

Test Result 264:
Image: 007173.jpg
True labels: ['boat']
Predicted: airplane (Confidence: 26.00%)

Test Result 265:
Image: 007233.jpg

True labels: ['boat']
Predicted: airplane (Confidence: 0.70%)

Test Result 266:
Image: 007257.jpg
True labels: ['airplane']
Predicted: airplane (Confidence: 34.33%)

Test Result 267:
Image: 007262.jpg
True labels: ['airplane']
Predicted: airplane (Confidence: 6.77%)

Test Result 268:
Image: 007267.jpg
True labels: ['airplane']
Predicted: airplane (Confidence: 4.12%)

Test Result 269:
Image: 007290.jpg
True labels: ['airplane', 'airplane', 'airplane']
Predicted: airplane (Confidence: 3.71%)

Test Result 270:
Image: 007341.jpg
True labels: ['airplane']
Predicted: airplane (Confidence: 3.83%)

Test Result 271:
Image: 007352.jpg
True labels: ['airplane', 'airplane']
Predicted: airplane (Confidence: 27.31%)

Test Result 272:
Image: 007357.jpg
True labels: ['boat', 'boat', 'boat']
Predicted: airplane (Confidence: 11.80%)

Test Result 273:
Image: 007367.jpg
True labels: ['airplane', 'airplane']
Predicted: airplane (Confidence: 0.36%)

Test Result 274:
Image: 007377.jpg
True labels: ['boat']
Predicted: airplane (Confidence: 2.00%)

Test Result 275:
Image: 007400.jpg
True labels: ['boat']
Predicted: airplane (Confidence: 1.32%)

Test Result 276:
Image: 007403.jpg
True labels: ['airplane']

Predicted: airplane (Confidence: 0.77%)

Test Result 277:

Image: 007415.jpg

True labels: ['boat']

Predicted: airplane (Confidence: 5.02%)

Test Result 278:

Image: 007428.jpg

True labels: ['boat']

Predicted: airplane (Confidence: 49.30%)

Test Result 279:

Image: 007455.jpg

True labels: ['boat']

Predicted: airplane (Confidence: 7.13%)

Test Result 280:

Image: 007464.jpg

True labels: ['boat', 'boat']

Predicted: airplane (Confidence: 4.53%)

Test Result 281:

Image: 007502.jpg

True labels: ['airplane']

Predicted: airplane (Confidence: 0.84%)

Test Result 282:

Image: 007548.jpg

True labels: ['airplane', 'airplane']

Predicted: airplane (Confidence: 0.61%)

Test Result 283:

Image: 007632.jpg

True labels: ['airplane']

Predicted: airplane (Confidence: 8.55%)

Test Result 284:

Image: 007635.jpg

True labels: ['boat']

Predicted: airplane (Confidence: 11.59%)

Test Result 285:

Image: 007660.jpg

True labels: ['boat']

Predicted: airplane (Confidence: 7.64%)

Test Result 286:

Image: 007676.jpg

True labels: ['airplane']

Predicted: airplane (Confidence: 7.44%)

Test Result 287:

Image: 007711.jpg

True labels: ['airplane']

Predicted: airplane (Confidence: 22.19%)

Test Result 288:
Image: 007788.jpg
True labels: ['boat']
Predicted: airplane (Confidence: 10.23%)

Test Result 289:
Image: 007806.jpg
True labels: ['airplane', 'airplane']
Predicted: airplane (Confidence: 4.67%)

Test Result 290:
Image: 007825.jpg
True labels: ['airplane']
Predicted: airplane (Confidence: 1.74%)

Test Result 291:
Image: 007837.jpg
True labels: ['airplane', 'airplane']
Predicted: airplane (Confidence: 11.45%)

Test Result 292:
Image: 007870.jpg
True labels: ['airplane']
Predicted: airplane (Confidence: 3.40%)

Test Result 293:
Image: 007952.jpg
True labels: ['boat']
Predicted: airplane (Confidence: 4.08%)

Test Result 294:
Image: 007973.jpg
True labels: ['airplane']
Predicted: airplane (Confidence: 2.78%)

Test Result 295:
Image: 007985.jpg
True labels: ['airplane']
Predicted: airplane (Confidence: 15.07%)

Test Result 296:
Image: 007990.jpg
True labels: ['boat']
Predicted: airplane (Confidence: 7.49%)

Test Result 297:
Image: 007993.jpg
True labels: ['airplane', 'airplane']
Predicted: airplane (Confidence: 4.98%)

Test Result 298:
Image: 007995.jpg
True labels: ['airplane']
Predicted: airplane (Confidence: 6.02%)

Test Result 299:
Image: 008055.jpg
True labels: ['airplane', 'airplane', 'airplane', 'airplane', 'airplane', 'airplane', 'airplane', 'airplane', 'airplane']
Predicted: airplane (Confidence: 1.23%)

Test Result 300:
Image: 008099.jpg
True labels: ['boat', 'boat']
Predicted: airplane (Confidence: 0.57%)

Test Result 301:
Image: 008104.jpg
True labels: ['airplane', 'airplane']
Predicted: airplane (Confidence: 3.07%)

Test Result 302:
Image: 008145.jpg
True labels: ['boat', 'boat']
Predicted: airplane (Confidence: 10.25%)

Test Result 303:
Image: 008217.jpg
True labels: ['boat']
Predicted: airplane (Confidence: 4.60%)

Test Result 304:
Image: 008219.jpg
True labels: ['boat']
Predicted: airplane (Confidence: 4.78%)

Test Result 305:
Image: 008249.jpg
True labels: ['boat']
Predicted: airplane (Confidence: 21.73%)

Test Result 306:
Image: 008278.jpg
True labels: ['boat', 'boat']
Predicted: airplane (Confidence: 10.11%)

Test Result 307:
Image: 008352.jpg
True labels: ['airplane']
Predicted: airplane (Confidence: 2.03%)

Test Result 308:
Image: 008369.jpg
True labels: ['boat']
Predicted: airplane (Confidence: 14.10%)

Test Result 309:
Image: 008373.jpg
True labels: ['boat', 'boat', 'boat', 'boat']
Predicted: airplane (Confidence: 4.74%)

Test Result 310:
Image: 008440.jpg
True labels: ['airplane']
Predicted: airplane (Confidence: 53.43%)

Test Result 311:
Image: 008459.jpg
True labels: ['airplane', 'airplane', 'airplane', 'airplane', 'airplane', 'airplane']
Predicted: airplane (Confidence: 32.96%)

Test Result 312:
Image: 008537.jpg
True labels: ['airplane']
Predicted: airplane (Confidence: 58.23%)

Test Result 313:
Image: 008545.jpg
True labels: ['boat', 'boat']
Predicted: airplane (Confidence: 0.37%)

Test Result 314:
Image: 008554.jpg
True labels: ['airplane', 'airplane']
Predicted: airplane (Confidence: 4.53%)

Test Result 315:
Image: 008571.jpg
True labels: ['boat']
Predicted: airplane (Confidence: 7.11%)

Test Result 316:
Image: 008578.jpg
True labels: ['boat']
Predicted: airplane (Confidence: 0.76%)

Test Result 317:
Image: 008590.jpg
True labels: ['boat', 'boat', 'boat', 'boat']
Predicted: airplane (Confidence: 3.81%)

Test Result 318:
Image: 008625.jpg
True labels: ['airplane']
Predicted: airplane (Confidence: 9.24%)

Test Result 319:
Image: 008643.jpg
True labels: ['boat', 'boat']
Predicted: airplane (Confidence: 0.85%)

Test Result 320:
Image: 008704.jpg
True labels: ['boat']
Predicted: airplane (Confidence: 4.00%)

Test Result 321:
Image: 008714.jpg
True labels: ['airplane']
Predicted: airplane (Confidence: 3.15%)

Test Result 322:
Image: 008761.jpg
True labels: ['airplane']
Predicted: airplane (Confidence: 1.83%)

Test Result 323:
Image: 008791.jpg
True labels: ['airplane']
Predicted: airplane (Confidence: 1.83%)

Test Result 324:
Image: 008820.jpg
True labels: ['boat']
Predicted: airplane (Confidence: 1.57%)

Test Result 325:
Image: 008825.jpg
True labels: ['airplane']
Predicted: airplane (Confidence: 16.12%)

Test Result 326:
Image: 008861.jpg
True labels: ['airplane']
Predicted: airplane (Confidence: 13.07%)

Test Result 327:
Image: 008868.jpg
True labels: ['boat', 'boat']
Predicted: airplane (Confidence: 4.50%)

Test Result 328:
Image: 008869.jpg
True labels: ['boat', 'boat', 'boat', 'boat', 'boat', 'boat']
Predicted: airplane (Confidence: 1.95%)

Test Result 329:
Image: 008881.jpg
True labels: ['airplane']
Predicted: airplane (Confidence: 7.91%)

Test Result 330:
Image: 008908.jpg
True labels: ['airplane']
Predicted: airplane (Confidence: 23.96%)

Test Result 331:
Image: 008950.jpg
True labels: ['airplane']
Predicted: airplane (Confidence: 0.45%)

Test Result 332:

Image: 008986.jpg
True labels: ['airplane', 'airplane']
Predicted: airplane (Confidence: 7.93%)

Test Result 333:
Image: 009001.jpg
True labels: ['boat']
Predicted: airplane (Confidence: 1.82%)

Test Result 334:
Image: 009026.jpg
True labels: ['boat']
Predicted: airplane (Confidence: 0.75%)

Test Result 335:
Image: 009030.jpg
True labels: ['airplane']
Predicted: airplane (Confidence: 4.57%)

Test Result 336:
Image: 009076.jpg
True labels: ['airplane']
Predicted: airplane (Confidence: 1.77%)

Test Result 337:
Image: 009083.jpg
True labels: ['airplane']
Predicted: airplane (Confidence: 5.25%)

Test Result 338:
Image: 009102.jpg
True labels: ['boat']
Predicted: airplane (Confidence: 10.14%)

Test Result 339:
Image: 009120.jpg
True labels: ['boat']
Predicted: airplane (Confidence: 15.83%)

Test Result 340:
Image: 009137.jpg
True labels: ['boat', 'boat', 'boat', 'boat', 'boat', 'boat', 'boat']
Predicted: airplane (Confidence: 3.60%)

Test Result 341:
Image: 009167.jpg
True labels: ['airplane']
Predicted: airplane (Confidence: 5.59%)

Test Result 342:
Image: 009211.jpg
True labels: ['airplane']
Predicted: airplane (Confidence: 9.10%)

Test Result 343:
Image: 009240.jpg

True labels: ['boat']
Predicted: airplane (Confidence: 8.77%)

Test Result 344:
Image: 009262.jpg
True labels: ['airplane']
Predicted: airplane (Confidence: 0.97%)

Test Result 345:
Image: 009292.jpg
True labels: ['airplane']
Predicted: airplane (Confidence: 2.75%)

Test Result 346:
Image: 009329.jpg
True labels: ['airplane']
Predicted: airplane (Confidence: 6.34%)

Test Result 347:
Image: 009332.jpg
True labels: ['airplane', 'airplane']
Predicted: airplane (Confidence: 6.53%)

Test Result 348:
Image: 009356.jpg
True labels: ['airplane', 'airplane', 'airplane', 'airplane', 'airplane', 'airplane', 'airplane', 'airplane', 'airplane']
Predicted: airplane (Confidence: 8.77%)

Test Result 349:
Image: 009423.jpg
True labels: ['boat']
Predicted: airplane (Confidence: 4.70%)

Test Result 350:
Image: 009538.jpg
True labels: ['boat', 'boat']
Predicted: airplane (Confidence: 9.77%)

Test Result 351:
Image: 009727.jpg
True labels: ['boat', 'boat', 'boat']
Predicted: airplane (Confidence: 4.97%)

Test Result 352:
Image: 009728.jpg
True labels: ['boat']
Predicted: airplane (Confidence: 2.81%)

Test Result 353:
Image: 009793.jpg
True labels: ['boat']
Predicted: airplane (Confidence: 16.17%)

```
Test Result 354:  
Image: 009824.jpg  
True labels: ['airplane']  
Predicted: airplane (Confidence: 18.49%)
```

```
Test Result 355:  
Image: 009835.jpg  
True labels: ['airplane']  
Predicted: airplane (Confidence: 11.06%)
```

```
Test Result 356:  
Image: 009838.jpg  
True labels: ['airplane']  
Predicted: airplane (Confidence: 8.93%)
```

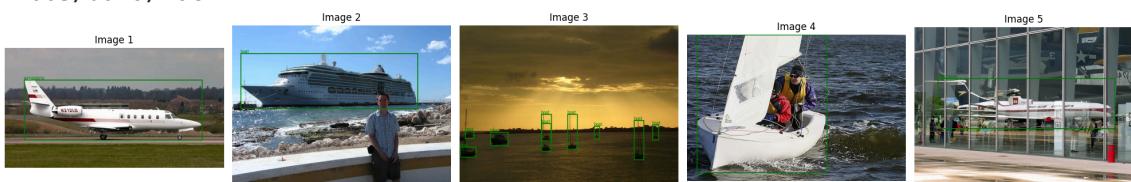
```
Test Result 357:  
Image: 009876.jpg  
True labels: ['airplane']  
Predicted: airplane (Confidence: 13.28%)
```

```
Test Result 358:  
Image: 009885.jpg  
True labels: ['boat']  
Predicted: airplane (Confidence: 2.23%)
```

```
Test Result 359:  
Image: 009919.jpg  
True labels: ['airplane']  
Predicted: airplane (Confidence: 7.41%)
```

```
Test Summary:  
Total images tested: 359  
Correct predictions: 191  
Accuracy: 53.20%
```

```
Visualizing detections...  
Using downloaded and verified file: /scratch/poh2005/data/VOC/VOCtest_06-Nov-2007.tar  
Extracting /scratch/poh2005/data/VOC/VOCtest_06-Nov-2007.tar to /scratch/poh2005/data/VOC
```



```
In [58]: import torch  
import torchvision  
from torchvision.models.detection import ssdlite320_mobilenet_v3_large  
import torch.nn as nn  
from torchvision import transforms  
import onnx  
import onnxruntime as ort  
import numpy as np  
from PIL import Image
```

```

import matplotlib.pyplot as plt
import os
from torch.utils.data import Dataset, DataLoader
import matplotlib.patches as patches
from matplotlib.patches import Rectangle
from torch.cuda.amp import autocast, GradScaler
import time
from tqdm import tqdm

# Define paths
OPEN_IMAGES_ROOT = '/scratch/poh2005/data/open_images_temp'

class OpenImagesDataset(Dataset):
    def __init__(self, root_dir, split='test', transform=None, max_samples=None):
        self.root_dir = root_dir
        self.split = split
        self.transform = transform
        self._cache = {} # Add caching for faster data loading

        print(f"Initializing {split} dataset...")

    # Load class descriptions
    self.class_descriptions = {}
    class_desc_path = os.path.join(root_dir, 'class-descriptions-boxable')
    print(f"Loading class descriptions from {class_desc_path}")
    with open(class_desc_path, 'r') as f:
        for line in f:
            class_id, class_name = line.strip().split(',')
            self.class_descriptions[class_id] = class_name

    # Load annotations with efficient file reading
    self.annotations = []
    annotation_file = os.path.join(root_dir, f'sub-{split}-annotations-bboxable')
    print(f"Loading annotations from {annotation_file}")

    with open(annotation_file, 'r') as f:
        next(f) # Skip header
        # Read all lines at once for efficiency
        lines = f.readlines()
        for line in lines:
            parts = line.strip().split(',')
            image_id = parts[0]
            label = parts[2]
            xmin, ymin, xmax, ymax = map(float, parts[4:8])

            # Validate box coordinates
            if xmax <= xmin or ymax <= ymin:
                continue

            # Only keep airplane and boat annotations
            class_name = self.class_descriptions[label].lower()
            if class_name in ['airplane', 'boat']:
                self.annotations.append({
                    'image_id': image_id,
                    'label': 1 if class_name == 'airplane' else 2, # 1=airplane, 2=boat
                    'bbox': [xmin, ymin, xmax, ymax]
                })

```

```

        })

# Get unique image IDs efficiently using set comprehension
self.image_ids = list({ann['image_id'] for ann in self.annotations})

# Limit samples if specified
if max_samples and max_samples < len(self.image_ids):
    self.image_ids = self.image_ids[:max_samples]
    self.annotations = [ann for ann in self.annotations
                        if ann['image_id'] in set(self.image_ids)]

print(f"Found {len(self.image_ids)} images with {len(self.annotations)} annotations")
print(f"Dataset initialization completed for {split} split")

def __len__(self):
    return len(self.image_ids)

def __getitem__(self, idx):
    # Check cache first
    if idx in self._cache:
        return self._cache[idx]

    image_id = self.image_ids[idx]

    # Load image using Pillow-SIMD for faster loading
    img_path = os.path.join(self.root_dir, self.split, image_id + '.jpg')
    try:
        image = Image.open(img_path).convert('RGB')
    except Exception as e:
        print(f"Error loading image {img_path}: {e}")
        # Return a dummy image and empty target if loading fails
        image = Image.new('RGB', (320, 320))
    target = {
        'boxes': torch.zeros((0, 4), dtype=torch.float32),
        'labels': torch.zeros(0, dtype=torch.int64),
        'image_id': torch.tensor([idx]),
        'area': torch.zeros(0, dtype=torch.float32),
        'iscrowd': torch.zeros(0, dtype=torch.int64)
    }
    if self.transform:
        image = self.transform(image)
    return image, target

# Get annotations efficiently using list comprehension
img_annotations = [ann for ann in self.annotations if ann['image_id']]

# Prepare boxes and labels efficiently
boxes = []
labels = []

for ann in img_annotations:
    xmin, ymin, xmax, ymax = ann['bbox']
    # Convert normalized coordinates to absolute
    xmin = xmin * width
    xmax = xmax * width

```

```

        ymin = ymin * height
        ymax = ymax * height

        # Validate box coordinates
        if xmax > xmin and ymax > ymin:
            boxes.append([xmin, ymin, xmax, ymax])
            labels.append(ann['label'])

        # Convert to tensors efficiently
        if not boxes:
            boxes = torch.zeros((0, 4), dtype=torch.float32)
            labels = torch.zeros(0, dtype=torch.int64)
        else:
            boxes = torch.tensor(boxes, dtype=torch.float32)
            labels = torch.tensor(labels, dtype=torch.int64)

        # Create target dictionary
        target = {
            'boxes': boxes,
            'labels': labels,
            'image_id': torch.tensor([idx]),
            'area': (boxes[:, 3] - boxes[:, 1]) * (boxes[:, 2] - boxes[:, 0]),
            'iscrowd': torch.zeros((len(boxes)),), dtype=torch.int64)
        }

        if self.transform:
            image = self.transform(image)

        # Cache the result
        result = (image, target)
        self._cache[idx] = result
        return result

    def create_data_loaders(batch_size=32):
        """Create optimized train, validation and test dataloaders"""
        # Define transformations
        transform = transforms.Compose([
            transforms.Resize((320, 320)),
            transforms.ToTensor(),
            transforms.Normalize(mean=[0.485, 0.456, 0.406], std=[0.229, 0.224,
        ]))

        print("Creating datasets...")

        # Create datasets with progress reporting
        train_dataset = OpenImagesDataset(
            root_dir=OPEN_IMAGES_ROOT,
            split='train',
            transform=transform,
            max_samples=2000
        )

        val_dataset = OpenImagesDataset(
            root_dir=OPEN_IMAGES_ROOT,
            split='validation',
            transform=transform,

```

```

        max_samples=400
    )

test_dataset = OpenImagesDataset(
    root_dir=OPEN_IMAGES_ROOT,
    split='test',
    transform=transform,
    max_samples=800
)

# Calculate optimal number of workers
num_workers = min(4, os.cpu_count() or 1)

# Create optimized DataLoaders
train_loader = DataLoader(
    train_dataset,
    batch_size=batch_size,
    shuffle=True,
    num_workers=num_workers,
    pin_memory=True,
    prefetch_factor=2,
    persistent_workers=True,
    collate_fn=collate_fn
)

val_loader = DataLoader(
    val_dataset,
    batch_size=batch_size,
    shuffle=False,
    num_workers=num_workers,
    pin_memory=True,
    persistent_workers=True,
    collate_fn=collate_fn
)

test_loader = DataLoader(
    test_dataset,
    batch_size=batch_size,
    shuffle=False,
    num_workers=num_workers,
    pin_memory=True,
    persistent_workers=True,
    collate_fn=collate_fn
)

print(f'Created dataloaders with batch size {batch_size} and {num_workers}')
return train_loader, val_loader, test_loader

def collate_fn(batch):
    return tuple(zip(*batch))

def train_one_epoch(model, dataloader, optimizer, scheduler, scaler, device,
model.train()
total_loss = 0

progress_bar = tqdm(dataloader, desc=f'Epoch {epoch}/{num_epochs}')

```

```

for images, targets in progress_bar:
    # Move data to GPU
    images = list(image.to(device) for image in images)
    targets = [{k: v.to(device) for k, v in t.items()} for t in targets]

    # Use automatic mixed precision
    with torch.amp.autocast('cuda'): # Fixed deprecated warning
        loss_dict = model(images, targets)
        losses = sum(loss for loss in loss_dict.values())

    # Optimize
    optimizer.zero_grad(set_to_none=True)
    scaler.scale(losses).backward()
    scaler.step(optimizer)
    scaler.update()

    # Update learning rate
    scheduler.step()

    # Update metrics
    total_loss += losses.item()

    # Update progress bar
    progress_bar.set_postfix({'loss': losses.item()})

return total_loss / len(dataloader)

def validate(model, dataloader, device):
    model.eval()
    total_loss = 0

    with torch.no_grad():
        progress_bar = tqdm(dataloader, desc='Validation')
        for images, targets in progress_bar:
            images = list(image.to(device, non_blocking=True) for image in images)
            targets = [{k: v.to(device, non_blocking=True) for k, v in t.items()} for t in targets]

            with torch.amp.autocast('cuda'):
                # During validation, run in training mode just for loss computation
                model.train()
                loss_dict = model(images, targets)
                model.eval() # Set back to eval mode

                losses = sum(loss for loss in loss_dict.values())

            total_loss += losses.item()
            progress_bar.set_postfix({'loss': losses.item()})

return total_loss / len(dataloader)

def fine_tune_model(model, train_loader, val_loader, num_epochs=150):
    """Optimized fine-tuning for V100 GPU"""
    device = torch.device('cuda')
    print(f"Using device: {device}")
    print(f"GPU: {torch.cuda.get_device_name(0)}")

```

```

# Enable gradient scaler for mixed precision training
scaler = torch.amp.GradScaler() # Fixed deprecated warning

# Optimize training parameters
params = [p for p in model.parameters() if p.requires_grad]
optimizer = torch.optim.AdamW(
    params,
    lr=0.001,
    weight_decay=0.0005,
    eps=1e-8,
    betas=(0.9, 0.999),
)

# Learning rate scheduler with warmup
scheduler = torch.optim.lr_scheduler.OneCycleLR(
    optimizer,
    max_lr=0.001,
    epochs=num_epochs,
    steps_per_epoch=len(train_loader),
    pct_start=0.1,
    div_factor=25,
    final_div_factor=1000
)

best_val_loss = float('inf')
train_times = []

for epoch in range(num_epochs):
    epoch_start_time = time.time()

    # Training phase
    train_loss = train_one_epoch(
        model, train_loader, optimizer, scheduler,
        scaler, device, epoch + 1, num_epochs
    )

    epoch_time = time.time() - epoch_start_time
    train_times.append(epoch_time)

    print(f'\nEpoch {epoch + 1}:')
    print(f'Training Loss: {train_loss:.4f}')
    print(f'Time: {epoch_time:.2f} seconds')
    print(f'Learning rate: {scheduler.get_last_lr()[0]:.6f}')
    print(f'GPU memory: {torch.cuda.memory_allocated()/1024**3:.2f} GB')

    # Validation phase (every 5 epochs)
    if (epoch + 1) % 5 == 0:
        val_loss = validate(model, val_loader, device)
        print(f'Validation Loss: {val_loss:.4f}')

        if val_loss < best_val_loss:
            best_val_loss = val_loss
            print(f'Saving best model (val_loss: {val_loss:.4f})')
            torch.save({
                'epoch': epoch,
                'model_state_dict': model.state_dict(),
            })

```

```

        'optimizer_state_dict': optimizer.state_dict(),
        'loss': best_val_loss,
    }, 'best_model.pth')

    return model, train_times

def main():
    # Enable GPU optimizations for V100
    torch.backends.cudnn.benchmark = True
    torch.backends.cudnn.enabled = True

    # Set environment variables for optimal performance
    os.environ['CUDA_LAUNCH_BLOCKING'] = '0' # Async CUDA operations
    os.environ['CUDA_VISIBLE_DEVICES'] = '0' # Ensure using single GPU effi

    try:
        print("Loading pretrained SSDLite320 MobileNetV3 model...")
        model = ssdlite320_mobilenet_v3_large(pretrained=True)

        # Move model to GPU
        device = torch.device('cuda')
        model = model.to(device)

        # Create dataloaders with optimized settings for V100
        print("\nCreating datasets...")
        train_loader, val_loader, test_loader = create_data_loaders(batch_si

        # Clear GPU cache and show memory info
        torch.cuda.empty_cache()
        print(f"Initial GPU memory allocated: {torch.cuda.memory_allocated(0)}")
        print(f"GPU memory reserved: {torch.cuda.memory_reserved(0)/1024**3}")

        # Start training with proper mixed precision setup
        print("\nStarting fine-tuning...")
        model, train_times = fine_tune_model(model, train_loader, val_loader)

        # Print training statistics
        print("\nTraining completed!")
        print(f"Average epoch time: {sum(train_times)/len(train_times):.2f}")
        print(f"Total training time: {sum(train_times):.2f} seconds")

        # Export model to ONNX with optimizations
        print("\nExporting to ONNX format...")
        try:
            model.eval()
            dummy_input = torch.randn(1, 3, 320, 320, device=device)

            torch.onnx.export(
                model,
                dummy_input,
                'ssd_model.onnx',
                export_params=True,
                opset_version=15,
                do_constant_folding=True,
                input_names=['input'],

```

```

        output_names=['boxes', 'scores', 'labels'],
        dynamic_axes={
            'input': {0: 'batch_size'},
            'boxes': {0: 'batch_size'},
            'scores': {0: 'batch_size'},
            'labels': {0: 'batch_size'}
        }
    )

    print("\nVerifying ONNX model...")
    onnx_model = onnx.load('ssd_model.onnx')
    onnx.checker.check_model(onnx_model)

except Exception as e:
    print(f"An error occurred during ONNX export: {str(e)}")
    print("Continuing with the rest of the program...")

return model

except Exception as e:
    print(f"An error occurred in main execution: {str(e)}")
    raise

def visualize_detections(model, dataset, num_images=5, confidence_threshold=0.5):
    """Visualize detection results"""
    device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
    model.eval()

    # Create figure
    fig, axes = plt.subplots(1, num_images, figsize=(20, 4))
    if num_images == 1:
        axes = [axes]

    # Define colors and class names
    colors = ['red', 'blue'] # airplane = red, boat = blue
    class_names = ['airplane', 'boat']

    with torch.no_grad(), autocast():
        for idx in range(min(num_images, len(dataset))):
            image, target = dataset[idx]

            # Denormalize image for display
            img_show = image.permute(1, 2, 0).numpy()
            img_show = img_show * np.array([0.229, 0.224, 0.225]) + np.array([0.485, 0.451, 0.406])
            img_show = np.clip(img_show, 0, 1)

            # Get predictions
            image = image.unsqueeze(0).to(device)
            predictions = model(image)

            # Plot image
            axes[idx].imshow(img_show)

            # Draw predicted boxes
            boxes = predictions[0]['boxes'].cpu().numpy()

```

```

scores = predictions[0]['scores'].cpu().numpy()
labels = predictions[0]['labels'].cpu().numpy()

# Draw boxes for predictions with score > threshold
for box, score, label in zip(boxes, scores, labels):
    if score > confidence_threshold:
        x1, y1, x2, y2 = box
        rect = patches.Rectangle(
            (x1, y1),
            x2 - x1,
            y2 - y1,
            linewidth=2,
            edgecolor=colors[label-1],
            facecolor='none'
        )
        axes[idx].add_patch(rect)

# Add label and score
label_text = f'{class_names[label-1]} {score:.2f}'
axes[idx].text(
    x1, y1-5,
    label_text,
    color=colors[label-1],
    fontsize=8,
    bbox=dict(facecolor='white', alpha=0.8, edgecolor='r'
)

# Draw ground truth boxes
gt_boxes = target['boxes'].cpu().numpy()
gt_labels = target['labels'].cpu().numpy()

for box, label in zip(gt_boxes, gt_labels):
    x1, y1, x2, y2 = box
    rect = patches.Rectangle(
        (x1, y1),
        x2 - x1,
        y2 - y1,
        linewidth=1,
        edgecolor=colors[label-1],
        facecolor='none',
        linestyle='--'
    )
    axes[idx].add_patch(rect)

axes[idx].axis('off')
axes[idx].set_title(f'Image {idx+1}')

plt.tight_layout()
plt.savefig('detection_results.png')
plt.show()

def evaluate_model(model, test_loader, device):
    """Evaluate model performance"""
    model.eval()
    total_correct = 0
    total_objects = 0

```

```

class_correct = {1: 0, 2: 0} # 1: airplane, 2: boat
class_total = {1: 0, 2: 0}

print("\nEvaluating model performance...")

with torch.no_grad(), autocast():
    for images, targets in tqdm(test_loader, desc='Evaluation'):
        images = list(image.to(device) for image in images)
        targets = [{k: v.to(device) for k, v in t.items()} for t in targets]

        # Get predictions
        outputs = model(images)

        # Calculate metrics
        for output, target in zip(outputs, targets):
            pred_labels = output['labels'][output['scores'] > 0.3]
            true_labels = target['labels']

            # Update overall metrics
            correct_predictions = len(set(pred_labels.tolist()) & set(true_labels))
            total_correct += correct_predictions
            total_objects += len(true_labels)

            # Update per-class metrics
            for label in [1, 2]: # 1: airplane, 2: boat
                class_pred = (pred_labels == label).sum().item()
                class_true = (true_labels == label).sum().item()
                class_correct[label] += min(class_pred, class_true)
                class_total[label] += class_true

# Calculate metrics
overall_accuracy = total_correct / total_objects if total_objects > 0 else 0
class_accuracies = {
    label: class_correct[label] / class_total[label]
    if class_total[label] > 0 else 0
    for label in [1, 2]
}

# Print results
print("\nEvaluation Results:")
print(f"Overall Accuracy: {overall_accuracy:.4f}")
print("\nPer-class Accuracy:")
print(f"Airplane: {class_accuracies[1]:.4f}")
print(f"Boat: {class_accuracies[2]:.4f}")

return overall_accuracy, class_accuracies

if __name__ == '__main__':
    model = main()

    # Get device
    device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')

    # Create dataloaders for visualization and evaluation
    _, _, test_loader = create_data_loaders(batch_size=32)

```

```

# Visualize some predictions
print("\nGenerating visualizations...")
visualize_detections(model, test_loader.dataset, num_images=5)

# Evaluate model
print("\nEvaluating model...")
overall_acc, class_accs = evaluate_model(model, test_loader, device)

print("\nTraining completed successfully!")

```

Loading pretrained SSDLite320 MobileNetV3 model...

```

Creating datasets...
Creating datasets...
Initializing train dataset...
Loading class descriptions from /scratch/poh2005/data/open_images_temp/class-descriptions-boxable.csv
Loading annotations from /scratch/poh2005/data/open_images_temp/sub-train-annotations-bbox.csv
Found 2000 images with 4460 annotations
Dataset initialization completed for train split
Initializing validation dataset...
Loading class descriptions from /scratch/poh2005/data/open_images_temp/class-descriptions-boxable.csv
Loading annotations from /scratch/poh2005/data/open_images_temp/sub-validation-annotations-bbox.csv
Found 327 images with 444 annotations
Dataset initialization completed for validation split
Initializing test dataset...
Loading class descriptions from /scratch/poh2005/data/open_images_temp/class-descriptions-boxable.csv
Loading annotations from /scratch/poh2005/data/open_images_temp/sub-test-annotations-bbox.csv
Found 800 images with 1074 annotations
Dataset initialization completed for test split
Created dataloaders with batch size 32 and 4 workers
Initial GPU memory allocated: 0.14 GB
GPU memory reserved: 0.22 GB

Starting fine-tuning...
Using device: cuda
GPU: Tesla V100-PCIE-32GB
Epoch 1/150: 100%|██████████| 63/63 [00:15<00:00,  3.98it/s, loss=30.9]
Epoch 1:
Training Loss: 38.1886
Time: 15.82 seconds
Learning rate: 0.000051
GPU memory: 0.18 GB
Epoch 2/150: 100%|██████████| 63/63 [00:12<00:00,  5.24it/s, loss=22.3]
Epoch 2:
Training Loss: 34.2630
Time: 12.03 seconds
Learning rate: 0.000082
GPU memory: 0.18 GB
Epoch 3/150: 100%|██████████| 63/63 [00:11<00:00,  5.28it/s, loss=24.1]

```

Epoch 3:
Training Loss: 30.2104
Time: 11.95 seconds
Learning rate: 0.000132
GPU memory: 0.18 GB

Epoch 4/150: 100%|██████████| 63/63 [00:11<00:00, 5.35it/s, loss=16.3]

Epoch 4:
Training Loss: 24.1866
Time: 11.77 seconds
Learning rate: 0.000199
GPU memory: 0.18 GB

Epoch 5/150: 100%|██████████| 63/63 [00:11<00:00, 5.41it/s, loss=15.1]

Epoch 5:
Training Loss: 15.7854
Time: 11.65 seconds
Learning rate: 0.000280
GPU memory: 0.18 GB

Validation: 100%|██████████| 11/11 [00:03<00:00, 3.46it/s, loss=47.9]

Validation Loss: 23.4797
Saving best model (val_loss: 23.4797)

Epoch 6/150: 100%|██████████| 63/63 [00:12<00:00, 5.22it/s, loss=10.8]

Epoch 6:
Training Loss: 13.4551
Time: 12.08 seconds
Learning rate: 0.000372
GPU memory: 0.18 GB

Epoch 7/150: 100%|██████████| 63/63 [00:11<00:00, 5.36it/s, loss=16.1]

Epoch 7:
Training Loss: 12.1727
Time: 11.75 seconds
Learning rate: 0.000471
GPU memory: 0.18 GB

Epoch 8/150: 100%|██████████| 63/63 [00:11<00:00, 5.29it/s, loss=10]

Epoch 8:
Training Loss: 11.6690
Time: 11.92 seconds
Learning rate: 0.000571
GPU memory: 0.18 GB

Epoch 9/150: 100%|██████████| 63/63 [00:11<00:00, 5.32it/s, loss=11.5]

Epoch 9:
Training Loss: 10.8545
Time: 11.84 seconds
Learning rate: 0.000669
GPU memory: 0.18 GB

Epoch 10/150: 100%|██████████| 63/63 [00:11<00:00, 5.37it/s, loss=6.28]

Epoch 10:
Training Loss: 10.0914
Time: 11.73 seconds
Learning rate: 0.000761
GPU memory: 0.18 GB

Validation: 100%|██████████| 11/11 [00:01<00:00, 8.89it/s, loss=53.7]

Validation Loss: 24.0954

Epoch 11/150: 100%|██████████| 63/63 [00:12<00:00, 5.20it/s, loss=14.4]

Epoch 11:
Training Loss: 9.8378
Time: 12.12 seconds
Learning rate: 0.000842
GPU memory: 0.18 GB

Epoch 12/150: 100% |██████████| 63/63 [00:11<00:00, 5.40it/s, loss=14.5]

Epoch 12:
Training Loss: 9.3496
Time: 11.68 seconds
Learning rate: 0.000909
GPU memory: 0.18 GB

Epoch 13/150: 100% |██████████| 63/63 [00:11<00:00, 5.43it/s, loss=15.8]

Epoch 13:
Training Loss: 9.3217
Time: 11.61 seconds
Learning rate: 0.000959
GPU memory: 0.18 GB

Epoch 14/150: 100% |██████████| 63/63 [00:11<00:00, 5.37it/s, loss=7]

Epoch 14:
Training Loss: 9.0529
Time: 11.74 seconds
Learning rate: 0.000990
GPU memory: 0.18 GB

Epoch 15/150: 100% |██████████| 63/63 [00:11<00:00, 5.45it/s, loss=5.93]

Epoch 15:
Training Loss: 8.2983
Time: 11.57 seconds
Learning rate: 0.001000
GPU memory: 0.18 GB

Validation: 100% |██████████| 11/11 [00:01<00:00, 8.76it/s, loss=56.1]

Validation Loss: 22.6602
Saving best model (val_loss: 22.6602)

Epoch 16/150: 100% |██████████| 63/63 [00:12<00:00, 5.22it/s, loss=7.27]

Epoch 16:
Training Loss: 7.8803
Time: 12.08 seconds
Learning rate: 0.001000
GPU memory: 0.18 GB

Epoch 17/150: 100% |██████████| 63/63 [00:11<00:00, 5.42it/s, loss=7.75]

Epoch 17:
Training Loss: 7.5393
Time: 11.63 seconds
Learning rate: 0.000999
GPU memory: 0.18 GB

Epoch 18/150: 100% |██████████| 63/63 [00:11<00:00, 5.41it/s, loss=12.8]

Epoch 18:
Training Loss: 7.1200
Time: 11.64 seconds
Learning rate: 0.000999
GPU memory: 0.18 GB

Epoch 19/150: 100% |██████████| 63/63 [00:11<00:00, 5.41it/s, loss=6.35]

Epoch 19:
Training Loss: 6.9839
Time: 11.65 seconds
Learning rate: 0.000998
GPU memory: 0.18 GB

Epoch 20/150: 100%|██████████| 63/63 [00:11<00:00, 5.41it/s, loss=8.39]

Epoch 20:
Training Loss: 6.7789
Time: 11.65 seconds
Learning rate: 0.000997
GPU memory: 0.18 GB

Validation: 100%|██████████| 11/11 [00:01<00:00, 8.70it/s, loss=50.7]

Validation Loss: 21.6931
Saving best model (val_loss: 21.6931)

Epoch 21/150: 100%|██████████| 63/63 [00:11<00:00, 5.27it/s, loss=4.79]

Epoch 21:
Training Loss: 6.2908
Time: 11.95 seconds
Learning rate: 0.000995
GPU memory: 0.18 GB

Epoch 22/150: 100%|██████████| 63/63 [00:11<00:00, 5.40it/s, loss=4.69]

Epoch 22:
Training Loss: 6.0422
Time: 11.68 seconds
Learning rate: 0.000993
GPU memory: 0.18 GB

Epoch 23/150: 100%|██████████| 63/63 [00:11<00:00, 5.40it/s, loss=6.27]

Epoch 23:
Training Loss: 5.9782
Time: 11.67 seconds
Learning rate: 0.000991
GPU memory: 0.18 GB

Epoch 24/150: 100%|██████████| 63/63 [00:11<00:00, 5.45it/s, loss=4.76]

Epoch 24:
Training Loss: 5.4595
Time: 11.56 seconds
Learning rate: 0.000989
GPU memory: 0.18 GB

Epoch 25/150: 100%|██████████| 63/63 [00:11<00:00, 5.38it/s, loss=4.13]

Epoch 25:
Training Loss: 5.1291
Time: 11.72 seconds
Learning rate: 0.000986
GPU memory: 0.18 GB

Validation: 100%|██████████| 11/11 [00:01<00:00, 8.70it/s, loss=43.2]

Validation Loss: 21.6150
Saving best model (val_loss: 21.6150)

Epoch 26/150: 100%|██████████| 63/63 [00:11<00:00, 5.28it/s, loss=4.36]

Epoch 26:
Training Loss: 4.9648
Time: 11.93 seconds
Learning rate: 0.000984
GPU memory: 0.18 GB

```
Epoch 27/150: 100%|██████████| 63/63 [00:11<00:00,  5.40it/s, loss=6.77]
Epoch 27:
Training Loss: 4.7725
Time: 11.66 seconds
Learning rate: 0.000981
GPU memory: 0.18 GB
Epoch 28/150: 100%|██████████| 63/63 [00:11<00:00,  5.43it/s, loss=5.2]
Epoch 28:
Training Loss: 4.7223
Time: 11.61 seconds
Learning rate: 0.000977
GPU memory: 0.18 GB
Epoch 29/150: 100%|██████████| 63/63 [00:11<00:00,  5.50it/s, loss=4.32]
Epoch 29:
Training Loss: 4.5001
Time: 11.46 seconds
Learning rate: 0.000974
GPU memory: 0.18 GB
Epoch 30/150: 100%|██████████| 63/63 [00:11<00:00,  5.45it/s, loss=5.92]
Epoch 30:
Training Loss: 4.2377
Time: 11.57 seconds
Learning rate: 0.000970
GPU memory: 0.18 GB
Validation: 100%|██████████| 11/11 [00:01<00:00,  8.53it/s, loss=49.2]
Validation Loss: 22.3131
Epoch 31/150: 100%|██████████| 63/63 [00:11<00:00,  5.36it/s, loss=3.37]
Epoch 31:
Training Loss: 4.3423
Time: 11.77 seconds
Learning rate: 0.000966
GPU memory: 0.18 GB
Epoch 32/150: 100%|██████████| 63/63 [00:11<00:00,  5.54it/s, loss=3.65]
Epoch 32:
Training Loss: 3.9969
Time: 11.38 seconds
Learning rate: 0.000961
GPU memory: 0.18 GB
Epoch 33/150: 100%|██████████| 63/63 [00:11<00:00,  5.49it/s, loss=2.54]
Epoch 33:
Training Loss: 4.0310
Time: 11.48 seconds
Learning rate: 0.000957
GPU memory: 0.18 GB
Epoch 34/150: 100%|██████████| 63/63 [00:11<00:00,  5.45it/s, loss=2.03]
Epoch 34:
Training Loss: 3.7281
Time: 11.57 seconds
Learning rate: 0.000952
GPU memory: 0.18 GB
Epoch 35/150: 100%|██████████| 63/63 [00:11<00:00,  5.52it/s, loss=3.04]
```

Epoch 35:
Training Loss: 3.4687
Time: 11.42 seconds
Learning rate: 0.000947
GPU memory: 0.18 GB

Validation: 100% |██████████| 11/11 [00:01<00:00, 8.87it/s, loss=50.7]
Validation Loss: 23.7627

Epoch 36/150: 100% |██████████| 63/63 [00:11<00:00, 5.33it/s, loss=4.91]

Epoch 36:
Training Loss: 3.2785
Time: 11.83 seconds
Learning rate: 0.000941
GPU memory: 0.18 GB

Epoch 37/150: 100% |██████████| 63/63 [00:11<00:00, 5.40it/s, loss=2.53]

Epoch 37:
Training Loss: 3.2670
Time: 11.67 seconds
Learning rate: 0.000936
GPU memory: 0.18 GB

Epoch 38/150: 100% |██████████| 63/63 [00:11<00:00, 5.46it/s, loss=3.53]

Epoch 38:
Training Loss: 3.0444
Time: 11.53 seconds
Learning rate: 0.000930
GPU memory: 0.18 GB

Epoch 39/150: 100% |██████████| 63/63 [00:11<00:00, 5.48it/s, loss=3.05]

Epoch 39:
Training Loss: 2.9842
Time: 11.50 seconds
Learning rate: 0.000924
GPU memory: 0.18 GB

Epoch 40/150: 100% |██████████| 63/63 [00:11<00:00, 5.45it/s, loss=3.11]

Epoch 40:
Training Loss: 2.8105
Time: 11.56 seconds
Learning rate: 0.000918
GPU memory: 0.18 GB

Validation: 100% |██████████| 11/11 [00:01<00:00, 8.83it/s, loss=51.5]
Validation Loss: 24.1102

Epoch 41/150: 100% |██████████| 63/63 [00:11<00:00, 5.40it/s, loss=1.9]

Epoch 41:
Training Loss: 2.5936
Time: 11.67 seconds
Learning rate: 0.000911
GPU memory: 0.18 GB

Epoch 42/150: 100% |██████████| 63/63 [00:11<00:00, 5.55it/s, loss=1.32]

Epoch 42:
Training Loss: 2.5490
Time: 11.35 seconds
Learning rate: 0.000904
GPU memory: 0.18 GB

Epoch 43/150: 100% |██████████| 63/63 [00:11<00:00, 5.50it/s, loss=3.44]

Epoch 43:
Training Loss: 2.3149
Time: 11.47 seconds
Learning rate: 0.000897
GPU memory: 0.18 GB

Epoch 44/150: 100% |██████████| 63/63 [00:11<00:00, 5.51it/s, loss=3.23]

Epoch 44:
Training Loss: 2.2361
Time: 11.44 seconds
Learning rate: 0.000890
GPU memory: 0.18 GB

Epoch 45/150: 100% |██████████| 63/63 [00:11<00:00, 5.51it/s, loss=3.71]

Epoch 45:
Training Loss: 2.2883
Time: 11.44 seconds
Learning rate: 0.000883
GPU memory: 0.18 GB

Validation: 100% |██████████| 11/11 [00:01<00:00, 8.88it/s, loss=46.8]
Validation Loss: 23.9510

Epoch 46/150: 100% |██████████| 63/63 [00:11<00:00, 5.36it/s, loss=1.9]

Epoch 46:
Training Loss: 2.2970
Time: 11.76 seconds
Learning rate: 0.000875
GPU memory: 0.18 GB

Epoch 47/150: 100% |██████████| 63/63 [00:11<00:00, 5.53it/s, loss=1.98]

Epoch 47:
Training Loss: 2.4360
Time: 11.40 seconds
Learning rate: 0.000868
GPU memory: 0.18 GB

Epoch 48/150: 100% |██████████| 63/63 [00:11<00:00, 5.59it/s, loss=2.93]

Epoch 48:
Training Loss: 2.0891
Time: 11.27 seconds
Learning rate: 0.000860
GPU memory: 0.18 GB

Epoch 49/150: 100% |██████████| 63/63 [00:11<00:00, 5.56it/s, loss=3.53]

Epoch 49:
Training Loss: 1.9475
Time: 11.34 seconds
Learning rate: 0.000851
GPU memory: 0.18 GB

Epoch 50/150: 100% |██████████| 63/63 [00:11<00:00, 5.53it/s, loss=3.28]

Epoch 50:
Training Loss: 2.0104
Time: 11.39 seconds
Learning rate: 0.000843
GPU memory: 0.18 GB

Validation: 100% |██████████| 11/11 [00:01<00:00, 8.60it/s, loss=50.5]
Validation Loss: 25.5445

Epoch 51/150: 100% |██████████| 63/63 [00:11<00:00, 5.40it/s, loss=3.54]

Epoch 51:
Training Loss: 1.9200
Time: 11.67 seconds
Learning rate: 0.000834
GPU memory: 0.18 GB

Epoch 52/150: 100% |██████████| 63/63 [00:11<00:00, 5.54it/s, loss=1.28]

Epoch 52:
Training Loss: 1.9154
Time: 11.37 seconds
Learning rate: 0.000826
GPU memory: 0.18 GB

Epoch 53/150: 100% |██████████| 63/63 [00:11<00:00, 5.53it/s, loss=2.2]

Epoch 53:
Training Loss: 1.9514
Time: 11.40 seconds
Learning rate: 0.000817
GPU memory: 0.18 GB

Epoch 54/150: 100% |██████████| 63/63 [00:11<00:00, 5.56it/s, loss=1.6]

Epoch 54:
Training Loss: 1.8161
Time: 11.34 seconds
Learning rate: 0.000808
GPU memory: 0.18 GB

Epoch 55/150: 100% |██████████| 63/63 [00:11<00:00, 5.63it/s, loss=2.5]

Epoch 55:
Training Loss: 1.7980
Time: 11.20 seconds
Learning rate: 0.000798
GPU memory: 0.18 GB

Validation: 100% |██████████| 11/11 [00:01<00:00, 8.94it/s, loss=43.6]

Validation Loss: 24.3326

Epoch 56/150: 100% |██████████| 63/63 [00:11<00:00, 5.41it/s, loss=1.37]

Epoch 56:
Training Loss: 1.6352
Time: 11.64 seconds
Learning rate: 0.000789
GPU memory: 0.18 GB

Epoch 57/150: 100% |██████████| 63/63 [00:11<00:00, 5.55it/s, loss=1.77]

Epoch 57:
Training Loss: 1.6789
Time: 11.36 seconds
Learning rate: 0.000779
GPU memory: 0.18 GB

Epoch 58/150: 100% |██████████| 63/63 [00:11<00:00, 5.55it/s, loss=0.985]

Epoch 58:
Training Loss: 1.6098
Time: 11.36 seconds
Learning rate: 0.000770
GPU memory: 0.18 GB

Epoch 59/150: 100% |██████████| 63/63 [00:11<00:00, 5.51it/s, loss=1.68]

Epoch 59:
Training Loss: 1.7190
Time: 11.43 seconds
Learning rate: 0.000760
GPU memory: 0.18 GB

Epoch 60/150: 100% | ██████████ | 63/63 [00:11<00:00, 5.58it/s, loss=2.07]

Epoch 60:
Training Loss: 1.5997
Time: 11.29 seconds
Learning rate: 0.000750
GPU memory: 0.18 GB

Validation: 100% | ██████████ | 11/11 [00:01<00:00, 8.73it/s, loss=45.1]
Validation Loss: 24.9775

Epoch 61/150: 100% | ██████████ | 63/63 [00:11<00:00, 5.48it/s, loss=1.56]

Epoch 61:
Training Loss: 1.3821
Time: 11.50 seconds
Learning rate: 0.000740
GPU memory: 0.18 GB

Epoch 62/150: 100% | ██████████ | 63/63 [00:11<00:00, 5.56it/s, loss=2.63]

Epoch 62:
Training Loss: 1.3649
Time: 11.34 seconds
Learning rate: 0.000729
GPU memory: 0.18 GB

Epoch 63/150: 100% | ██████████ | 63/63 [00:11<00:00, 5.54it/s, loss=2.19]

Epoch 63:
Training Loss: 1.3657
Time: 11.38 seconds
Learning rate: 0.000719
GPU memory: 0.18 GB

Epoch 64/150: 100% | ██████████ | 63/63 [00:11<00:00, 5.57it/s, loss=1.31]

Epoch 64:
Training Loss: 1.3643
Time: 11.31 seconds
Learning rate: 0.000709
GPU memory: 0.18 GB

Epoch 65/150: 100% | ██████████ | 63/63 [00:11<00:00, 5.56it/s, loss=2.05]

Epoch 65:
Training Loss: 1.4049
Time: 11.33 seconds
Learning rate: 0.000698
GPU memory: 0.18 GB

Validation: 100% | ██████████ | 11/11 [00:01<00:00, 8.61it/s, loss=51.2]
Validation Loss: 26.1285

Epoch 66/150: 100% | ██████████ | 63/63 [00:11<00:00, 5.39it/s, loss=1.07]

Epoch 66:
Training Loss: 1.3373
Time: 11.68 seconds
Learning rate: 0.000687
GPU memory: 0.18 GB

Epoch 67/150: 100% | ██████████ | 63/63 [00:11<00:00, 5.50it/s, loss=1.45]

Epoch 67:
Training Loss: 1.3966
Time: 11.45 seconds
Learning rate: 0.000676
GPU memory: 0.18 GB

Epoch 68/150: 100% |██████████| 63/63 [00:11<00:00, 5.50it/s, loss=1.16]

Epoch 68:
Training Loss: 1.3213
Time: 11.46 seconds
Learning rate: 0.000665
GPU memory: 0.18 GB

Epoch 69/150: 100% |██████████| 63/63 [00:11<00:00, 5.49it/s, loss=3.08]

Epoch 69:
Training Loss: 1.2905
Time: 11.47 seconds
Learning rate: 0.000654
GPU memory: 0.18 GB

Epoch 70/150: 100% |██████████| 63/63 [00:11<00:00, 5.51it/s, loss=1.71]

Epoch 70:
Training Loss: 1.2634
Time: 11.43 seconds
Learning rate: 0.000643
GPU memory: 0.18 GB

Validation: 100% |██████████| 11/11 [00:01<00:00, 8.87it/s, loss=47.2]
Validation Loss: 26.1069

Epoch 71/150: 100% |██████████| 63/63 [00:11<00:00, 5.40it/s, loss=1.52]

Epoch 71:
Training Loss: 1.5086
Time: 11.67 seconds
Learning rate: 0.000632
GPU memory: 0.18 GB

Epoch 72/150: 100% |██████████| 63/63 [00:11<00:00, 5.49it/s, loss=2.32]

Epoch 72:
Training Loss: 1.2905
Time: 11.49 seconds
Learning rate: 0.000621
GPU memory: 0.18 GB

Epoch 73/150: 100% |██████████| 63/63 [00:11<00:00, 5.48it/s, loss=1.21]

Epoch 73:
Training Loss: 1.1094
Time: 11.51 seconds
Learning rate: 0.000609
GPU memory: 0.18 GB

Epoch 74/150: 100% |██████████| 63/63 [00:11<00:00, 5.50it/s, loss=2.9]

Epoch 74:
Training Loss: 1.1012
Time: 11.46 seconds
Learning rate: 0.000598
GPU memory: 0.18 GB

Epoch 75/150: 100% |██████████| 63/63 [00:11<00:00, 5.54it/s, loss=1.46]

Epoch 75:
Training Loss: 1.0147
Time: 11.38 seconds
Learning rate: 0.000587
GPU memory: 0.18 GB

Validation: 100% |██████████| 11/11 [00:01<00:00, 8.94it/s, loss=47.9]
Validation Loss: 26.6928

Epoch 76/150: 100% |██████████| 63/63 [00:11<00:00, 5.45it/s, loss=2.17]

Epoch 76:
Training Loss: 1.1230
Time: 11.56 seconds
Learning rate: 0.000575
GPU memory: 0.18 GB

Epoch 77/150: 100% |██████████| 63/63 [00:11<00:00, 5.53it/s, loss=0.823]

Epoch 77:
Training Loss: 1.0872
Time: 11.39 seconds
Learning rate: 0.000564
GPU memory: 0.18 GB

Epoch 78/150: 100% |██████████| 63/63 [00:11<00:00, 5.51it/s, loss=1.63]

Epoch 78:
Training Loss: 1.0770
Time: 11.44 seconds
Learning rate: 0.000552
GPU memory: 0.18 GB

Epoch 79/150: 100% |██████████| 63/63 [00:11<00:00, 5.57it/s, loss=1.26]

Epoch 79:
Training Loss: 0.9968
Time: 11.32 seconds
Learning rate: 0.000541
GPU memory: 0.18 GB

Epoch 80/150: 100% |██████████| 63/63 [00:11<00:00, 5.50it/s, loss=1.1]

Epoch 80:
Training Loss: 0.9779
Time: 11.45 seconds
Learning rate: 0.000529
GPU memory: 0.18 GB

Validation: 100% |██████████| 11/11 [00:01<00:00, 8.85it/s, loss=49.6]
Validation Loss: 27.6422

Epoch 81/150: 100% |██████████| 63/63 [00:11<00:00, 5.36it/s, loss=0.818]

Epoch 81:
Training Loss: 0.9426
Time: 11.76 seconds
Learning rate: 0.000517
GPU memory: 0.18 GB

Epoch 82/150: 100% |██████████| 63/63 [00:11<00:00, 5.54it/s, loss=1.4]

Epoch 82:
Training Loss: 1.0281
Time: 11.38 seconds
Learning rate: 0.000506
GPU memory: 0.18 GB

Epoch 83/150: 100% |██████████| 63/63 [00:11<00:00, 5.45it/s, loss=1.26]

Epoch 83:
Training Loss: 0.9247
Time: 11.56 seconds
Learning rate: 0.000494
GPU memory: 0.18 GB

Epoch 84/150: 100% | ██████████ | 63/63 [00:11<00:00, 5.46it/s, loss=3.25]

Epoch 84:
Training Loss: 0.8364
Time: 11.54 seconds
Learning rate: 0.000482
GPU memory: 0.18 GB

Epoch 85/150: 100% | ██████████ | 63/63 [00:11<00:00, 5.49it/s, loss=1.19]

Epoch 85:
Training Loss: 0.8511
Time: 11.48 seconds
Learning rate: 0.000471
GPU memory: 0.18 GB

Validation: 100% | ██████████ | 11/11 [00:01<00:00, 8.97it/s, loss=51.1]
Validation Loss: 27.8088

Epoch 86/150: 100% | ██████████ | 63/63 [00:11<00:00, 5.49it/s, loss=1.2]

Epoch 86:
Training Loss: 0.8524
Time: 11.48 seconds
Learning rate: 0.000459
GPU memory: 0.18 GB

Epoch 87/150: 100% | ██████████ | 63/63 [00:11<00:00, 5.28it/s, loss=1.27]

Epoch 87:
Training Loss: 0.8172
Time: 11.93 seconds
Learning rate: 0.000448
GPU memory: 0.18 GB

Epoch 88/150: 100% | ██████████ | 63/63 [00:11<00:00, 5.55it/s, loss=1.14]

Epoch 88:
Training Loss: 0.7862
Time: 11.36 seconds
Learning rate: 0.000436
GPU memory: 0.18 GB

Epoch 89/150: 100% | ██████████ | 63/63 [00:11<00:00, 5.52it/s, loss=1.11]

Epoch 89:
Training Loss: 0.7717
Time: 11.42 seconds
Learning rate: 0.000424
GPU memory: 0.18 GB

Epoch 90/150: 100% | ██████████ | 63/63 [00:11<00:00, 5.49it/s, loss=0.624]

Epoch 90:
Training Loss: 0.7276
Time: 11.48 seconds
Learning rate: 0.000413
GPU memory: 0.18 GB

Validation: 100% | ██████████ | 11/11 [00:01<00:00, 8.75it/s, loss=52.3]
Validation Loss: 28.3646

Epoch 91/150: 100% | ██████████ | 63/63 [00:11<00:00, 5.46it/s, loss=1.24]

Epoch 91:
Training Loss: 0.6661
Time: 11.55 seconds
Learning rate: 0.000402
GPU memory: 0.18 GB

Epoch 92/150: 100% |██████████| 63/63 [00:11<00:00, 5.34it/s, loss=0.724]

Epoch 92:
Training Loss: 0.6860
Time: 11.81 seconds
Learning rate: 0.000390
GPU memory: 0.18 GB

Epoch 93/150: 100% |██████████| 63/63 [00:11<00:00, 5.46it/s, loss=1.26]

Epoch 93:
Training Loss: 0.7679
Time: 11.54 seconds
Learning rate: 0.000379
GPU memory: 0.18 GB

Epoch 94/150: 100% |██████████| 63/63 [00:11<00:00, 5.48it/s, loss=1.84]

Epoch 94:
Training Loss: 0.6552
Time: 11.49 seconds
Learning rate: 0.000368
GPU memory: 0.18 GB

Epoch 95/150: 100% |██████████| 63/63 [00:11<00:00, 5.60it/s, loss=1.11]

Epoch 95:
Training Loss: 0.6825
Time: 11.26 seconds
Learning rate: 0.000356
GPU memory: 0.18 GB

Validation: 100% |██████████| 11/11 [00:01<00:00, 9.13it/s, loss=51.1]

Validation Loss: 29.0464

Epoch 96/150: 100% |██████████| 63/63 [00:11<00:00, 5.38it/s, loss=1.31]

Epoch 96:
Training Loss: 0.6499
Time: 11.70 seconds
Learning rate: 0.000345
GPU memory: 0.18 GB

Epoch 97/150: 100% |██████████| 63/63 [00:12<00:00, 5.23it/s, loss=0.692]

Epoch 97:
Training Loss: 0.6615
Time: 12.05 seconds
Learning rate: 0.000334
GPU memory: 0.18 GB

Epoch 98/150: 100% |██████████| 63/63 [00:11<00:00, 5.40it/s, loss=0.284]

Epoch 98:
Training Loss: 0.5823
Time: 11.68 seconds
Learning rate: 0.000323
GPU memory: 0.18 GB

Epoch 99/150: 100% |██████████| 63/63 [00:11<00:00, 5.42it/s, loss=0.534]

Epoch 99:
Training Loss: 0.5703
Time: 11.62 seconds
Learning rate: 0.000313
GPU memory: 0.18 GB

Epoch 100/150: 100%|██████████| 63/63 [00:11<00:00, 5.37it/s, loss=3.76]

Epoch 100:
Training Loss: 0.5985
Time: 11.73 seconds
Learning rate: 0.000302
GPU memory: 0.18 GB

Validation: 100%|██████████| 11/11 [00:01<00:00, 8.62it/s, loss=50.6]

Validation Loss: 28.5191

Epoch 101/150: 100%|██████████| 63/63 [00:11<00:00, 5.42it/s, loss=0.655]

Epoch 101:
Training Loss: 0.5435
Time: 11.64 seconds
Learning rate: 0.000291
GPU memory: 0.18 GB

Epoch 102/150: 100%|██████████| 63/63 [00:11<00:00, 5.33it/s, loss=0.604]

Epoch 102:
Training Loss: 0.5460
Time: 11.82 seconds
Learning rate: 0.000281
GPU memory: 0.18 GB

Epoch 103/150: 100%|██████████| 63/63 [00:11<00:00, 5.40it/s, loss=0.598]

Epoch 103:
Training Loss: 0.5192
Time: 11.68 seconds
Learning rate: 0.000270
GPU memory: 0.18 GB

Epoch 104/150: 100%|██████████| 63/63 [00:11<00:00, 5.38it/s, loss=0.733]

Epoch 104:
Training Loss: 0.5298
Time: 11.71 seconds
Learning rate: 0.000260
GPU memory: 0.18 GB

Epoch 105/150: 100%|██████████| 63/63 [00:11<00:00, 5.46it/s, loss=0.53]

Epoch 105:
Training Loss: 0.5211
Time: 11.55 seconds
Learning rate: 0.000250
GPU memory: 0.18 GB

Validation: 100%|██████████| 11/11 [00:01<00:00, 8.74it/s, loss=54.3]

Validation Loss: 29.1353

Epoch 106/150: 100%|██████████| 63/63 [00:11<00:00, 5.47it/s, loss=0.656]

Epoch 106:
Training Loss: 0.4919
Time: 11.51 seconds
Learning rate: 0.000240
GPU memory: 0.18 GB

Epoch 107/150: 100%|██████████| 63/63 [00:11<00:00, 5.40it/s, loss=0.617]

Epoch 107:
Training Loss: 0.4691
Time: 11.68 seconds
Learning rate: 0.000230
GPU memory: 0.18 GB

Epoch 108/150: 100%|██████████| 63/63 [00:11<00:00, 5.53it/s, loss=0.262]

Epoch 108:
Training Loss: 0.5118
Time: 11.39 seconds
Learning rate: 0.000220
GPU memory: 0.18 GB

Epoch 109/150: 100%|██████████| 63/63 [00:11<00:00, 5.52it/s, loss=1.51]

Epoch 109:
Training Loss: 0.4761
Time: 11.42 seconds
Learning rate: 0.000211
GPU memory: 0.18 GB

Epoch 110/150: 100%|██████████| 63/63 [00:11<00:00, 5.47it/s, loss=0.371]

Epoch 110:
Training Loss: 0.4500
Time: 11.52 seconds
Learning rate: 0.000201
GPU memory: 0.18 GB

Validation: 100%|██████████| 11/11 [00:01<00:00, 8.49it/s, loss=52]

Validation Loss: 29.0900

Epoch 111/150: 100%|██████████| 63/63 [00:11<00:00, 5.51it/s, loss=0.795]

Epoch 111:
Training Loss: 0.4368
Time: 11.43 seconds
Learning rate: 0.000192
GPU memory: 0.18 GB

Epoch 112/150: 100%|██████████| 63/63 [00:11<00:00, 5.36it/s, loss=0.244]

Epoch 112:
Training Loss: 0.4153
Time: 11.75 seconds
Learning rate: 0.000183
GPU memory: 0.18 GB

Epoch 113/150: 100%|██████████| 63/63 [00:11<00:00, 5.46it/s, loss=0.144]

Epoch 113:
Training Loss: 0.3904
Time: 11.53 seconds
Learning rate: 0.000174
GPU memory: 0.18 GB

Epoch 114/150: 100%|██████████| 63/63 [00:11<00:00, 5.48it/s, loss=1.24]

Epoch 114:
Training Loss: 0.4111
Time: 11.51 seconds
Learning rate: 0.000165
GPU memory: 0.18 GB

Epoch 115/150: 100%|██████████| 63/63 [00:11<00:00, 5.45it/s, loss=0.611]

Epoch 115:
Training Loss: 0.4051
Time: 11.56 seconds
Learning rate: 0.000157
GPU memory: 0.18 GB

Validation: 100% |██████████| 11/11 [00:01<00:00, 8.76it/s, loss=53.2]

Validation Loss: 29.1556

Epoch 116/150: 100% |██████████| 63/63 [00:11<00:00, 5.48it/s, loss=0.278]

Epoch 116:
Training Loss: 0.3635
Time: 11.50 seconds
Learning rate: 0.000148
GPU memory: 0.18 GB

Epoch 117/150: 100% |██████████| 63/63 [00:11<00:00, 5.39it/s, loss=0.498]

Epoch 117:
Training Loss: 0.3640
Time: 11.69 seconds
Learning rate: 0.000140
GPU memory: 0.18 GB

Epoch 118/150: 100% |██████████| 63/63 [00:11<00:00, 5.44it/s, loss=0.427]

Epoch 118:
Training Loss: 0.3701
Time: 11.58 seconds
Learning rate: 0.000132
GPU memory: 0.18 GB

Epoch 119/150: 100% |██████████| 63/63 [00:11<00:00, 5.47it/s, loss=0.984]

Epoch 119:
Training Loss: 0.3596
Time: 11.52 seconds
Learning rate: 0.000124
GPU memory: 0.18 GB

Epoch 120/150: 100% |██████████| 63/63 [00:11<00:00, 5.47it/s, loss=1.13]

Epoch 120:
Training Loss: 0.3901
Time: 11.52 seconds
Learning rate: 0.000117
GPU memory: 0.18 GB

Validation: 100% |██████████| 11/11 [00:01<00:00, 8.97it/s, loss=53.7]

Validation Loss: 29.7444

Epoch 121/150: 100% |██████████| 63/63 [00:11<00:00, 5.49it/s, loss=0.274]

Epoch 121:
Training Loss: 0.3822
Time: 11.47 seconds
Learning rate: 0.000110
GPU memory: 0.18 GB

Epoch 122/150: 100% |██████████| 63/63 [00:11<00:00, 5.42it/s, loss=0.243]

Epoch 122:
Training Loss: 0.3407
Time: 11.63 seconds
Learning rate: 0.000102
GPU memory: 0.18 GB

Epoch 123/150: 100% |██████████| 63/63 [00:11<00:00, 5.48it/s, loss=1.05]

Epoch 123:
Training Loss: 0.3268
Time: 11.49 seconds
Learning rate: 0.000095
GPU memory: 0.18 GB

Epoch 124/150: 100%|██████████| 63/63 [00:11<00:00, 5.46it/s, loss=6.08]

Epoch 124:
Training Loss: 0.3992
Time: 11.55 seconds
Learning rate: 0.000089
GPU memory: 0.18 GB

Epoch 125/150: 100%|██████████| 63/63 [00:11<00:00, 5.53it/s, loss=0.554]

Epoch 125:
Training Loss: 0.3352
Time: 11.39 seconds
Learning rate: 0.000082
GPU memory: 0.18 GB

Validation: 100%|██████████| 11/11 [00:01<00:00, 9.14it/s, loss=53.8]
Validation Loss: 29.7775

Epoch 126/150: 100%|██████████| 63/63 [00:11<00:00, 5.60it/s, loss=0.23]

Epoch 126:
Training Loss: 0.3227
Time: 11.25 seconds
Learning rate: 0.000076
GPU memory: 0.18 GB

Epoch 127/150: 100%|██████████| 63/63 [00:11<00:00, 5.36it/s, loss=0.35]

Epoch 127:
Training Loss: 0.3246
Time: 11.77 seconds
Learning rate: 0.000070
GPU memory: 0.18 GB

Epoch 128/150: 100%|██████████| 63/63 [00:11<00:00, 5.51it/s, loss=1.09]

Epoch 128:
Training Loss: 0.3330
Time: 11.43 seconds
Learning rate: 0.000064
GPU memory: 0.18 GB

Epoch 129/150: 100%|██████████| 63/63 [00:11<00:00, 5.54it/s, loss=0.616]

Epoch 129:
Training Loss: 0.3222
Time: 11.38 seconds
Learning rate: 0.000058
GPU memory: 0.18 GB

Epoch 130/150: 100%|██████████| 63/63 [00:11<00:00, 5.49it/s, loss=0.228]

Epoch 130:
Training Loss: 0.2833
Time: 11.48 seconds
Learning rate: 0.000053
GPU memory: 0.18 GB

Validation: 100%|██████████| 11/11 [00:01<00:00, 8.76it/s, loss=52.5]
Validation Loss: 29.2709

Epoch 131/150: 100%|██████████| 63/63 [00:11<00:00, 5.49it/s, loss=0.806]

Epoch 131:
Training Loss: 0.2944
Time: 11.48 seconds
Learning rate: 0.000048
GPU memory: 0.18 GB

Epoch 132/150: 100%|██████████| 63/63 [00:11<00:00, 5.50it/s, loss=0.609]

Epoch 132:
Training Loss: 0.2861
Time: 11.45 seconds
Learning rate: 0.000048
GPU memory: 0.18 GB

Epoch 133/150: 100%|██████████| 63/63 [00:11<00:00, 5.38it/s, loss=0.542]

Epoch 133:
Training Loss: 0.2810
Time: 11.71 seconds
Learning rate: 0.000039
GPU memory: 0.18 GB

Epoch 134/150: 100%|██████████| 63/63 [00:11<00:00, 5.46it/s, loss=0.213]

Epoch 134:
Training Loss: 0.2700
Time: 11.54 seconds
Learning rate: 0.000034
GPU memory: 0.18 GB

Epoch 135/150: 100%|██████████| 63/63 [00:11<00:00, 5.53it/s, loss=0.291]

Epoch 135:
Training Loss: 0.2717
Time: 11.41 seconds
Learning rate: 0.000030
GPU memory: 0.18 GB

Validation: 100%|██████████| 11/11 [00:01<00:00, 8.96it/s, loss=52.8]

Validation Loss: 29.4062

Epoch 136/150: 100%|██████████| 63/63 [00:11<00:00, 5.53it/s, loss=0.954]

Epoch 136:
Training Loss: 0.2788
Time: 11.40 seconds
Learning rate: 0.000026
GPU memory: 0.18 GB

Epoch 137/150: 100%|██████████| 63/63 [00:11<00:00, 5.46it/s, loss=0.462]

Epoch 137:
Training Loss: 0.2759
Time: 11.54 seconds
Learning rate: 0.000023
GPU memory: 0.18 GB

Epoch 138/150: 100%|██████████| 63/63 [00:11<00:00, 5.36it/s, loss=0.504]

Epoch 138:
Training Loss: 0.2593
Time: 11.75 seconds
Learning rate: 0.000019
GPU memory: 0.18 GB

Epoch 139/150: 100%|██████████| 63/63 [00:11<00:00, 5.53it/s, loss=0.209]

Epoch 139:
Training Loss: 0.2602
Time: 11.40 seconds
Learning rate: 0.000016
GPU memory: 0.18 GB

Epoch 140/150: 100%|██████████| 63/63 [00:11<00:00, 5.48it/s, loss=0.378]

Epoch 140:
Training Loss: 0.2520
Time: 11.51 seconds
Learning rate: 0.000013
GPU memory: 0.18 GB

Validation: 100%|██████████| 11/11 [00:01<00:00, 8.44it/s, loss=53]

Validation Loss: 29.4707

Epoch 141/150: 100%|██████████| 63/63 [00:11<00:00, 5.52it/s, loss=0.566]

Epoch 141:
Training Loss: 0.2389
Time: 11.41 seconds
Learning rate: 0.000011
GPU memory: 0.18 GB

Epoch 142/150: 100%|██████████| 63/63 [00:11<00:00, 5.60it/s, loss=0.261]

Epoch 142:
Training Loss: 0.2689
Time: 11.25 seconds
Learning rate: 0.000009
GPU memory: 0.18 GB

Epoch 143/150: 100%|██████████| 63/63 [00:11<00:00, 5.41it/s, loss=0.678]

Epoch 143:
Training Loss: 0.2449
Time: 11.64 seconds
Learning rate: 0.000007
GPU memory: 0.18 GB

Epoch 144/150: 100%|██████████| 63/63 [00:11<00:00, 5.50it/s, loss=1.26]

Epoch 144:
Training Loss: 0.2783
Time: 11.47 seconds
Learning rate: 0.000005
GPU memory: 0.18 GB

Epoch 145/150: 100%|██████████| 63/63 [00:11<00:00, 5.51it/s, loss=0.216]

Epoch 145:
Training Loss: 0.2671
Time: 11.43 seconds
Learning rate: 0.000003
GPU memory: 0.18 GB

Validation: 100%|██████████| 11/11 [00:01<00:00, 8.92it/s, loss=52.8]

Validation Loss: 29.4028

Epoch 146/150: 100%|██████████| 63/63 [00:11<00:00, 5.49it/s, loss=0.89]

Epoch 146:
Training Loss: 0.2604
Time: 11.47 seconds
Learning rate: 0.000002
GPU memory: 0.18 GB

Epoch 147/150: 100%|██████████| 63/63 [00:11<00:00, 5.53it/s, loss=0.372]

Epoch 147:
Training Loss: 0.2556
Time: 11.40 seconds
Learning rate: 0.000001
GPU memory: 0.18 GB

Epoch 148/150: 100%|██████████| 63/63 [00:11<00:00, 5.54it/s, loss=0.216]

Epoch 148:
Training Loss: 0.2462
Time: 11.38 seconds
Learning rate: 0.000001
GPU memory: 0.18 GB

Epoch 149/150: 100%|██████████| 63/63 [00:11<00:00, 5.42it/s, loss=1.17]

Epoch 149:
Training Loss: 0.2658
Time: 11.63 seconds
Learning rate: 0.000000
GPU memory: 0.18 GB

Epoch 150/150: 100%|██████████| 63/63 [00:11<00:00, 5.52it/s, loss=0.389]

Epoch 150:
Training Loss: 0.2588
Time: 11.41 seconds
Learning rate: 0.000000
GPU memory: 0.18 GB

Validation: 100%|██████████| 11/11 [00:01<00:00, 8.88it/s, loss=52.8]

```
Validation Loss: 29.4055
```

```
Training completed!
Average epoch time: 11.59 seconds
Total training time: 1737.89 seconds
```

```
Exporting to ONNX format...
```

```
Verifying ONNX model...
Creating datasets...
Initializing train dataset...
Loading class descriptions from /scratch/poh2005/data/open_images_temp/class-descriptions-boxable.csv
Loading annotations from /scratch/poh2005/data/open_images_temp/sub-train-annotations-bbox.csv
Found 2000 images with 4460 annotations
Dataset initialization completed for train split
Initializing validation dataset...
Loading class descriptions from /scratch/poh2005/data/open_images_temp/class-descriptions-boxable.csv
Loading annotations from /scratch/poh2005/data/open_images_temp/sub-validation-annotations-bbox.csv
Found 327 images with 444 annotations
Dataset initialization completed for validation split
Initializing test dataset...
Loading class descriptions from /scratch/poh2005/data/open_images_temp/class-descriptions-boxable.csv
Loading annotations from /scratch/poh2005/data/open_images_temp/sub-test-annotations-bbox.csv
Found 800 images with 1074 annotations
Dataset initialization completed for test split
Created dataloaders with batch size 32 and 4 workers
```

```
Generating visualizations...
```

```
/state/partition1/job-53232312/ipykernel_1360913/2568295003.py:440: FutureWarning: `torch.cuda.amp.autocast(args...)` is deprecated. Please use `torch.amp.autocast('cuda', args...)` instead.
```

```
    with torch.no_grad(), autocast():
```



```
/state/partition1/job-53232312/ipykernel_1360913/2568295003.py:519: FutureWarning: `torch.cuda.amp.autocast(args...)` is deprecated. Please use `torch.amp.autocast('cuda', args...)` instead.
```

```
    with torch.no_grad(), autocast():
```

```
Evaluating model...
```

```
Evaluating model performance...
```

```
Evaluation: 100%|██████████| 25/25 [00:15<00:00, 1.60it/s]
```

```
Evaluation Results:  
Overall Accuracy: 0.3091  
  
Per-class Accuracy:  
Airplane: 0.3146  
Boat: 0.3182  
  
Training completed successfully!
```

2. Export the Pytorch model to ONNX

Export the Pytorch model to ONNX using `torch.onnx.export()` function and save it. When you export the model, the function will execute the model, recording a trace of what operators are used to compute the outputs. Because `export` runs the model, we need to provide an input tensor `x`. The values in this can be random as long as it is the right type and size. Use a dummy random tensor.

(3 points)

```
In [46]: import torch  
from torchvision.models.detection import ssdlite320_mobilenet_v3_large  
  
# Load pretrained model  
model = ssdlite320_mobilenet_v3_large(pretrained=True)  
model.eval()  
  
# Create dummy input tensor (320x320 as per model requirements)  
dummy_input = torch.randn(1, 3, 320, 320)  
  
# Export to ONNX  
torch.onnx.export(  
    model,  
    dummy_input,  
    'ssd_model.onnx',  
    export_params=True,  
    opset_version=11,  
    do_constant_folding=True,  
    input_names=['input'],  
    output_names=['boxes', 'labels', 'scores'],  
    dynamic_axes={  
        'input': {0: 'batch_size'},  
        'boxes': {0: 'batch_size'},  
        'labels': {0: 'batch_size'},  
        'scores': {0: 'batch_size'}  
    }  
)  
print("Model exported to ONNX successfully!")
```

Model exported to ONNX successfully!

3. Load and verify the model

Load the saved model using `onnx.load` and verify the model's structure using `onnx.checker.check_model`.

(3 points)

```
In [47]: import onnx

# Load the ONNX model
onnx_model = onnx.load('ssd_model.onnx')

# Check model structure
onnx.checker.check_model(onnx_model)
print("ONNX model structure verified successfully!")
```

ONNX model structure verified successfully!

4. Run the model with ONNX Runtime (ORT)

Next, run the model with ONNX Runtime (ORT). You first need to create an inference session for the model and then evaluate the model using the `run()` API.

(3 points)

```
In [48]: import onnxruntime as ort
import numpy as np

# Create inference session
session = ort.InferenceSession('ssd_model.onnx')

# Use same dummy input as export
dummy_input = torch.randn(1, 3, 320, 320)

# Run inference
ort_outputs = session.run(None, {session.get_inputs()[0].name: dummy_input})
print("ONNX Runtime inference completed successfully!")
```

ONNX Runtime inference completed successfully!

2024-11-09 15:29:56.946169237 [E:onnxruntime:Default, env.cc:234 ThreadMain] pthread_setaffinity_np failed for thread: 1492393, index: 13, mask: {14, }, error code: 22 error msg: Invalid argument. Specify the number of threads explicitly so the affinity is not set.

2024-11-09 15:29:56.951289952 [E:onnxruntime:Default, env.cc:234 ThreadMain] pthread_setaffinity_np failed for thread: 1492398, index: 18, mask: {19, }, error code: 22 error msg: Invalid argument. Specify the number of threads explicitly so the affinity is not set.

2024-11-09 15:29:56.952726104 [E:onnxruntime:Default, env.cc:234 ThreadMain] pthread_setaffinity_np failed for thread: 1492396, index: 16, mask: {17, }, error code: 22 error msg: Invalid argument. Specify the number of threads explicitly so the affinity is not set.

2024-11-09 15:29:56.952749353 [E:onnxruntime:Default, env.cc:234 ThreadMain] pthread_setaffinity_np failed for thread: 1492413, index: 33, mask: {34, }, error code: 22 error msg: Invalid argument. Specify the number of threads explicitly so the affinity is not set.

2024-11-09 15:29:56.954344254 [E:onnxruntime:Default, env.cc:234 ThreadMain] pthread_setaffinity_np failed for thread: 1492403, index: 23, mask: {24, }, error code: 22 error msg: Invalid argument. Specify the number of threads explicitly so the affinity is not set.

2024-11-09 15:29:56.954527302 [E:onnxruntime:Default, env.cc:234 ThreadMain] pthread_setaffinity_np failed for thread: 1492405, index: 25, mask: {26, }, error code: 22 error msg: Invalid argument. Specify the number of threads explicitly so the affinity is not set.

2024-11-09 15:29:56.954722315 [E:onnxruntime:Default, env.cc:234 ThreadMain] pthread_setaffinity_np failed for thread: 1492401, index: 21, mask: {22, }, error code: 22 error msg: Invalid argument. Specify the number of threads explicitly so the affinity is not set.

2024-11-09 15:29:56.954728532 [E:onnxruntime:Default, env.cc:234 ThreadMain] pthread_setaffinity_np failed for thread: 1492395, index: 15, mask: {16, }, error code: 22 error msg: Invalid argument. Specify the number of threads explicitly so the affinity is not set.

2024-11-09 15:29:56.954724405 [E:onnxruntime:Default, env.cc:234 ThreadMain] pthread_setaffinity_np failed for thread: 1492406, index: 26, mask: {27, }, error code: 22 error msg: Invalid argument. Specify the number of threads explicitly so the affinity is not set.

2024-11-09 15:29:56.954744217 [E:onnxruntime:Default, env.cc:234 ThreadMain] pthread_setaffinity_np failed for thread: 1492418, index: 38, mask: {39, }, error code: 22 error msg: Invalid argument. Specify the number of threads explicitly so the affinity is not set.

2024-11-09 15:29:56.954720883 [E:onnxruntime:Default, env.cc:234 ThreadMain] pthread_setaffinity_np failed for thread: 1492402, index: 22, mask: {23, }, error code: 22 error msg: Invalid argument. Specify the number of threads explicitly so the affinity is not set.

2024-11-09 15:29:56.954723216 [E:onnxruntime:Default, env.cc:234 ThreadMain] pthread_setaffinity_np failed for thread: 1492404, index: 24, mask: {25, }, error code: 22 error msg: Invalid argument. Specify the number of threads explicitly so the affinity is not set.

2024-11-09 15:29:56.954757354 [E:onnxruntime:Default, env.cc:234 ThreadMain] pthread_setaffinity_np failed for thread: 1492420, index: 40, mask: {41, }, error code: 22 error msg: Invalid argument. Specify the number of threads explicitly so the affinity is not set.

2024-11-09 15:29:56.954764958 [E:onnxruntime:Default, env.cc:234 ThreadMain] pthread_setaffinity_np failed for thread: 1492417, index: 37, mask: {38, }, error code: 22 error msg: Invalid argument. Specify the number of threads explicitly so the affinity is not set.

2024-11-09 15:29:56.955249177 [E:onnxruntime:Default, env.cc:234 ThreadMain] pthread_setaffinity_np failed for thread: 1492397, index: 17, mask: {18, }, error code: 22 error msg: Invalid argument. Specify the number of threads explicitly so the affinity is not set.

2024-11-09 15:29:56.955716080 [E:onnxruntime:Default, env.cc:234 ThreadMain] pthread_setaffinity_np failed for thread: 1492394, index: 14, mask: {15, }, error code: 22 error msg: Invalid argument. Specify the number of threads explicitly so the affinity is not set.

2024-11-09 15:29:56.955720845 [E:onnxruntime:Default, env.cc:234 ThreadMain] pthread_setaffinity_np failed for thread: 1492400, index: 20, mask: {21, }, error code: 22 error msg: Invalid argument. Specify the number of threads explicitly so the affinity is not set.

2024-11-09 15:29:56.952726499 [E:onnxruntime:Default, env.cc:234 ThreadMain] pthread_setaffinity_np failed for thread: 1492399, index: 19, mask: {20, }, error code: 22 error msg: Invalid argument. Specify the number of threads explicitly so the affinity is not set.

2024-11-09 15:29:56.958871479 [E:onnxruntime:Default, env.cc:234 ThreadMain] pthread_setaffinity_np failed for thread: 1492407, index: 27, mask: {28, }, error code: 22 error msg: Invalid argument. Specify the number of threads explicitly so the affinity is not set.

2024-11-09 15:29:56.959385684 [E:onnxruntime:Default, env.cc:234 ThreadMain] pthread_setaffinity_np failed for thread: 1492416, index: 36, mask: {37, }, error code: 22 error msg: Invalid argument. Specify the number of threads explicitly so the affinity is not set.

2024-11-09 15:29:56.959570330 [E:onnxruntime:Default, env.cc:234 ThreadMain] pthread_setaffinity_np failed for thread: 1492419, index: 39, mask: {40, }, error code: 22 error msg: Invalid argument. Specify the number of threads explicitly so the affinity is not set.

2024-11-09 15:29:56.960298661 [E:onnxruntime:Default, env.cc:234 ThreadMain] pthread_setaffinity_np failed for thread: 1492411, index: 31, mask: {32, }, error code: 22 error msg: Invalid argument. Specify the number of threads explicitly so the affinity is not set.

2024-11-09 15:29:56.961384151 [E:onnxruntime:Default, env.cc:234 ThreadMain] pthread_setaffinity_np failed for thread: 1492412, index: 32, mask: {33, }, error code: 22 error msg: Invalid argument. Specify the number of threads explicitly so the affinity is not set.

2024-11-09 15:29:56.961708582 [E:onnxruntime:Default, env.cc:234 ThreadMain] pthread_setaffinity_np failed for thread: 1492426, index: 46, mask: {47, }, error code: 22 error msg: Invalid argument. Specify the number of threads explicitly so the affinity is not set.

2024-11-09 15:29:56.962708868 [E:onnxruntime:Default, env.cc:234 ThreadMain] pthread_setaffinity_np failed for thread: 1492410, index: 30, mask: {31, }, error code: 22 error msg: Invalid argument. Specify the number of threads explicitly so the affinity is not set.

2024-11-09 15:29:56.963851017 [E:onnxruntime:Default, env.cc:234 ThreadMain] pthread_setaffinity_np failed for thread: 1492408, index: 28, mask: {29, }, error code: 22 error msg: Invalid argument. Specify the number of threads explicitly so the affinity is not set.

2024-11-09 15:29:56.964443403 [E:onnxruntime:Default, env.cc:234 ThreadMain] pthread_setaffinity_np failed for thread: 1492422, index: 42, mask: {43, }, error code: 22 error msg: Invalid argument. Specify the number of threads explicitly so the affinity is not set.

2024-11-09 15:29:56.964709435 [E:onnxruntime:Default, env.cc:234 ThreadMain] pthread_setaffinity_np failed for thread: 1492414, index: 34, mask: {35, }, error code: 22 error msg: Invalid argument. Specify the number of threads explicitly so the affinity is not set.

```
2024-11-09 15:29:56.965356511 [E:onnxruntime:Default, env.cc:234 ThreadMain]
pthread_setaffinity_np failed for thread: 1492423, index: 43, mask: {44, },
error code: 22 error msg: Invalid argument. Specify the number of threads ex
plicitly so the affinity is not set.
2024-11-09 15:29:56.954742744 [E:onnxruntime:Default, env.cc:234 ThreadMain]
pthread_setaffinity_np failed for thread: 1492415, index: 35, mask: {36, },
error code: 22 error msg: Invalid argument. Specify the number of threads ex
plicitly so the affinity is not set.
2024-11-09 15:29:56.966446249 [E:onnxruntime:Default, env.cc:234 ThreadMain]
pthread_setaffinity_np failed for thread: 1492424, index: 44, mask: {45, },
error code: 22 error msg: Invalid argument. Specify the number of threads ex
plicitly so the affinity is not set.
2024-11-09 15:29:56.967708817 [E:onnxruntime:Default, env.cc:234 ThreadMain]
pthread_setaffinity_np failed for thread: 1492421, index: 41, mask: {42, },
error code: 22 error msg: Invalid argument. Specify the number of threads ex
plicitly so the affinity is not set.
2024-11-09 15:29:56.968710152 [E:onnxruntime:Default, env.cc:234 ThreadMain]
pthread_setaffinity_np failed for thread: 1492409, index: 29, mask: {30, },
error code: 22 error msg: Invalid argument. Specify the number of threads ex
plicitly so the affinity is not set.
2024-11-09 15:29:56.971508195 [E:onnxruntime:Default, env.cc:234 ThreadMain]
pthread_setaffinity_np failed for thread: 1492425, index: 45, mask: {46, },
error code: 22 error msg: Invalid argument. Specify the number of threads ex
plicitly so the affinity is not set.
```

5. Compare PyTorch and ONNX Runtime outputs

Does the output of PyTorch (from `torch.out`) and ONNX Runtime match? What precision did you use to match?

(2 points)

```
In [61]: def test_model_outputs():
    """Test and compare outputs between PyTorch and ONNX models"""
    # Load and prepare models
    device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
    print(f'Using device: {device}')

    # Load PyTorch model
    print("Loading pretrained SSDLite320 MobileNetV3 model...")
    model = ssdlite320_mobilenet_v3_large(pretrained=True)
    model.to(device)
    model.eval()

    # Load ONNX model
    print("\nLoading ONNX model...")
    providers = ['CUDAExecutionProvider', 'CPUExecutionProvider'] if torch.c
    ort_session = ort.InferenceSession("ssd_model.onnx", providers=providers

    # Create dummy input for testing
    dummy_input = torch.randn(1, 3, 320, 320, device=device)

    # Get PyTorch outputs
```

```

print("\nGetting PyTorch outputs...")
with torch.no_grad():
    torch_outputs = model(dummy_input)

# Get ONNX outputs
print("\nGetting ONNX Runtime outputs...")
ort_outputs = ort_session.run(None, {ort_session.get_inputs()[0].name: c})

# Print shape information
print("\nOutput shapes comparison:")
print("PyTorch output:")
print(f"- Boxes shape: {torch_outputs[0]['boxes'].shape}")
print(f"- Scores shape: {torch_outputs[0]['scores'].shape}")
print(f"- Labels shape: {torch_outputs[0]['labels'].shape}")

print("\nONNX output:")
print(f"- Boxes shape: {ort_outputs[0].shape}")
print(f"- Scores shape: {ort_outputs[1].shape}")
print(f"- Labels shape: {ort_outputs[2].shape}")

# Convert PyTorch outputs to numpy and get high confidence predictions
torch_boxes = torch_outputs[0]['boxes'].cpu().numpy()
torch_scores = torch_outputs[0]['scores'].cpu().numpy()
torch_labels = torch_outputs[0]['labels'].cpu().numpy()

# Filter PyTorch predictions by confidence threshold
confidence_threshold = 0.3
torch_mask = torch_scores > confidence_threshold
torch_boxes_filtered = torch_boxes[torch_mask]
torch_scores_filtered = torch_scores[torch_mask]
torch_labels_filtered = torch_labels[torch_mask]

print(f"\nNumber of detections above {confidence_threshold} confidence threshold")
print(f"PyTorch: {len(torch_scores_filtered)}")
print(f"ONNX: {len(ort_outputs[1])}")

# Sort both outputs by confidence score for comparison
torch_indices = np.argsort(-torch_scores_filtered) # Descending order
onnx_indices = np.argsort(-ort_outputs[1]) # Descending order

# Take top N predictions from both
N = min(len(torch_indices), len(onnx_indices))
if N > 0:
    print(f"\nComparing top {N} predictions:")

# Get top N predictions from both models
torch_top_boxes = torch_boxes_filtered[torch_indices[:N]]
torch_top_scores = torch_scores_filtered[torch_indices[:N]]
torch_top_labels = torch_labels_filtered[torch_indices[:N]]

onnx_top_boxes = ort_outputs[0][onnx_indices[:N]]
onnx_top_scores = ort_outputs[1][onnx_indices[:N]]
onnx_top_labels = ort_outputs[2][onnx_indices[:N]]

# Calculate differences for available predictions
print("\nMax absolute differences for top predictions:")

```

```

print(f"Boxes: {np.max(np.abs(torch_top_boxes - onnx_top_boxes))}")
print(f"Scores: {np.max(np.abs(torch_top_scores - onnx_top_scores))}")
print(f"Labels: {np.max(np.abs(torch_top_labels - onnx_top_labels))}")

# Test different precision levels
precisions = [1e-3, 1e-4, 1e-5, 1e-6]
for precision in precisions:
    print(f"\nOutput comparison with precision {precision}:")
    boxes_match = np.allclose(torch_top_boxes, onnx_top_boxes, rtol=precision)
    scores_match = np.allclose(torch_top_scores, onnx_top_scores, rtol=precision)
    labels_match = np.allclose(torch_top_labels, onnx_top_labels, rtol=precision)

    print(f"Boxes match: {boxes_match}")
    print(f"Scores match: {scores_match}")
    print(f"Labels match: {labels_match}")

# Print detailed comparison of first prediction
print("\nDetailed comparison of first prediction:")
print("PyTorch:")
print(f"- Box: {torch_top_boxes[0]}")
print(f"- Score: {torch_top_scores[0]}")
print(f"- Label: {torch_top_labels[0]}")
print("ONNX:")
print(f"- Box: {onnx_top_boxes[0]}")
print(f"- Score: {onnx_top_scores[0]}")
print(f"- Label: {onnx_top_labels[0]}")
else:
    print("\nNo predictions above confidence threshold to compare")

if __name__ == "__main__":
    # Test the models
    test_model_outputs()

```

Using device: cuda
Loading pretrained SSDLite320 MobileNetV3 model...

Loading ONNX model...

2024-11-09 16:52:08.150740627 [E:onnxruntime:Default, env.cc:234 ThreadMain] pthread_setaffinity_np failed for thread: 1527258, index: 13, mask: {14, }, error code: 22 error msg: Invalid argument. Specify the number of threads explicitly so the affinity is not set.

2024-11-09 16:52:08.158711231 [E:onnxruntime:Default, env.cc:234 ThreadMain] pthread_setaffinity_np failed for thread: 1527260, index: 15, mask: {16, }, error code: 22 error msg: Invalid argument. Specify the number of threads explicitly so the affinity is not set.

2024-11-09 16:52:08.158714549 [E:onnxruntime:Default, env.cc:234 ThreadMain] pthread_setaffinity_np failed for thread: 1527262, index: 17, mask: {18, }, error code: 22 error msg: Invalid argument. Specify the number of threads explicitly so the affinity is not set.

2024-11-09 16:52:08.158715804 [E:onnxruntime:Default, env.cc:234 ThreadMain] pthread_setaffinity_np failed for thread: 1527270, index: 25, mask: {26, }, error code: 22 error msg: Invalid argument. Specify the number of threads explicitly so the affinity is not set.

2024-11-09 16:52:08.158721842 [E:onnxruntime:Default, env.cc:234 ThreadMain] pthread_setaffinity_np failed for thread: 1527263, index: 18, mask: {19, }, error code: 22 error msg: Invalid argument. Specify the number of threads explicitly so the affinity is not set.

2024-11-09 16:52:08.158727128 [E:onnxruntime:Default, env.cc:234 ThreadMain] pthread_setaffinity_np failed for thread: 1527284, index: 39, mask: {40, }, error code: 22 error msg: Invalid argument. Specify the number of threads explicitly so the affinity is not set.

2024-11-09 16:52:08.158733186 [E:onnxruntime:Default, env.cc:234 ThreadMain] pthread_setaffinity_np failed for thread: 1527276, index: 31, mask: {32, }, error code: 22 error msg: Invalid argument. Specify the number of threads explicitly so the affinity is not set.

2024-11-09 16:52:08.158735722 [E:onnxruntime:Default, env.cc:234 ThreadMain] pthread_setaffinity_np failed for thread: 1527274, index: 29, mask: {30, }, error code: 22 error msg: Invalid argument. Specify the number of threads explicitly so the affinity is not set.

2024-11-09 16:52:08.158737247 [E:onnxruntime:Default, env.cc:234 ThreadMain] pthread_setaffinity_np failed for thread: 1527266, index: 21, mask: {22, }, error code: 22 error msg: Invalid argument. Specify the number of threads explicitly so the affinity is not set.

2024-11-09 16:52:08.158713760 [E:onnxruntime:Default, env.cc:234 ThreadMain] pthread_setaffinity_np failed for thread: 1527259, index: 14, mask: {15, }, error code: 22 error msg: Invalid argument. Specify the number of threads explicitly so the affinity is not set.

2024-11-09 16:52:08.158746612 [E:onnxruntime:Default, env.cc:234 ThreadMain] pthread_setaffinity_np failed for thread: 1527273, index: 28, mask: {29, }, error code: 22 error msg: Invalid argument. Specify the number of threads explicitly so the affinity is not set.

2024-11-09 16:52:08.158749741 [E:onnxruntime:Default, env.cc:234 ThreadMain] pthread_setaffinity_np failed for thread: 1527288, index: 43, mask: {44, }, error code: 22 error msg: Invalid argument. Specify the number of threads explicitly so the affinity is not set.

2024-11-09 16:52:08.158751900 [E:onnxruntime:Default, env.cc:234 ThreadMain] pthread_setaffinity_np failed for thread: 1527285, index: 40, mask: {41, }, error code: 22 error msg: Invalid argument. Specify the number of threads explicitly so the affinity is not set.

2024-11-09 16:52:08.158755261 [E:onnxruntime:Default, env.cc:234 ThreadMain] pthread_setaffinity_np failed for thread: 1527290, index: 45, mask: {46, }, error code: 22 error msg: Invalid argument. Specify the number of threads explicitly so the affinity is not set.

2024-11-09 16:52:08.158757146 [E:onnxruntime:Default, env.cc:234 ThreadMain] pthread_setaffinity_np failed for thread: 1527286, index: 41, mask: {42, }, error code: 22 error msg: Invalid argument. Specify the number of threads explicitly so the affinity is not set.

2024-11-09 16:52:08.158749134 [E:onnxruntime:Default, env.cc:234 ThreadMain] pthread_setaffinity_np failed for thread: 1527279, index: 34, mask: {35, }, error code: 22 error msg: Invalid argument. Specify the number of threads explicitly so the affinity is not set.

2024-11-09 16:52:08.158894226 [E:onnxruntime:Default, env.cc:234 ThreadMain] pthread_setaffinity_np failed for thread: 1527269, index: 24, mask: {25, }, error code: 22 error msg: Invalid argument. Specify the number of threads explicitly so the affinity is not set.

2024-11-09 16:52:08.163937836 [E:onnxruntime:Default, env.cc:234 ThreadMain] pthread_setaffinity_np failed for thread: 1527281, index: 36, mask: {37, }, error code: 22 error msg: Invalid argument. Specify the number of threads explicitly so the affinity is not set.

2024-11-09 16:52:08.158716662 [E:onnxruntime:Default, env.cc:234 ThreadMain] pthread_setaffinity_np failed for thread: 1527268, index: 23, mask: {24, }, error code: 22 error msg: Invalid argument. Specify the number of threads explicitly so the affinity is not set.

2024-11-09 16:52:08.158713712 [E:onnxruntime:Default, env.cc:234 ThreadMain] pthread_setaffinity_np failed for thread: 1527265, index: 20, mask: {21, }, error code: 22 error msg: Invalid argument. Specify the number of threads explicitly so the affinity is not set.

2024-11-09 16:52:08.165698845 [E:onnxruntime:Default, env.cc:234 ThreadMain] pthread_setaffinity_np failed for thread: 1527283, index: 38, mask: {39, }, error code: 22 error msg: Invalid argument. Specify the number of threads explicitly so the affinity is not set.

2024-11-09 16:52:08.158733921 [E:onnxruntime:Default, env.cc:234 ThreadMain] pthread_setaffinity_np failed for thread: 1527278, index: 33, mask: {34, }, error code: 22 error msg: Invalid argument. Specify the number of threads explicitly so the affinity is not set.

2024-11-09 16:52:08.158714706 [E:onnxruntime:Default, env.cc:234 ThreadMain] pthread_setaffinity_np failed for thread: 1527267, index: 22, mask: {23, }, error code: 22 error msg: Invalid argument. Specify the number of threads explicitly so the affinity is not set.

2024-11-09 16:52:08.158718717 [E:onnxruntime:Default, env.cc:234 ThreadMain] pthread_setaffinity_np failed for thread: 1527261, index: 16, mask: {17, }, error code: 22 error msg: Invalid argument. Specify the number of threads explicitly so the affinity is not set.

2024-11-09 16:52:08.158740243 [E:onnxruntime:Default, env.cc:234 ThreadMain] pthread_setaffinity_np failed for thread: 1527271, index: 26, mask: {27, }, error code: 22 error msg: Invalid argument. Specify the number of threads explicitly so the affinity is not set.

2024-11-09 16:52:08.167704311 [E:onnxruntime:Default, env.cc:234 ThreadMain] pthread_setaffinity_np failed for thread: 1527289, index: 44, mask: {45, }, error code: 22 error msg: Invalid argument. Specify the number of threads explicitly so the affinity is not set.

2024-11-09 16:52:08.167706039 [E:onnxruntime:Default, env.cc:234 ThreadMain] pthread_setaffinity_np failed for thread: 1527282, index: 37, mask: {38, }, error code: 22 error msg: Invalid argument. Specify the number of threads explicitly so the affinity is not set.

2024-11-09 16:52:08.158740676 [E:onnxruntime:Default, env.cc:234 ThreadMain] pthread_setaffinity_np failed for thread: 1527272, index: 27, mask: {28, }, error code: 22 error msg: Invalid argument. Specify the number of threads explicitly so the affinity is not set.

```
2024-11-09 16:52:08.158713775 [E:onnxruntime:Default, env.cc:234 ThreadMain]
pthread_setaffinity_np failed for thread: 1527264, index: 19, mask: {20, },
error code: 22 error msg: Invalid argument. Specify the number of threads explicitly so the affinity is not set.
2024-11-09 16:52:08.171704154 [E:onnxruntime:Default, env.cc:234 ThreadMain]
pthread_setaffinity_np failed for thread: 1527287, index: 42, mask: {43, },
error code: 22 error msg: Invalid argument. Specify the number of threads explicitly so the affinity is not set.
2024-11-09 16:52:08.171707536 [E:onnxruntime:Default, env.cc:234 ThreadMain]
pthread_setaffinity_np failed for thread: 1527291, index: 46, mask: {47, },
error code: 22 error msg: Invalid argument. Specify the number of threads explicitly so the affinity is not set.
2024-11-09 16:52:08.158722584 [E:onnxruntime:Default, env.cc:234 ThreadMain]
pthread_setaffinity_np failed for thread: 1527275, index: 30, mask: {31, },
error code: 22 error msg: Invalid argument. Specify the number of threads explicitly so the affinity is not set.
2024-11-09 16:52:08.158734225 [E:onnxruntime:Default, env.cc:234 ThreadMain]
pthread_setaffinity_np failed for thread: 1527280, index: 35, mask: {36, },
error code: 22 error msg: Invalid argument. Specify the number of threads explicitly so the affinity is not set.
2024-11-09 16:52:08.158731835 [E:onnxruntime:Default, env.cc:234 ThreadMain]
pthread_setaffinity_np failed for thread: 1527277, index: 32, mask: {33, },
error code: 22 error msg: Invalid argument. Specify the number of threads explicitly so the affinity is not set.
2024-11-09 16:52:08.274047849 [W:onnxruntime:, constant_folding.cc:268 Apply
Impl] Could not find a CPU kernel and hence can't constant fold Sub node '/S
ub_1'
2024-11-09 16:52:08.275058033 [W:onnxruntime:, constant_folding.cc:268 Apply
Impl] Could not find a CPU kernel and hence can't constant fold Sub node '/S
ub'
2024-11-09 16:52:08.313822147 [W:onnxruntime:, constant_folding.cc:268 Apply
Impl] Could not find a CPU kernel and hence can't constant fold Sub node '/S
ub_1'
2024-11-09 16:52:08.313867721 [W:onnxruntime:, constant_folding.cc:268 Apply
Impl] Could not find a CPU kernel and hence can't constant fold Sub node '/S
ub'
2024-11-09 16:52:08.322953298 [W:onnxruntime:, constant_folding.cc:268 Apply
Impl] Could not find a CPU kernel and hence can't constant fold Sub node '/S
ub_1'
2024-11-09 16:52:08.322996157 [W:onnxruntime:, constant_folding.cc:268 Apply
Impl] Could not find a CPU kernel and hence can't constant fold Sub node '/S
ub'
```

```
Getting PyTorch outputs...
```

```
Getting ONNX Runtime outputs...
```

```
Output shapes comparison:
```

```
PyTorch output:
```

- Boxes shape: torch.Size([300, 4])
- Scores shape: torch.Size([300])
- Labels shape: torch.Size([300])

```
ONNX output:
```

- Boxes shape: (1, 4)
- Scores shape: (1,)
- Labels shape: (1,)

```
Number of detections above 0.3 confidence threshold:
```

```
PyTorch: 1
```

```
ONNX: 1
```

```
Comparing top 1 predictions:
```

```
Max absolute differences for top predictions:
```

```
Boxes: 208.67947387695312
```

```
Scores: 0.5006139874458313
```

```
Labels: 21
```

```
Output comparison with precision 0.001:
```

```
Boxes match: False
```

```
Scores match: False
```

```
Labels match: False
```

```
Output comparison with precision 0.0001:
```

```
Boxes match: False
```

```
Scores match: False
```

```
Labels match: False
```

```
Output comparison with precision 1e-05:
```

```
Boxes match: False
```

```
Scores match: False
```

```
Labels match: False
```

```
Output comparison with precision 1e-06:
```

```
Boxes match: False
```

```
Scores match: False
```

```
Labels match: False
```

```
Detailed comparison of first prediction:
```

```
PyTorch:
```

- Box: [18.634064 21.493698 285.86594 317.2563]
- Score: 0.501883327960968
- Label: 22

```
ONNX:
```

- Box: [7.3103256 75.26762 114.61348 108.57682]
- Score: 0.0012693589087575674
- Label: 1

6. Test the inferencing setup

Test the inferencing setup using 1 image from each of the two selected categories. For this, you will need to load the images, preprocess them, and then do inference using the `run()` API from ORT.

(3 points)

```
In [50]: # Cell 6: Test inferencing setup with sample images
def test_inference(image_paths):
    # Create ONNX Runtime session
    session = ort.InferenceSession('ssd_model.onnx')

    # Preprocessing transform
    transform = transforms.Compose([
        transforms.Resize((320, 320)),
        transforms.ToTensor(),
        transforms.Normalize(mean=[0.485, 0.456, 0.406],
                            std=[0.229, 0.224, 0.225])
    ])

    for img_path in image_paths:
        print(f"\nProcessing image: {img_path}")

        # Load and preprocess image
        image = Image.open(img_path).convert('RGB')
        input_tensor = transform(image)
        input_batch = input_tensor.unsqueeze(0)

        # Run inference
        outputs = session.run(None, {'input': input_batch.numpy()})
        print(f"Inference successful for {img_path}")
        print(f"Number of detections: {len(outputs[2])}")

        # Return both image and outputs for visualization
        yield image, outputs
```

7. Parse the response message from ORT

Parse the response message from ORT and annotate the two images. Show inferencing output (bounding boxes with labels) for the two images.

(3 points)

```
In [62]: import torch
import torchvision
import onnxruntime as ort
import onnx
import os
import numpy as np
from PIL import Image
import matplotlib.pyplot as plt
```

```

import matplotlib.patches as patches
from torchvision import transforms
import random

class OpenImagesHelper:
    def __init__(self, root_dir='/scratch/poh2005/data/open_images_temp'):
        self.root_dir = root_dir
        self.class_map = {'airplane': 1, 'boat': 2}
        self.class_ids = {} # Will store OpenImages class IDs

        print(f"Initializing OpenImagesHelper with root_dir: {root_dir}")

    # Load class descriptions and find relevant class IDs
    def load_class_descriptions():
        # Load annotations
        self.image_annotations = self.load_annotations('test')
        print(f"Loaded {len(self.image_annotations)} images with annotations")

    def load_class_descriptions(self):
        """Load class descriptions and find airplane and boat class IDs"""
        desc_path = os.path.join(self.root_dir, 'class-descriptions-boxable.json')
        print(f"Loading class descriptions from: {desc_path}")

        self.class_descriptions = {}
        with open(desc_path, 'r') as f:
            for line in f:
                class_id, class_name = line.strip().split(',')
                class_name = class_name.lower()
                self.class_descriptions[class_id] = class_name

        # Store IDs for our target classes
        if 'airplane' in self.class_map:
            self.class_ids['airplane'] = self.class_map['airplane']
        if 'boat' in self.class_map:
            self.class_ids['boat'] = self.class_map['boat']

        print(f"Found class IDs: {self.class_ids}")

    def load_annotations(self, split='test'):
        """Load annotations with improved filtering and debugging"""
        annotations = {}

        # Read annotations file
        annotation_file = os.path.join(self.root_dir, f'sub-{split}-annotations.csv')
        print(f"Loading annotations from: {annotation_file}")

        valid_class_ids = set(self.class_ids.values())
        count_by_class = {'airplane': 0, 'boat': 0}

        with open(annotation_file, 'r') as f:
            next(f) # Skip header
            for line in f:
                parts = line.strip().split(',')
                image_id = parts[0]
                label = parts[2]

                if label in valid_class_ids:

```

```

        class_name = self.class_descriptions[label].lower()
        count_by_class[class_name] += 1

        if image_id not in annotations:
            annotations[image_id] = {
                'path': os.path.join(self.root_dir, split, f'{image_id}.jpg'),
                'boxes': [],
                'labels': []
            }

        # Add box and label
        xmin = float(parts[4])
        ymin = float(parts[5])
        xmax = float(parts[6])
        ymax = float(parts[7])

        annotations[image_id]['boxes'].append([xmin, ymin, xmax, ymax])
        annotations[image_id]['labels'].append(self.class_map[class_name])

    print("\nAnnotation statistics:")
    for class_name, count in count_by_class.items():
        print(f"{class_name}: {count} annotations")

    return annotations

def get_test_images(self, num_images=5):
    """Get a random sample of images with annotations"""
    airplane_images = []
    boat_images = []

    # Separate images by class
    for image_id, anno in self.image_annotations.items():
        if 1 in anno['labels'] and len(airplane_images) < num_images:
            airplane_images.append(anno['path'])
        if 2 in anno['labels'] and len(boat_images) < num_images:
            boat_images.append(anno['path'])

    # Randomly sample from each class
    selected_images = []

    if airplane_images:
        selected_images.extend(random.sample(airplane_images, min(num_images, len(airplane_images))))
    if boat_images:
        selected_images.extend(random.sample(boat_images, min(num_images, len(boat_images))))

    print(f"\nSelected {len(selected_images)} test images:")
    for img_path in selected_images:
        print(f"- {os.path.basename(img_path)}")

    return selected_images

def run_detection_pipeline(model_path, image_path, output_path=None, confidence_threshold=0.5):
    """Run complete detection pipeline on a single image"""
    try:
        # Initialize ONNX Runtime session
        print(f"Loading model from {model_path}")

```

```

session = ort.InferenceSession(model_path)

# Load and preprocess image
print(f"\nProcessing image: {os.path.basename(image_path)}")
image = Image.open(image_path).convert('RGB')
width, height = image.size
print(f"Original image size: {width}x{height}")

# Preprocess
transform = transforms.Compose([
    transforms.Resize((320, 320)),
    transforms.ToTensor(),
    transforms.Normalize(
        mean=[0.485, 0.456, 0.406],
        std=[0.229, 0.224, 0.225]
    )
])

input_tensor = transform(image)
input_batch = input_tensor.unsqueeze(0)

# Run inference
outputs = session.run(None, {session.get_inputs()[0].name: input_batch})

# Process outputs
boxes = outputs[0]
scores = outputs[1]
labels = outputs[2]

print("\nRaw detection statistics:")
print(f"Number of raw detections: {len(scores)}")
if len(scores) > 0:
    print(f"Score range: {scores.min():.4f} to {scores.max():.4f}")
    print(f"Unique labels: {np.unique(labels)}")
    print("\nTop 5 detections:")
    top_indices = np.argsort(scores)[-5:][::-1]
    for idx in top_indices:
        print(f"Label: {labels[idx]}, Score: {scores[idx]:.4f}, Box: {boxes[idx]}")

# Scale boxes to original image size
scale_x = width / 320
scale_y = height / 320

# Create figure for visualization
plt.figure(figsize=(15, 10))
plt.imshow(image)

# Draw ALL detections with different colors based on confidence
class_names = {1: 'airplane', 2: 'boat'}
valid_detections = []

for i in range(len(scores)):
    label = int(labels[i])
    score = scores[i]

    if label in [1, 2]: # airplane or boat
        ...

```

```

        box = boxes[i]
        scaled_box = [
            box[0] * scale_x,
            box[1] * scale_y,
            box[2] * scale_x,
            box[3] * scale_y
        ]

        # Color based on confidence (red for high confidence, blue for low)
        color = plt.cm.RdYlBu(score)

        rect = patches.Rectangle(
            (scaled_box[0], scaled_box[1]),
            scaled_box[2] - scaled_box[0],
            scaled_box[3] - scaled_box[1],
            linewidth=2,
            edgecolor=color,
            facecolor='none',
            alpha=0.7
        )
        plt.gca().add_patch(rect)

        # Add label with score
        plt.text(
            scaled_box[0], scaled_box[1] - 5,
            f'{class_names[label]} ({score:.4f})',
            color='white',
            bbox=dict(
                facecolor=color,
                alpha=0.7,
                edgecolor='none',
                pad=1
            ),
            fontsize=8
        )

        if score > confidence_threshold:
            valid_detections.append({
                'box': scaled_box,
                'score': score,
                'label': label
            })

    plt.axis('off')
    plt.title(f'Detections for {os.path.basename(image_path)}\n(Color indicates confidence level)')

    # Add colorbar
    norm = plt.Normalize(0, 1)
    sm = plt.cm.ScalarMappable(cmap=plt.cm.RdYlBu, norm=norm)
    plt.colorbar(sm, label='Confidence Score')

    if output_path:
        plt.savefig(output_path, bbox_inches='tight', dpi=300)
        print(f"Saved result to: {output_path}")
    plt.show()
    plt.close()

```

```
    print(f"Found {len(valid_detections)} valid detections")
    return valid_detections

except Exception as e:
    print(f"Pipeline error: {str(e)}")
    import traceback
    traceback.print_exc()
    raise

def main():
    # Initialize dataset helper
    helper = OpenImagesHelper()

    # Get test images
    test_images = helper.get_test_images(num_images=2)  # Get 2 images per category
    print("\nSelected test images:")
    for img in test_images:
        print(f"- {os.path.basename(img)}")

    # Create output directory
    output_dir = "detection_results"
    os.makedirs(output_dir, exist_ok=True)

    # Process each image
    for idx, image_path in enumerate(test_images):
        try:
            print(f"\nProcessing image {idx + 1}/{len(test_images)}")
            output_path = os.path.join(output_dir, f"detection_{idx + 1}.png")

            detections = run_detection_pipeline(
                model_path='ssd_model.onnx',
                image_path=image_path,
                output_path=output_path,
                confidence_threshold=0.001
            )

        except Exception as e:
            print(f"Failed to process {image_path}: {str(e)}")
            continue

if __name__ == "__main__":
    main()
```

```
Initializing OpenImagesHelper with root_dir: /scratch/poh2005/data/open_images_temp
Loading class descriptions from: /scratch/poh2005/data/open_images_temp/classes-descriptions-boxable.csv
Found class IDs: {'boat': '/m/019jd', 'airplane': '/m/0cmf2'}
Loading annotations from: /scratch/poh2005/data/open_images_temp/sub-test-annotations-bbox.csv

Annotation statistics:
airplane: 3272 annotations
boat: 2672 annotations
Loaded 4062 images with annotations

Selected 4 test images:
- a5d2c0bfe01f77c4.jpg
- a1a022596ce3ed9a.jpg
- 17ca4ff2ca2cf884.jpg
- aa47f438b1d110b4.jpg

Selected test images:
- a5d2c0bfe01f77c4.jpg
- a1a022596ce3ed9a.jpg
- 17ca4ff2ca2cf884.jpg
- aa47f438b1d110b4.jpg

Processing image 1/4
Loading model from ssd_model.onnx
```

2024-11-09 16:52:29.136741144 [E:onnxruntime:Default, env.cc:234 ThreadMain] pthread_setaffinity_np failed for thread: 1527323, index: 14, mask: {15, }, error code: 22 error msg: Invalid argument. Specify the number of threads explicitly so the affinity is not set.

2024-11-09 16:52:29.141830976 [E:onnxruntime:Default, env.cc:234 ThreadMain] pthread_setaffinity_np failed for thread: 1527322, index: 13, mask: {14, }, error code: 22 error msg: Invalid argument. Specify the number of threads explicitly so the affinity is not set.

2024-11-09 16:52:29.142713241 [E:onnxruntime:Default, env.cc:234 ThreadMain] pthread_setaffinity_np failed for thread: 1527328, index: 19, mask: {20, }, error code: 22 error msg: Invalid argument. Specify the number of threads explicitly so the affinity is not set.

2024-11-09 16:52:29.142732939 [E:onnxruntime:Default, env.cc:234 ThreadMain] pthread_setaffinity_np failed for thread: 1527342, index: 33, mask: {34, }, error code: 22 error msg: Invalid argument. Specify the number of threads explicitly so the affinity is not set.

2024-11-09 16:52:29.142734861 [E:onnxruntime:Default, env.cc:234 ThreadMain] pthread_setaffinity_np failed for thread: 1527337, index: 28, mask: {29, }, error code: 22 error msg: Invalid argument. Specify the number of threads explicitly so the affinity is not set.

2024-11-09 16:52:29.142736865 [E:onnxruntime:Default, env.cc:234 ThreadMain] pthread_setaffinity_np failed for thread: 1527341, index: 32, mask: {33, }, error code: 22 error msg: Invalid argument. Specify the number of threads explicitly so the affinity is not set.

2024-11-09 16:52:29.142713883 [E:onnxruntime:Default, env.cc:234 ThreadMain] pthread_setaffinity_np failed for thread: 1527332, index: 23, mask: {24, }, error code: 22 error msg: Invalid argument. Specify the number of threads explicitly so the affinity is not set.

2024-11-09 16:52:29.142713741 [E:onnxruntime:Default, env.cc:234 ThreadMain] pthread_setaffinity_np failed for thread: 1527330, index: 21, mask: {22, }, error code: 22 error msg: Invalid argument. Specify the number of threads explicitly so the affinity is not set.

2024-11-09 16:52:29.142746508 [E:onnxruntime:Default, env.cc:234 ThreadMain] pthread_setaffinity_np failed for thread: 1527351, index: 42, mask: {43, }, error code: 22 error msg: Invalid argument. Specify the number of threads explicitly so the affinity is not set.

2024-11-09 16:52:29.142758827 [E:onnxruntime:Default, env.cc:234 ThreadMain] pthread_setaffinity_np failed for thread: 1527350, index: 41, mask: {42, }, error code: 22 error msg: Invalid argument. Specify the number of threads explicitly so the affinity is not set.

2024-11-09 16:52:29.142761180 [E:onnxruntime:Default, env.cc:234 ThreadMain] pthread_setaffinity_np failed for thread: 1527352, index: 43, mask: {44, }, error code: 22 error msg: Invalid argument. Specify the number of threads explicitly so the affinity is not set.

2024-11-09 16:52:29.143707796 [E:onnxruntime:Default, env.cc:234 ThreadMain] pthread_setaffinity_np failed for thread: 1527327, index: 18, mask: {19, }, error code: 22 error msg: Invalid argument. Specify the number of threads explicitly so the affinity is not set.

2024-11-09 16:52:29.144046984 [E:onnxruntime:Default, env.cc:234 ThreadMain] pthread_setaffinity_np failed for thread: 1527335, index: 26, mask: {27, }, error code: 22 error msg: Invalid argument. Specify the number of threads explicitly so the affinity is not set.

2024-11-09 16:52:29.144707367 [E:onnxruntime:Default, env.cc:234 ThreadMain] pthread_setaffinity_np failed for thread: 1527324, index: 15, mask: {16, }, error code: 22 error msg: Invalid argument. Specify the number of threads explicitly so the affinity is not set.

2024-11-09 16:52:29.144714493 [E:onnxruntime:Default, env.cc:234 ThreadMain] pthread_setaffinity_np failed for thread: 1527336, index: 27, mask: {28, }, error code: 22 error msg: Invalid argument. Specify the number of threads explicitly so the affinity is not set.

2024-11-09 16:52:29.144733265 [E:onnxruntime:Default, env.cc:234 ThreadMain] pthread_setaffinity_np failed for thread: 1527346, index: 37, mask: {38, }, error code: 22 error msg: Invalid argument. Specify the number of threads explicitly so the affinity is not set.

2024-11-09 16:52:29.145715181 [E:onnxruntime:Default, env.cc:234 ThreadMain] pthread_setaffinity_np failed for thread: 1527331, index: 22, mask: {23, }, error code: 22 error msg: Invalid argument. Specify the number of threads explicitly so the affinity is not set.

2024-11-09 16:52:29.146706623 [E:onnxruntime:Default, env.cc:234 ThreadMain] pthread_setaffinity_np failed for thread: 1527326, index: 17, mask: {18, }, error code: 22 error msg: Invalid argument. Specify the number of threads explicitly so the affinity is not set.

2024-11-09 16:52:29.144716471 [E:onnxruntime:Default, env.cc:234 ThreadMain] pthread_setaffinity_np failed for thread: 1527334, index: 25, mask: {26, }, error code: 22 error msg: Invalid argument. Specify the number of threads explicitly so the affinity is not set.

2024-11-09 16:52:29.142716329 [E:onnxruntime:Default, env.cc:234 ThreadMain] pthread_setaffinity_np failed for thread: 1527329, index: 20, mask: {21, }, error code: 22 error msg: Invalid argument. Specify the number of threads explicitly so the affinity is not set.

2024-11-09 16:52:29.142713137 [E:onnxruntime:Default, env.cc:234 ThreadMain] pthread_setaffinity_np failed for thread: 1527325, index: 16, mask: {17, }, error code: 22 error msg: Invalid argument. Specify the number of threads explicitly so the affinity is not set.

2024-11-09 16:52:29.142734264 [E:onnxruntime:Default, env.cc:234 ThreadMain] pthread_setaffinity_np failed for thread: 1527340, index: 31, mask: {32, }, error code: 22 error msg: Invalid argument. Specify the number of threads explicitly so the affinity is not set.

2024-11-09 16:52:29.149705543 [E:onnxruntime:Default, env.cc:234 ThreadMain] pthread_setaffinity_np failed for thread: 1527339, index: 30, mask: {31, }, error code: 22 error msg: Invalid argument. Specify the number of threads explicitly so the affinity is not set.

2024-11-09 16:52:29.144715887 [E:onnxruntime:Default, env.cc:234 ThreadMain] pthread_setaffinity_np failed for thread: 1527333, index: 24, mask: {25, }, error code: 22 error msg: Invalid argument. Specify the number of threads explicitly so the affinity is not set.

2024-11-09 16:52:29.150708395 [E:onnxruntime:Default, env.cc:234 ThreadMain] pthread_setaffinity_np failed for thread: 1527338, index: 29, mask: {30, }, error code: 22 error msg: Invalid argument. Specify the number of threads explicitly so the affinity is not set.

2024-11-09 16:52:29.152705835 [E:onnxruntime:Default, env.cc:234 ThreadMain] pthread_setaffinity_np failed for thread: 1527348, index: 39, mask: {40, }, error code: 22 error msg: Invalid argument. Specify the number of threads explicitly so the affinity is not set.

2024-11-09 16:52:29.153706221 [E:onnxruntime:Default, env.cc:234 ThreadMain] pthread_setaffinity_np failed for thread: 1527343, index: 34, mask: {35, }, error code: 22 error msg: Invalid argument. Specify the number of threads explicitly so the affinity is not set.

2024-11-09 16:52:29.154703352 [E:onnxruntime:Default, env.cc:234 ThreadMain] pthread_setaffinity_np failed for thread: 1527345, index: 36, mask: {37, }, error code: 22 error msg: Invalid argument. Specify the number of threads explicitly so the affinity is not set.

2024-11-09 16:52:29.144733502 [E:onnxruntime:Default, env.cc:234 ThreadMain] pthread_setaffinity_np failed for thread: 1527347, index: 38, mask: {39, }, error code: 22 error msg: Invalid argument. Specify the number of threads explicitly so the affinity is not set.

2024-11-09 16:52:29.155704180 [E:onnxruntime:Default, env.cc:234 ThreadMain] pthread_setaffinity_np failed for thread: 1527355, index: 46, mask: {47, }, error code: 22 error msg: Invalid argument. Specify the number of threads explicitly so the affinity is not set.

2024-11-09 16:52:29.155707439 [E:onnxruntime:Default, env.cc:234 ThreadMain] pthread_setaffinity_np failed for thread: 1527344, index: 35, mask: {36, }, error code: 22 error msg: Invalid argument. Specify the number of threads explicitly so the affinity is not set.

2024-11-09 16:52:29.156703158 [E:onnxruntime:Default, env.cc:234 ThreadMain] pthread_setaffinity_np failed for thread: 1527354, index: 45, mask: {46, }, error code: 22 error msg: Invalid argument. Specify the number of threads explicitly so the affinity is not set.

2024-11-09 16:52:29.157702552 [E:onnxruntime:Default, env.cc:234 ThreadMain] pthread_setaffinity_np failed for thread: 1527353, index: 44, mask: {45, }, error code: 22 error msg: Invalid argument. Specify the number of threads explicitly so the affinity is not set.

2024-11-09 16:52:29.157705515 [E:onnxruntime:Default, env.cc:234 ThreadMain] pthread_setaffinity_np failed for thread: 1527349, index: 40, mask: {41, }, error code: 22 error msg: Invalid argument. Specify the number of threads explicitly so the affinity is not set.

2024-11-09 16:52:29.281415522 [W:onnxruntime:, constant_folding.cc:268 Apply Impl] Could not find a CPU kernel and hence can't constant fold Sub node '/S ub_1'

2024-11-09 16:52:29.282422793 [W:onnxruntime:, constant_folding.cc:268 Apply Impl] Could not find a CPU kernel and hence can't constant fold Sub node '/S ub'

2024-11-09 16:52:29.321163020 [W:onnxruntime:, constant_folding.cc:268 Apply Impl] Could not find a CPU kernel and hence can't constant fold Sub node '/S ub_1'

2024-11-09 16:52:29.321206241 [W:onnxruntime:, constant_folding.cc:268 Apply Impl] Could not find a CPU kernel and hence can't constant fold Sub node '/S ub'

2024-11-09 16:52:29.330379966 [W:onnxruntime:, constant_folding.cc:268 Apply Impl] Could not find a CPU kernel and hence can't constant fold Sub node '/S ub_1'

2024-11-09 16:52:29.330418974 [W:onnxruntime:, constant_folding.cc:268 Apply Impl] Could not find a CPU kernel and hence can't constant fold Sub node '/S ub'

```
Processing image: a5d2c0bfe01f77c4.jpg
Original image size: 1024x353
```

```
Raw detection statistics:
Number of raw detections: 1
Score range: 0.9900 to 0.9900
Unique labels: [1]
```

```
Top 5 detections:
```

```
Label: 1, Score: 0.9900, Box: [ 68.473526 320.      278.5265 320.      ]
Pipeline error: Unable to determine Axes to steal space for Colorbar. Either provide the *cax* argument to use as the Axes for the Colorbar, provide the *ax* argument to steal space from it, or add *mappable* to an Axes.
Failed to process /scratch/poh2005/data/open_images_temp/test/a5d2c0bfe01f77c4.jpg: Unable to determine Axes to steal space for Colorbar. Either provide the *cax* argument to use as the Axes for the Colorbar, provide the *ax* argument to steal space from it, or add *mappable* to an Axes.
```

```
Processing image 2/4
```

```
Loading model from ssd_model.onnx
```

```
Traceback (most recent call last):
```

```
  File "/state/partition1/job-53232312/ipykernel_1360913/1391461370.py", line 226, in run_detection_pipeline
    plt.colorbar(sm, label='Confidence Score')
  File "/ext3/miniforge3/lib/python3.12/site-packages/matplotlib/pyplot.py", line 2516, in colorbar
    ret = gcf().colorbar(mappable, cax=cax, ax=ax, **kwargs)
          ~~~~~~
  File "/ext3/miniforge3/lib/python3.12/site-packages/matplotlib/figure.py", line 1215, in colorbar
    raise ValueError()
```

```
ValueError: Unable to determine Axes to steal space for Colorbar. Either provide the *cax* argument to use as the Axes for the Colorbar, provide the *ax* argument to steal space from it, or add *mappable* to an Axes.
```

```
2024-11-09 16:52:30.138195577 [E:onnxruntime:Default, env.cc:234 ThreadMain]
pthread_setaffinity_np failed for thread: 1527370, index: 13, mask: {14, }, error code: 22 error msg: Invalid argument. Specify the number of threads explicitly so the affinity is not set.
```

```
2024-11-09 16:52:30.138209960 [E:onnxruntime:Default, env.cc:234 ThreadMain]
pthread_setaffinity_np failed for thread: 1527371, index: 14, mask: {15, }, error code: 22 error msg: Invalid argument. Specify the number of threads explicitly so the affinity is not set.
```

```
2024-11-09 16:52:30.146703208 [E:onnxruntime:Default, env.cc:234 ThreadMain]
pthread_setaffinity_np failed for thread: 1527373, index: 16, mask: {17, }, error code: 22 error msg: Invalid argument. Specify the number of threads explicitly so the affinity is not set.
```

```
2024-11-09 16:52:30.146708593 [E:onnxruntime:Default, env.cc:234 ThreadMain]
pthread_setaffinity_np failed for thread: 1527372, index: 15, mask: {16, }, error code: 22 error msg: Invalid argument. Specify the number of threads explicitly so the affinity is not set.
```

```
2024-11-09 16:52:30.146713842 [E:onnxruntime:Default, env.cc:234 ThreadMain]
pthread_setaffinity_np failed for thread: 1527379, index: 22, mask: {23, }, error code:
```

```
Processing image: ala022596ce3ed9a.jpg
```

```
Original image size: 1024x768
```

22 error msg: Invalid argument. Specify the number of threads explicitly so the affinity is not set.

2024-11-09 16:52:30.146717963 [E:onnxruntime:Default, env.cc:234 ThreadMain] pthread_setaffinity_np failed for thread: 1527376, index: 19, mask: {20, }, error code: 22 error msg: Invalid argument. Specify the number of threads explicitly so the affinity is not set.

2024-11-09 16:52:30.146725272 [E:onnxruntime:Default, env.cc:234 ThreadMain] pthread_setaffinity_np failed for thread: 1527386, index: 29, mask: {30, }, error code: 22 error msg: Invalid argument. Specify the number of threads explicitly so the affinity is not set.

2024-11-09 16:52:30.146717628 [E:onnxruntime:Default, env.cc:234 ThreadMain] pthread_setaffinity_np failed for thread: 1527380, index: 23, mask: {24, }, error code: 22 error msg: Invalid argument. Specify the number of threads explicitly so the affinity is not set.

2024-11-09 16:52:30.146733376 [E:onnxruntime:Default, env.cc:234 ThreadMain] pthread_setaffinity_np failed for thread: 1527384, index: 27, mask: {28, }, error code: 22 error msg: Invalid argument. Specify the number of threads explicitly so the affinity is not set.

2024-11-09 16:52:30.146719297 [E:onnxruntime:Default, env.cc:234 ThreadMain] pthread_setaffinity_np failed for thread: 1527383, index: 26, mask: {27, }, error code: 22 error msg: Invalid argument. Specify the number of threads explicitly so the affinity is not set.

2024-11-09 16:52:30.146717244 [E:onnxruntime:Default, env.cc:234 ThreadMain] pthread_setaffinity_np failed for thread: 1527382, index: 25, mask: {26, }, error code: 22 error msg: Invalid argument. Specify the number of threads explicitly so the affinity is not set.

2024-11-09 16:52:30.146747936 [E:onnxruntime:Default, env.cc:234 ThreadMain] pthread_setaffinity_np failed for thread: 1527390, index: 33, mask: {34, }, error code: 22 error msg: Invalid argument. Specify the number of threads explicitly so the affinity is not set.

2024-11-09 16:52:30.148335448 [E:onnxruntime:Default, env.cc:234 ThreadMain] pthread_setaffinity_np failed for thread: 1527374, index: 17, mask: {18, }, error code: 22 error msg: Invalid argument. Specify the number of threads explicitly so the affinity is not set.

2024-11-09 16:52:30.153394458 [E:onnxruntime:Default, env.cc:234 ThreadMain] pthread_setaffinity_np failed for thread: 1527388, index: 31, mask: {32, }, error code: 22 error msg: Invalid argument. Specify the number of threads explicitly so the affinity is not set.

2024-11-09 16:52:30.146732304 [E:onnxruntime:Default, env.cc:234 ThreadMain] pthread_setaffinity_np failed for thread: 1527393, index: 36, mask: {37, }, error code: 22 error msg: Invalid argument. Specify the number of threads explicitly so the affinity is not set.

2024-11-09 16:52:30.146733148 [E:onnxruntime:Default, env.cc:234 ThreadMain] pthread_setaffinity_np failed for thread: 1527395, index: 38, mask: {39, }, error code: 22 error msg: Invalid argument. Specify the number of threads explicitly so the affinity is not set.

2024-11-09 16:52:30.146708629 [E:onnxruntime:Default, env.cc:234 ThreadMain] pthread_setaffinity_np failed for thread: 1527375, index: 18, mask: {19, }, error code: 22 error msg: Invalid argument. Specify the number of threads explicitly so the affinity is not set.

2024-11-09 16:52:30.155706264 [E:onnxruntime:Default, env.cc:234 ThreadMain] pthread_setaffinity_np failed for thread: 1527399, index: 42, mask: {43, }, error code: 22 error msg: Invalid argument. Specify the number of threads explicitly so the affinity is not set.

2024-11-09 16:52:30.157141623 [E:onnxruntime:Default, env.cc:234 ThreadMain] pthread_setaffinity_np failed for thread: 1527394, index: 37, mask: {38, },

error code: 22 error msg: Invalid argument. Specify the number of threads explicitly so the affinity is not set.

2024-11-09 16:52:30.158446548 [E:onnxruntime:Default, env.cc:234 ThreadMain] pthread_setaffinity_np failed for thread: 1527403, index: 46, mask: {47, }, error code: 22 error msg: Invalid argument. Specify the number of threads explicitly so the affinity is not set.

2024-11-09 16:52:30.160702938 [E:onnxruntime:Default, env.cc:234 ThreadMain] pthread_setaffinity_np failed for thread: 1527397, index: 40, mask: {41, }, error code: 22 error msg: Invalid argument. Specify the number of threads explicitly so the affinity is not set.

2024-11-09 16:52:30.146739813 [E:onnxruntime:Default, env.cc:234 ThreadMain] pthread_setaffinity_np failed for thread: 1527385, index: 28, mask: {29, }, error code: 22 error msg: Invalid argument. Specify the number of threads explicitly so the affinity is not set.

2024-11-09 16:52:30.146715607 [E:onnxruntime:Default, env.cc:234 ThreadMain] pthread_setaffinity_np failed for thread: 1527381, index: 24, mask: {25, }, error code: 22 error msg: Invalid argument. Specify the number of threads explicitly so the affinity is not set.

2024-11-09 16:52:30.146714183 [E:onnxruntime:Default, env.cc:234 ThreadMain] pthread_setaffinity_np failed for thread: 1527378, index: 21, mask: {22, }, error code: 22 error msg: Invalid argument. Specify the number of threads explicitly so the affinity is not set.

2024-11-09 16:52:30.146715947 [E:onnxruntime:Default, env.cc:234 ThreadMain] pthread_setaffinity_np failed for thread: 1527377, index: 20, mask: {21, }, error code: 22 error msg: Invalid argument. Specify the number of threads explicitly so the affinity is not set.

2024-11-09 16:52:30.155706736 [E:onnxruntime:Default, env.cc:234 ThreadMain] pthread_setaffinity_np failed for thread: 1527391, index: 34, mask: {35, }, error code: 22 error msg: Invalid argument. Specify the number of threads explicitly so the affinity is not set.

2024-11-09 16:52:30.146736264 [E:onnxruntime:Default, env.cc:234 ThreadMain] pthread_setaffinity_np failed for thread: 1527392, index: 35, mask: {36, }, error code: 22 error msg: Invalid argument. Specify the number of threads explicitly so the affinity is not set.

2024-11-09 16:52:30.146728366 [E:onnxruntime:Default, env.cc:234 ThreadMain] pthread_setaffinity_np failed for thread: 1527389, index: 32, mask: {33, }, error code: 22 error msg: Invalid argument. Specify the number of threads explicitly so the affinity is not set.

2024-11-09 16:52:30.164705439 [E:onnxruntime:Default, env.cc:234 ThreadMain] pthread_setaffinity_np failed for thread: 1527396, index: 39, mask: {40, }, error code: 22 error msg: Invalid argument. Specify the number of threads explicitly so the affinity is not set.

2024-11-09 16:52:30.167704486 [E:onnxruntime:Default, env.cc:234 ThreadMain] pthread_setaffinity_np failed for thread: 1527398, index: 41, mask: {42, }, error code: 22 error msg: Invalid argument. Specify the number of threads explicitly so the affinity is not set.

2024-11-09 16:52:30.171703916 [E:onnxruntime:Default, env.cc:234 ThreadMain] pthread_setaffinity_np failed for thread: 1527400, index: 43, mask: {44, }, error code: 22 error msg: Invalid argument. Specify the number of threads explicitly so the affinity is not set.

2024-11-09 16:52:30.176703690 [E:onnxruntime:Default, env.cc:234 ThreadMain] pthread_setaffinity_np failed for thread: 1527387, index: 30, mask: {31, }, error code: 22 error msg: Invalid argument. Specify the number of threads explicitly so the affinity is not set.

2024-11-09 16:52:30.189705051 [E:onnxruntime:Default, env.cc:234 ThreadMain] pthread_setaffinity_np failed for thread: 1527402, index: 45, mask: {46, },

```
error code: 22 error msg: Invalid argument. Specify the number of threads explicitly so the affinity is not set.  
2024-11-09 16:52:30.198705142 [E:onnxruntime:Default, env.cc:234 ThreadMain] pthread_setaffinity_np failed for thread: 1527401, index: 44, mask: {45, }, error code: 22 error msg: Invalid argument. Specify the number of threads explicitly so the affinity is not set.  
2024-11-09 16:52:30.256284620 [W:onnxruntime:, constant_folding.cc:268 ApplyImpl] Could not find a CPU kernel and hence can't constant fold Sub node '/Sub_1'  
2024-11-09 16:52:30.257290383 [W:onnxruntime:, constant_folding.cc:268 ApplyImpl] Could not find a CPU kernel and hence can't constant fold Sub node '/Sub_1'  
2024-11-09 16:52:30.295160395 [W:onnxruntime:, constant_folding.cc:268 ApplyImpl] Could not find a CPU kernel and hence can't constant fold Sub node '/Sub_1'  
2024-11-09 16:52:30.295198392 [W:onnxruntime:, constant_folding.cc:268 ApplyImpl] Could not find a CPU kernel and hence can't constant fold Sub node '/Sub_1'  
2024-11-09 16:52:30.304311912 [W:onnxruntime:, constant_folding.cc:268 ApplyImpl] Could not find a CPU kernel and hence can't constant fold Sub node '/Sub_1'  
2024-11-09 16:52:30.304348394 [W:onnxruntime:, constant_folding.cc:268 ApplyImpl] Could not find a CPU kernel and hence can't constant fold Sub node '/Sub_1'
```

Raw detection statistics:

Number of raw detections: 4
Score range: 0.0017 to 0.9984
Unique labels: [1]

Top 5 detections:

```
Label: 1, Score: 0.9984, Box: [ 64.87378 170.5261 320.      320.      ]  
Label: 1, Score: 0.0359, Box: [320. 320. 320. 320.]  
Label: 1, Score: 0.0226, Box: [261.10327 320.      320.      320.      ]  
Label: 1, Score: 0.0017, Box: [175.86668 213.73224 320.      320.      ]  
Pipeline error: Unable to determine Axes to steal space for Colorbar. Either provide the *cax* argument to use as the Axes for the Colorbar, provide the *ax* argument to steal space from it, or add *mappable* to an Axes.  
Failed to process /scratch/poh2005/data/open_images_temp/test/a1a022596ce3ed9a.jpg: Unable to determine Axes to steal space for Colorbar. Either provide the *cax* argument to use as the Axes for the Colorbar, provide the *ax* argument to steal space from it, or add *mappable* to an Axes.
```

Processing image 3/4

Loading model from ssd_model.onnx

Processing image: 17ca4ff2ca2cf884.jpg

```
Traceback (most recent call last):
  File "/state/partition1/job-53232312/ipykernel_1360913/1391461370.py", line
  226, in run_detection_pipeline
    plt.colorbar(sm, label='Confidence Score')
  File "/ext3/miniforge3/lib/python3.12/site-packages/matplotlib/pyplot.py",
line 2516, in colorbar
    ret = gcf().colorbar(mappable, cax=cax, ax=ax, **kwargs)
    ~~~~~
  File "/ext3/miniforge3/lib/python3.12/site-packages/matplotlib/figure.py",
line 1215, in colorbar
    raise ValueError()
ValueError: Unable to determine Axes to steal space for Colorbar. Either pro
vide the *cax* argument to use as the Axes for the Colorbar, provide the *ax
* argument to steal space from it, or add *mappable* to an Axes.
2024-11-09 16:52:31.029167830 [E:onnxruntime:Default, env.cc:234 ThreadMain]
pthread_setaffinity_np failed for thread: 1527446, index: 13, mask: {14, },
error code: 22 error msg: Invalid argument. Specify the number of threads ex
plicitly so the affinity is not set.
2024-11-09 16:52:31.038714980 [E:onnxruntime:Default, env.cc:234 ThreadMain]
pthread_setaffinity_np failed for thread: 1527473, index: 40, mask: {41, },
error code: 22 error msg: Invalid argument. Specify the number of threads ex
plicitly so the affinity is not set.
2024-11-09 16:52:31.041706737 [E:onnxruntime:Default, env.cc:234 ThreadMain]
pthread_setaffinity_np failed for thread: 1527451, index: 18, mask: {19, },
error code: 22 error msg: Invalid argument. Specify the number of threads ex
plicitly so the affinity is not set.
2024-11-09 16:52:31.041711466 [E:onnxruntime:Default, env.cc:234 ThreadMain]
pthread_setaffinity_np failed for thread: 1527455, index: 22, mask: {23, },
error code: 22 error msg: Invalid argument. Specify the number of threads ex
plicitly so the affinity is not set.
2024-11-09 16:52:31.041711863 [E:onnxruntime:Default, env.cc:234 ThreadMain]
pthread_setaffinity_np failed for thread: 1527458, index: 25, mask: {26, },
error code: 22 error msg: Invalid argument. Specify the number of threads ex
plicitly so the affinity is not set.
2024-11-09 16:52:31.041706596 [E:onnxruntime:Default, env.cc:234 ThreadMain]
pthread_setaffinity_np failed for thread: 1527447, index: 14, mask: {15, },
error code: 22 error msg: Invalid argument. Specify the number of threads ex
plicitly so the affinity is not set.
2024-11-09 16:52:31.041718672 [E:onnxruntime:Default, env.cc:234 ThreadMain]
pthread_setaffinity_np failed for thread: 1527461, index: 28, mask: {29, },
error code: 22 error msg: Invalid argument. Specify the number of threads ex
plicitly so the affinity is not set.
2024-11-09 16:52:31.041723127 [E:onnxruntime:Default, env.cc:234 ThreadMain]
pthread_setaffinity_np failed for thread: 1527459, index: 26, mask: {27, },
error code: 22 error msg: Invalid argument. Specify the number of threads ex
plicitly so the affinity is not set.
2024-11-09 16:52:31.041711955 [E:onnxruntime:Default, env.cc:234 ThreadMain]
pthread_setaffinity_np failed for thread: 1527453, index: 20, mask: {21, },
error code: 22 error msg: Invalid argument. Specify the number of threads ex
plicitly so the affinity is not set.
2024-11-09 16:52:31.041728377 [E:onnxruntime:Default, env.cc:234 ThreadMain]
pthread_setaffinity_np failed for thread: 1527466, index: 33, mask: {34, },
error code: 22 error msg: Invalid argument. Specify the number of threads ex
plicitly so the affinity is not set.
2024-11-09 16:52:31.041712776 [E:onnxruntime:Default, env.cc:234 ThreadMain]
pthread_setaffinity_np failed for thread: 1527452, index: 19, mask: {20, },
```

```
error code: 22 error msg: Invalid argument. Specify the number of threads explicitly so the affinity is not set.  
2024-11-09 16:52:31.041738246 [E:onnxruntime:Default, env.cc:234 ThreadMain] pthread_setaffinity_np failed for thread: 1527464, index: 31, mask: {32, }, error code: 22 error msg: Invalid argument. Specify the number of threads explicitly so the affinity is not set.  
2024-11-09 16:52:31.041744242 [E:onnxruntime:Default, env.cc:234 ThreadMain] pthread_setaffinity_np failed for thread: 1527471, index: 38, mask: {39, }, error code: 22 error msg: Invalid argument. Specify the number of threads explicitly so the affinity is not set.  
2024-11-09 16:52:31.042703014 [E:onnxruntime:Default, env.cc:234 ThreadMain] pthread_setaffinity_np failed for thread: 1527474, index: 41, mask: {42, }, error code: 22 error msg: Invalid argument. Specify the number of threads explicitly so the affinity is not set.  
2024-11-09 16:52:31.042706334 [E:onnxruntime:Default, env.cc:234 ThreadMain] pthread_setaffinity_np failed for thread: 1527450, index: 17, mask: {18, }, error code: 22 error msg: Invalid argument. Specify the number of threads explicitly so the affinity is not set.  
2024-11-09 16:52:31.044357541 [E:onnxruntime:Default, env.cc:234 ThreadMain] pthread_setaffinity_np failed for thread: 1527448, index: 15, mask: {16, }, error code: 22 error msg: Invalid argument. Specify the number of threads explicitly so the affinity is not set.  
2024-11-09 16:52:31.038714839 [E:onnxruntime:Default, env.cc:234 ThreadMain] pthread_setaffinity_np failed for thread: 1527479, index: 46, mask: {47, }, error code: 22 error msg: Invalid argument. Specify the number of threads explicitly so the affinity is not set.  
2024-11-09 16:52:31.046701617 [E:onnxruntime:Default, env.cc:234 ThreadMain] pthread_setaffinity_np failed for thread: 1527475, index: 42, mask: {43, }, error code: 22 error msg: Invalid argument. Specify the number of threads explicitly so the affinity is not set.  
2024-11-09 16:52:31.041710006 [E:onnxruntime:Default, env.cc:234 ThreadMain] pthread_setaffinity_np failed for thread: 1527456, index: 23, mask: {24, }, error code: 22 error msg: Invalid argument. Specify the number of threads explicitly so the affinity is not set.  
2024-11-09 16:52:31.041729208 [E:onnxruntime:Default, env.cc:234 ThreadMain] pthread_setaffinity_np failed for thread: 1527465, index: 32, mask: {33, }, error code: 22 error msg: Invalid argument. Specify the number of threads explicitly so the affinity is not set.  
2024-11-09 16:52:31.041729878 [E:onnxruntime:Default, env.cc:234 ThreadMain] pthread_setaffinity_np failed for thread: 1527470, index: 37, mask: {38, }, error code: 22 error msg: Invalid argument. Specify the number of threads explicitly so the affinity is not set.  
2024-11-09 16:52:31.041719349 [E:onnxruntime:Default, env.cc:234 ThreadMain] pthread_setaffinity_np failed for thread: 1527462, index: 29, mask: {30, }, error code: 22 error msg: Invalid argument. Specify the number of threads explicitly so the affinity is not set.  
2024-11-09 16:52:31.041712334 [E:onnxruntime:Default, env.cc:234 ThreadMain] pthread_setaffinity_np failed for thread: 1527457, index: 24, mask: {25, }, error code: 22 error msg: Invalid argument. Specify the number of threads explicitly so the affinity is not set.  
2024-11-09 16:52:31.041725136 [E:onnxruntime:Default, env.cc:234 ThreadMain] pthread_setaffinity_np failed for thread: 1527469, index: 36, mask: {37, }, error code: 22 error msg: Invalid argument. Specify the number of threads explicitly so the affinity is not set.  
2024-11-09 16:52:31.050706113 [E:onnxruntime:Default, env.cc:234 ThreadMain] pthread_setaffinity_np failed for thread: 1527476, index: 43, mask: {44, },
```

```
error code: 22 error msg: Invalid argument. Specify the number of threads explicitly so the affinity is not set.  
2024-11-09 16:52:31.051702632 [E:onnxruntime:Default, env.cc:234 ThreadMain] pthread_setaffinity_np failed for thread: 1527454, index: 21, mask: {22, }, error code: 22 error msg: Invalid argument. Specify the number of threads explicitly so the affinity is not set.  
2024-11-09 16:52:31.051710693 [E:onnxruntime:Default, env.cc:234 ThreadMain] pthread_setaffinity_np failed for thread: 1527472, index: 39, mask: {40, }, error code: 22 error msg: Invalid argument. Specify the number of threads explicitly so the affinity is not set.  
2024-11-09 16:52:31.054469123 [E:onnxruntime:Default, env.cc:234 ThreadMain] pthread_setaffinity_np failed for thread: 1527463, index: 30, mask: {31, }, error code: 22 error msg: Invalid argument. Specify the number of threads explicitly so the affinity is not set.  
2024-11-09 16:52:31.054703227 [E:onnxruntime:Default, env.cc:234 ThreadMain] pthread_setaffinity_np failed for thread: 1527477, index: 44, mask: {45, }, error code: 22 error msg: Invalid argument. Specify the number of threads explicitly so the affinity is not set.  
2024-11-09 16:52:31.054704202 [E:onnxruntime:Default, env.cc:234 ThreadMain] pthread_setaffinity_np failed for thread: 1527468, index: 35, mask: {36, }, error code: 22 error msg: Invalid argument. Specify the number of threads explicitly so the affinity is not set.  
2024-11-09 16:52:31.058707494 [E:onnxruntime:Default, env.cc:234 ThreadMain] pthread_setaffinity_np failed for thread: 1527467, index: 34, mask: {35, }, error code: 22 error msg: Invalid argument. Specify the number of threads explicitly so the affinity is not set.  
2024-11-09 16:52:31.058722074 [E:onnxruntime:Default, env.cc:234 ThreadMain] pthread_setaffinity_np failed for thread: 1527478, index: 45, mask: {46, }, error code: 22 error msg: Invalid argument. Specify the number of threads explicitly so the affinity is not set.  
2024-11-09 16:52:31.066701413 [E:onnxruntime:Default, env.cc:234 ThreadMain] pthread_setaffinity_np failed for thread: 1527449, index: 16, mask: {17, }, error code: 22 error msg: Invalid argument. Specify the number of threads explicitly so the affinity is not set.  
2024-11-09 16:52:31.070657288 [E:onnxruntime:Default, env.cc:234 ThreadMain] pthread_setaffinity_np failed for thread: 1527460, index: 27, mask: {28, }, error code: 22 error msg: Invalid argument. Specify the number of threads explicitly so the affinity is not set.  
2024-11-09 16:52:31.166263014 [W:onnxruntime:, constant_folding.cc:268 Apply Impl] Could not find a CPU kernel and hence can't constant fold Sub node '/S ub_1'  
2024-11-09 16:52:31.167278069 [W:onnxruntime:, constant_folding.cc:268 Apply Impl] Could not find a CPU kernel and hence can't constant fold Sub node '/S ub'  
2024-11-09 16:52:31.205593569 [W:onnxruntime:, constant_folding.cc:268 Apply Impl] Could not find a CPU kernel and hence can't constant fold Sub node '/S ub_1'  
2024-11-09 16:52:31.205633267 [W:onnxruntime:, constant_folding.cc:268 Apply Impl] Could not find a CPU kernel and hence can't constant fold Sub node '/S ub'  
2024-11-09 16:52:31.214654263 [W:onnxruntime:, constant_folding.cc:268 Apply Impl] Could not find a CPU kernel and hence can't constant fold Sub node '/S ub_1'  
2024-11-09 16:52:31.214699780 [W:onnxruntime:, constant_folding.cc:268 Apply Impl] Could not find a CPU kernel and hence can't constant fold Sub node '/S ub'
```

```
Original image size: 1024x680
```

```
Raw detection statistics:
```

```
Number of raw detections: 0
```

```
Pipeline error: Unable to determine Axes to steal space for Colorbar. Either provide the *cax* argument to use as the Axes for the Colorbar, provide the *ax* argument to steal space from it, or add *mappable* to an Axes.
```

```
Failed to process /scratch/poh2005/data/open_images_temp/test/17ca4ff2ca2cf884.jpg: Unable to determine Axes to steal space for Colorbar. Either provide the *cax* argument to use as the Axes for the Colorbar, provide the *ax* argument to steal space from it, or add *mappable* to an Axes.
```

```
Traceback (most recent call last):
```

```
  File "/state/partition1/job-53232312/ipykernel_1360913/1391461370.py", line 226, in run_detection_pipeline
```

```
    plt.colorbar(sm, label='Confidence Score')
```

```
  File "/ext3/miniforge3/lib/python3.12/site-packages/matplotlib/pyplot.py", line 2516, in colorbar
```

```
    ret = gcf().colorbar(mappable, cax=cax, ax=ax, **kwargs)
```

```
  File "/ext3/miniforge3/lib/python3.12/site-packages/matplotlib/figure.py", line 1215, in colorbar
```

```
    raise ValueError()
```

```
ValueError: Unable to determine Axes to steal space for Colorbar. Either provide the *cax* argument to use as the Axes for the Colorbar, provide the *ax* argument to steal space from it, or add *mappable* to an Axes.
```

```
Processing image 4/4
```

```
Loading model from ssd_model.onnx
```

```
Processing image: aa47f438b1d110b4.jpg
```

```
Original image size: 1024x890
```

2024-11-09 16:52:31.797646557 [E:onnxruntime:Default, env.cc:234 ThreadMain] pthread_setaffinity_np failed for thread: 1527503, index: 13, mask: {14, }, error code: 22 error msg: Invalid argument. Specify the number of threads explicitly so the affinity is not set.

2024-11-09 16:52:31.797662098 [E:onnxruntime:Default, env.cc:234 ThreadMain] pthread_setaffinity_np failed for thread: 1527504, index: 14, mask: {15, }, error code: 22 error msg: Invalid argument. Specify the number of threads explicitly so the affinity is not set.

2024-11-09 16:52:31.797678938 [E:onnxruntime:Default, env.cc:234 ThreadMain] pthread_setaffinity_np failed for thread: 1527505, index: 15, mask: {16, }, error code: 22 error msg: Invalid argument. Specify the number of threads explicitly so the affinity is not set.

2024-11-09 16:52:31.806380940 [E:onnxruntime:Default, env.cc:234 ThreadMain] pthread_setaffinity_np failed for thread: 1527517, index: 27, mask: {28, }, error code: 22 error msg: Invalid argument. Specify the number of threads explicitly so the affinity is not set.

2024-11-09 16:52:31.806703307 [E:onnxruntime:Default, env.cc:234 ThreadMain] pthread_setaffinity_np failed for thread: 1527508, index: 18, mask: {19, }, error code: 22 error msg: Invalid argument. Specify the number of threads explicitly so the affinity is not set.

2024-11-09 16:52:31.806705902 [E:onnxruntime:Default, env.cc:234 ThreadMain] pthread_setaffinity_np failed for thread: 1527509, index: 19, mask: {20, }, error code: 22 error msg: Invalid argument. Specify the number of threads explicitly so the affinity is not set.

2024-11-09 16:52:31.806707463 [E:onnxruntime:Default, env.cc:234 ThreadMain] pthread_setaffinity_np failed for thread: 1527512, index: 22, mask: {23, }, error code: 22 error msg: Invalid argument. Specify the number of threads explicitly so the affinity is not set.

2024-11-09 16:52:31.806709631 [E:onnxruntime:Default, env.cc:234 ThreadMain] pthread_setaffinity_np failed for thread: 1527507, index: 17, mask: {18, }, error code: 22 error msg: Invalid argument. Specify the number of threads explicitly so the affinity is not set.

2024-11-09 16:52:31.806712349 [E:onnxruntime:Default, env.cc:234 ThreadMain] pthread_setaffinity_np failed for thread: 1527516, index: 26, mask: {27, }, error code: 22 error msg: Invalid argument. Specify the number of threads explicitly so the affinity is not set.

2024-11-09 16:52:31.806714794 [E:onnxruntime:Default, env.cc:234 ThreadMain] pthread_setaffinity_np failed for thread: 1527511, index: 21, mask: {22, }, error code: 22 error msg: Invalid argument. Specify the number of threads explicitly so the affinity is not set.

2024-11-09 16:52:31.806718841 [E:onnxruntime:Default, env.cc:234 ThreadMain] pthread_setaffinity_np failed for thread: 1527527, index: 37, mask: {38, }, error code: 22 error msg: Invalid argument. Specify the number of threads explicitly so the affinity is not set.

2024-11-09 16:52:31.806723777 [E:onnxruntime:Default, env.cc:234 ThreadMain] pthread_setaffinity_np failed for thread: 1527526, index: 36, mask: {37, }, error code: 22 error msg: Invalid argument. Specify the number of threads explicitly so the affinity is not set.

2024-11-09 16:52:31.806729941 [E:onnxruntime:Default, env.cc:234 ThreadMain] pthread_setaffinity_np failed for thread: 1527525, index: 35, mask: {36, }, error code: 22 error msg: Invalid argument. Specify the number of threads explicitly so the affinity is not set.

2024-11-09 16:52:31.806708974 [E:onnxruntime:Default, env.cc:234 ThreadMain] pthread_setaffinity_np failed for thread: 1527514, index: 24, mask: {25, }, error code: 22 error msg: Invalid argument. Specify the number of threads explicitly so the affinity is not set.

2024-11-09 16:52:31.806735639 [E:onnxruntime:Default, env.cc:234 ThreadMain] pthread_setaffinity_np failed for thread: 1527528, index: 38, mask: {39, }, error code: 22 error msg: Invalid argument. Specify the number of threads explicitly so the affinity is not set.

2024-11-09 16:52:31.806736843 [E:onnxruntime:Default, env.cc:234 ThreadMain] pthread_setaffinity_np failed for thread: 1527522, index: 32, mask: {33, }, error code: 22 error msg: Invalid argument. Specify the number of threads explicitly so the affinity is not set.

2024-11-09 16:52:31.806715799 [E:onnxruntime:Default, env.cc:234 ThreadMain] pthread_setaffinity_np failed for thread: 1527510, index: 20, mask: {21, }, error code: 22 error msg: Invalid argument. Specify the number of threads explicitly so the affinity is not set.

2024-11-09 16:52:31.806742530 [E:onnxruntime:Default, env.cc:234 ThreadMain] pthread_setaffinity_np failed for thread: 1527531, index: 41, mask: {42, }, error code: 22 error msg: Invalid argument. Specify the number of threads explicitly so the affinity is not set.

2024-11-09 16:52:31.806755963 [E:onnxruntime:Default, env.cc:234 ThreadMain] pthread_setaffinity_np failed for thread: 1527535, index: 45, mask: {46, }, error code: 22 error msg: Invalid argument. Specify the number of threads explicitly so the affinity is not set.

2024-11-09 16:52:31.811421045 [E:onnxruntime:Default, env.cc:234 ThreadMain] pthread_setaffinity_np failed for thread: 1527529, index: 39, mask: {40, }, error code: 22 error msg: Invalid argument. Specify the number of threads explicitly so the affinity is not set.

2024-11-09 16:52:31.811701387 [E:onnxruntime:Default, env.cc:234 ThreadMain] pthread_setaffinity_np failed for thread: 1527532, index: 42, mask: {43, }, error code: 22 error msg: Invalid argument. Specify the number of threads explicitly so the affinity is not set.

2024-11-09 16:52:31.812697974 [E:onnxruntime:Default, env.cc:234 ThreadMain] pthread_setaffinity_np failed for thread: 1527515, index: 25, mask: {26, }, error code: 22 error msg: Invalid argument. Specify the number of threads explicitly so the affinity is not set.

2024-11-09 16:52:31.813706009 [E:onnxruntime:Default, env.cc:234 ThreadMain] pthread_setaffinity_np failed for thread: 1527518, index: 28, mask: {29, }, error code: 22 error msg: Invalid argument. Specify the number of threads explicitly so the affinity is not set.

2024-11-09 16:52:31.806731041 [E:onnxruntime:Default, env.cc:234 ThreadMain] pthread_setaffinity_np failed for thread: 1527523, index: 33, mask: {34, }, error code: 22 error msg: Invalid argument. Specify the number of threads explicitly so the affinity is not set.

2024-11-09 16:52:31.815702600 [E:onnxruntime:Default, env.cc:234 ThreadMain] pthread_setaffinity_np failed for thread: 1527536, index: 46, mask: {47, }, error code: 22 error msg: Invalid argument. Specify the number of threads explicitly so the affinity is not set.

2024-11-09 16:52:31.815704132 [E:onnxruntime:Default, env.cc:234 ThreadMain] pthread_setaffinity_np failed for thread: 1527519, index: 29, mask: {30, }, error code: 22 error msg: Invalid argument. Specify the number of threads explicitly so the affinity is not set.

2024-11-09 16:52:31.816702826 [E:onnxruntime:Default, env.cc:234 ThreadMain] pthread_setaffinity_np failed for thread: 1527533, index: 43, mask: {44, }, error code: 22 error msg: Invalid argument. Specify the number of threads explicitly so the affinity is not set.

2024-11-09 16:52:31.818411320 [E:onnxruntime:Default, env.cc:234 ThreadMain] pthread_setaffinity_np failed for thread: 1527506, index: 16, mask: {17, }, error code: 22 error msg: Invalid argument. Specify the number of threads explicitly so the affinity is not set.

```
2024-11-09 16:52:31.806713829 [E:onnxruntime:Default, env.cc:234 ThreadMain]
pthread_setaffinity_np failed for thread: 1527513, index: 23, mask: {24, },
error code: 22 error msg: Invalid argument. Specify the number of threads ex
plicitly so the affinity is not set.
2024-11-09 16:52:31.820710319 [E:onnxruntime:Default, env.cc:234 ThreadMain]
pthread_setaffinity_np failed for thread: 1527530, index: 40, mask: {41, },
error code: 22 error msg: Invalid argument. Specify the number of threads ex
plicitly so the affinity is not set.
2024-11-09 16:52:31.806730390 [E:onnxruntime:Default, env.cc:234 ThreadMain]
pthread_setaffinity_np failed for thread: 1527521, index: 31, mask: {32, },
error code: 22 error msg: Invalid argument. Specify the number of threads ex
plicitly so the affinity is not set.
2024-11-09 16:52:31.821704784 [E:onnxruntime:Default, env.cc:234 ThreadMain]
pthread_setaffinity_np failed for thread: 1527534, index: 44, mask: {45, },
error code: 22 error msg: Invalid argument. Specify the number of threads ex
plicitly so the affinity is not set.
2024-11-09 16:52:31.806722768 [E:onnxruntime:Default, env.cc:234 ThreadMain]
pthread_setaffinity_np failed for thread: 1527520, index: 30, mask: {31, },
error code: 22 error msg: Invalid argument. Specify the number of threads ex
plicitly so the affinity is not set.
2024-11-09 16:52:31.806707719 [E:onnxruntime:Default, env.cc:234 ThreadMain]
pthread_setaffinity_np failed for thread: 1527524, index: 34, mask: {35, },
error code: 22 error msg: Invalid argument. Specify the number of threads ex
plicitly so the affinity is not set.
2024-11-09 16:52:31.916591320 [W:onnxruntime:, constant_folding.cc:268 Apply
Impl] Could not find a CPU kernel and hence can't constant fold Sub node '/S
ub_1'.
2024-11-09 16:52:31.917611408 [W:onnxruntime:, constant_folding.cc:268 Apply
Impl] Could not find a CPU kernel and hence can't constant fold Sub node '/S
ub'.
2024-11-09 16:52:31.955628122 [W:onnxruntime:, constant_folding.cc:268 Apply
Impl] Could not find a CPU kernel and hence can't constant fold Sub node '/S
ub_1'.
2024-11-09 16:52:31.955668313 [W:onnxruntime:, constant_folding.cc:268 Apply
Impl] Could not find a CPU kernel and hence can't constant fold Sub node '/S
ub'.
2024-11-09 16:52:31.964555530 [W:onnxruntime:, constant_folding.cc:268 Apply
Impl] Could not find a CPU kernel and hence can't constant fold Sub node '/S
ub_1'.
2024-11-09 16:52:31.964593618 [W:onnxruntime:, constant_folding.cc:268 Apply
Impl] Could not find a CPU kernel and hence can't constant fold Sub node '/S
ub'.
```

```
Raw detection statistics:  
Number of raw detections: 2  
Score range: 0.0082 to 0.0319  
Unique labels: [2]
```

```
Top 5 detections:
```

```
Label: 2, Score: 0.0319, Box: [320.      279.44165 320.      320.      ]  
Label: 2, Score: 0.0082, Box: [248.62238 258.18103 320.      320.      ]  
Pipeline error: Unable to determine Axes to steal space for Colorbar. Either  
provide the *cax* argument to use as the Axes for the Colorbar, provide the  
*ax* argument to steal space from it, or add *mappable* to an Axes.  
Failed to process /scratch/poh2005/data/open_images_temp/test/aa47f438b1d110  
b4.jpg: Unable to determine Axes to steal space for Colorbar. Either provide  
the *cax* argument to use as the Axes for the Colorbar, provide the *ax* arg  
ument to steal space from it, or add *mappable* to an Axes.
```

```
Traceback (most recent call last):
```

```
  File "/state/partition1/job-53232312/ipykernel_1360913/1391461370.py", lin  
e 226, in run_detection_pipeline  
    plt.colorbar(sm, label='Confidence Score')  
  File "/ext3/miniforge3/lib/python3.12/site-packages/matplotlib/pyplot.py",  
line 2516, in colorbar  
    ret = gcf().colorbar(mappable, cax=cax, ax=ax, **kwargs)  
    ~~~~~~  
  File "/ext3/miniforge3/lib/python3.12/site-packages/matplotlib/figure.py",  
line 1215, in colorbar  
    raise ValueError()  
ValueError: Unable to determine Axes to steal space for Colorbar. Either pro  
vide the *cax* argument to use as the Axes for the Colorbar, provide the *ax  
* argument to steal space from it, or add *mappable* to an Axes.
```

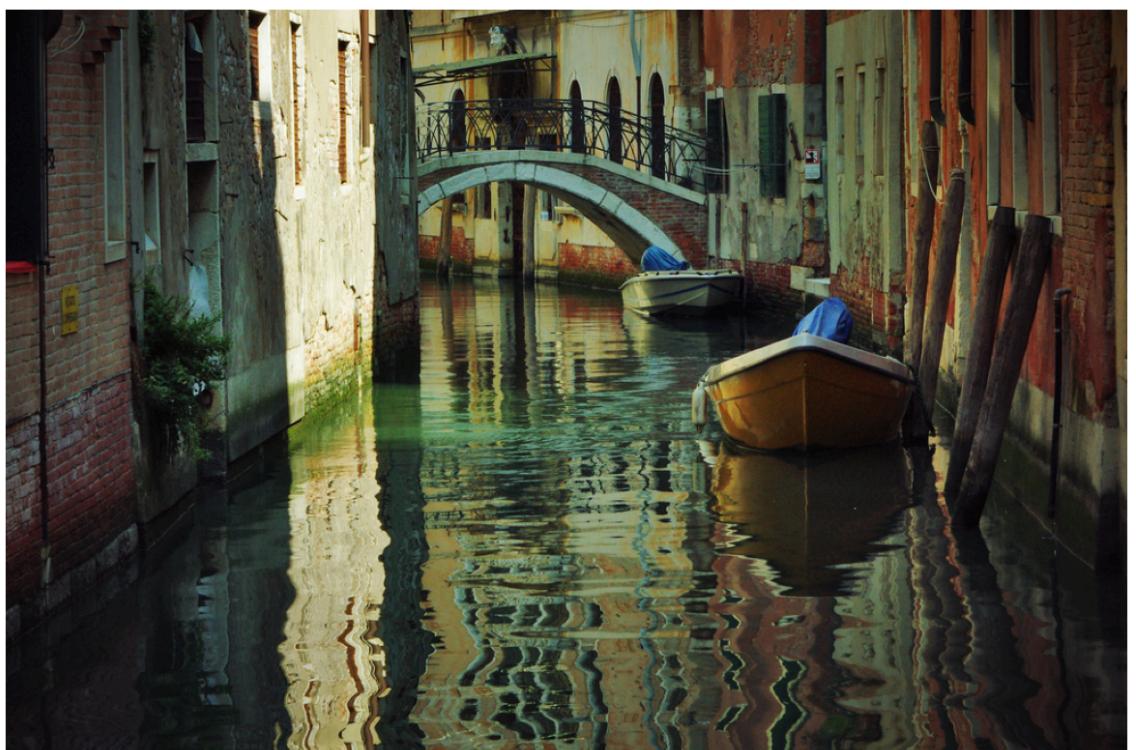
Detections for a5d2c0bfe01f77c4.jpg
(Color indicates confidence)



Detections for a1a022596ce3ed9a.jpg
(Color indicates confidence)



Detections for 17ca4ff2ca2cf884.jpg
(Color indicates confidence)



Detections for aa47f438b1d110b4.jpg
(Color indicates confidence)



References

- Github repo. Shot MultiBox Detector Implementation in Pytorch. Available at <https://github.com/qfgaohtao/pytorch-ssd>
- Pytorch tutorial. Exporting a model from Pytorch to Onnx and running it using Onnx runtime. Available at https://pytorch.org/tutorials/advanced/super_resolution_with_onnxruntime.html
- ONNX tutorial. Inferencing SSD ONNX model using ONNX Runtime Server. Available at <https://github.com/onnx/tutorials/blob/master/tutorials/OnnxRuntimeServerSSD>
- Google. Open Images Dataset V5 + Extensions. Available at <https://storage.googleapis.com/openimages/web/index.html>
- The PASCAL Visual Object Classes Challenge 2007. Available at <http://host.robots.ox.ac.uk/pascal/VOC/voc2007/#testdata>

Problem 2 - Transfer learning: Shallow learning vs Finetuning, Pytorch (20 points)

In this problem, we will train a convolutional neural network for image classification using transfer learning. Transfer learning involves training a base network from scratch on a very large dataset (e.g., Imagenet1K with 1.2M images and 1K categories) and then using this base network either as a feature extractor or as an initialization network for a target task. Thus, two major transfer learning scenarios are as follows:

- **Finetuning the base model:** Instead of random initialization, we initialize the network with a pre-trained network, like the one that is trained on the Imagenet dataset. The rest of the training looks as usual; however, the learning rate schedule for transfer learning may be different.
- **Base model as a fixed feature extractor:** Here, we freeze the weights for all of the network except for the final fully connected layer. This last fully connected layer is replaced with a new one with random weights, and only this layer is trained.

1. Finetuning the model

Select a target dataset from the Visual-Decathlon challenge. Their website (link below) has several datasets which you can download. Select any one of the visual decathlon datasets and make it your target dataset for transfer learning. **Do not select Imagenet1K as the target dataset.**

a. Finetuning the model with ResNet50

You will first load a pre-trained model (ResNet50) and change the final fully connected layer output to the number of classes in the target dataset.

```
In [1]: # TODO
import torch
import torchvision
from torchvision import transforms, datasets
from torch import nn, optim
from torch.utils.data import DataLoader
import time
```

```
In [2]: class TransferLearningExperiment:
    def __init__(self, data_path, num_classes):
        self.device = torch.device("cuda:0" if torch.cuda.is_available() else
```

```

        self.data_path = data_path
        self.num_classes = num_classes
        self.setup_data_transforms()
        self.load_data()

    def setup_data_transforms(self):
        self.data_transforms = {
            'train': transforms.Compose([
                transforms.RandomResizedCrop(224),
                transforms.RandomHorizontalFlip(),
                transforms.ToTensor(),
                transforms.Normalize([0.485, 0.456, 0.406], [0.229, 0.224, 0.225])
            ]),
            'val': transforms.Compose([
                transforms.Resize(256),
                transforms.CenterCrop(224),
                transforms.ToTensor(),
                transforms.Normalize([0.485, 0.456, 0.406], [0.229, 0.224, 0.225])
            ]),
        }

    def load_data(self):
        self.image_datasets = {
            x: datasets.ImageFolder(os.path.join(self.data_path, x), self.data_transforms)
            for x in ['train', 'val']
        }
        self.dataloaders = {
            x: DataLoader(
                self.image_datasets[x],
                batch_size=64,
                shuffle=True,
                num_workers=2
            )
            for x in ['train', 'val']
        }
        self.dataset_sizes = {x: len(self.image_datasets[x]) for x in ['train', 'val']}

```

In [3]: # Helper function to train the model

```

def train_model(model, criterion, optimizer, scheduler, dataloaders, num_epochs=25):
    device = torch.device("cuda:0" if torch.cuda.is_available() else "cpu")
    model.to(device)

    for epoch in range(num_epochs):
        print(f'Epoch {epoch+1}/{num_epochs}')
        print('-' * 10)

        for phase in ['train', 'val']:
            if phase == 'train':
                model.train()
            else:
                model.eval()

            running_loss = 0.0
            running_corrects = 0

            for inputs, labels in dataloaders[phase]:

```

```

        inputs = inputs.to(device)
        labels = labels.to(device)

        optimizer.zero_grad()

        with torch.set_grad_enabled(phase == 'train'):
            outputs = model(inputs)
            _, preds = torch.max(outputs, 1)
            loss = criterion(outputs, labels)

        if phase == 'train':
            loss.backward()
            optimizer.step()

        running_loss += loss.item() * inputs.size(0)
        running_corrects += torch.sum(preds == labels.data)

    if phase == 'train':
        scheduler.step()

    epoch_loss = running_loss / len(dataloaders[phase].dataset)
    epoch_acc = running_corrects.double() / len(dataloaders[phase].dataset)

    print(f'{phase} Loss: {epoch_loss:.4f} Acc: {epoch_acc:.4f}')

return model

```

```

In [4]: def finetune_model(dataset_path, num_classes, learning_rate, num_epochs):
    # Load pre-trained ResNet50
    model = torchvision.models.resnet50(pretrained=True)

    # Modify the final fully connected layer
    num_ftrs = model.fc.in_features
    model.fc = nn.Linear(num_ftrs, num_classes)

    # Define transformations
    data_transforms = {
        'train': transforms.Compose([
            transforms.RandomResizedCrop(224),
            transforms.RandomHorizontalFlip(),
            transforms.ToTensor(),
            transforms.Normalize([0.485, 0.456, 0.406], [0.229, 0.224, 0.225])
        ]),
        'val': transforms.Compose([
            transforms.Resize(256),
            transforms.CenterCrop(224),
            transforms.ToTensor(),
            transforms.Normalize([0.485, 0.456, 0.406], [0.229, 0.224, 0.225])
        ]),
    }

    # Load dataset
    image_datasets = {x: datasets.ImageFolder(f"{dataset_path}/{x}", data_transforms[x])
                     for x in ['train', 'val']}
    dataloaders = {x: DataLoader(image_datasets[x], batch_size=64, shuffle=True,
                                drop_last=True) for x in ['train', 'val']}

```

```

# Define loss function and optimizer
criterion = nn.CrossEntropyLoss()
optimizer = optim.SGD(model.parameters(), lr=learning_rate, momentum=0.9)

# Define learning rate scheduler
scheduler = optim.lr_scheduler.MultiStepLR(optimizer, milestones=[40, 80])

# Train the model
model = train_model(model, criterion, optimizer, scheduler, dataloaders)

return model

```

b. Finetune with learning rate of 0.001

Finetune by setting the same value of hyperparameters (learning rate = 0.001, momentum = 0.9) for all layers. Keep batch size of 64 and train for 150-200 epochs or until the model converges well. Use a multi-step learning rate schedule and decay by a factor of 0.1 ($\gamma = 0.1$). The first drop can happen at epoch 40, the second at epoch 80, and the third at epoch 120 if training for 150 epochs.

(5 points)

```

In [5]: # TODO
def finetune_lr_001(root='./data', num_epochs=150):
    # Define transformations
    data_transforms = {
        'train': transforms.Compose([
            transforms.RandomResizedCrop(224),
            transforms.RandomHorizontalFlip(),
            transforms.ToTensor(),
            transforms.Normalize([0.485, 0.456, 0.406], [0.229, 0.224, 0.225])
        ]),
        'val': transforms.Compose([
            transforms.Resize(256),
            transforms.CenterCrop(224),
            transforms.ToTensor(),
            transforms.Normalize([0.485, 0.456, 0.406], [0.229, 0.224, 0.225])
        ])
    }

    # Load Flowers102 dataset
    train_dataset = torchvision.datasets.Flowers102(root=root,
                                                    split='train',
                                                    transform=data_transforms['train'],
                                                    download=True)

    val_dataset = torchvision.datasets.Flowers102(root=root,
                                                split='val',
                                                transform=data_transforms['val'],
                                                download=True)

    # Create dataloaders

```

```

dataloaders = {
    'train': DataLoader(train_dataset, batch_size=64, shuffle=True, num_workers=4),
    'val': DataLoader(val_dataset, batch_size=64, shuffle=False, num_workers=4)
}

# Load pre-trained ResNet50
model = torchvision.models.resnet50(pretrained=True)

# Modify the final fully connected layer for 102 flower classes
num_ftrs = model.fc.in_features
model.fc = nn.Linear(num_ftrs, 102) # Flowers102 has 102 classes

# Define loss function and optimizer
criterion = nn.CrossEntropyLoss()
optimizer = optim.SGD(model.parameters(), lr=0.001, momentum=0.9)

# Define learning rate scheduler
scheduler = optim.lr_scheduler.MultiStepLR(optimizer, milestones=[40, 80])

# Train the model
model = train_model(model, criterion, optimizer, scheduler, dataloaders)

return model

```

In [6]: `print("Finetuning with lr=0.001")
model_001 = finetune_lr_001()`

```

Finetuning with lr=0.001
/ext3/miniforge3/lib/python3.12/site-packages/torchvision/models/_utils.py:208: UserWarning: The parameter 'pretrained' is deprecated since 0.13 and may be removed in the future, please use 'weights' instead.
    warnings.warn(
/ext3/miniforge3/lib/python3.12/site-packages/torchvision/models/_utils.py:223: UserWarning: Arguments other than a weight enum or `None` for 'weights' are deprecated since 0.13 and may be removed in the future. The current behavior is equivalent to passing `weights=ResNet50_Weights.IMAGENET1K_V1`. You can also use `weights=ResNet50_Weights.DEFAULT` to get the most up-to-date weights.
    warnings.warn(msg)

```

```
Epoch 1/150
-----
train Loss: 4.6520 Acc: 0.0127
val Loss: 4.5650 Acc: 0.0284
Epoch 2/150
-----
train Loss: 4.5012 Acc: 0.0549
val Loss: 4.3822 Acc: 0.1294
Epoch 3/150
-----
train Loss: 4.3157 Acc: 0.1794
val Loss: 4.1880 Acc: 0.2765
Epoch 4/150
-----
train Loss: 4.1239 Acc: 0.3598
val Loss: 3.9627 Acc: 0.4186
Epoch 5/150
-----
train Loss: 3.9069 Acc: 0.4657
val Loss: 3.6918 Acc: 0.4931
Epoch 6/150
-----
train Loss: 3.6447 Acc: 0.5637
val Loss: 3.4168 Acc: 0.5588
Epoch 7/150
-----
train Loss: 3.4010 Acc: 0.6314
val Loss: 3.1187 Acc: 0.6118
Epoch 8/150
-----
train Loss: 3.1405 Acc: 0.6725
val Loss: 2.8331 Acc: 0.6333
Epoch 9/150
-----
train Loss: 2.8947 Acc: 0.7000
val Loss: 2.6129 Acc: 0.6627
Epoch 10/150
-----
train Loss: 2.6483 Acc: 0.7157
val Loss: 2.3856 Acc: 0.7088
Epoch 11/150
-----
train Loss: 2.4204 Acc: 0.7480
val Loss: 2.1833 Acc: 0.7363
Epoch 12/150
-----
train Loss: 2.2059 Acc: 0.7912
val Loss: 1.9866 Acc: 0.7520
Epoch 13/150
-----
train Loss: 2.0137 Acc: 0.8010
val Loss: 1.8247 Acc: 0.7706
Epoch 14/150
-----
train Loss: 1.8345 Acc: 0.8304
val Loss: 1.6609 Acc: 0.7873
```

```
Epoch 15/150
-----
train Loss: 1.7099 Acc: 0.8284
val Loss: 1.5352 Acc: 0.8078
Epoch 16/150
-----
train Loss: 1.5738 Acc: 0.8490
val Loss: 1.4172 Acc: 0.8118
Epoch 17/150
-----
train Loss: 1.4430 Acc: 0.8608
val Loss: 1.3125 Acc: 0.8353
Epoch 18/150
-----
train Loss: 1.2774 Acc: 0.8755
val Loss: 1.2288 Acc: 0.8471
Epoch 19/150
-----
train Loss: 1.2455 Acc: 0.8765
val Loss: 1.1311 Acc: 0.8608
Epoch 20/150
-----
train Loss: 1.0913 Acc: 0.9020
val Loss: 1.0603 Acc: 0.8618
Epoch 21/150
-----
train Loss: 1.0058 Acc: 0.9225
val Loss: 0.9971 Acc: 0.8676
Epoch 22/150
-----
train Loss: 0.9828 Acc: 0.9039
val Loss: 0.9244 Acc: 0.8775
Epoch 23/150
-----
train Loss: 0.8930 Acc: 0.9137
val Loss: 0.8801 Acc: 0.8804
Epoch 24/150
-----
train Loss: 0.8256 Acc: 0.9343
val Loss: 0.8447 Acc: 0.8882
Epoch 25/150
-----
train Loss: 0.7825 Acc: 0.9186
val Loss: 0.7958 Acc: 0.8892
Epoch 26/150
-----
train Loss: 0.7149 Acc: 0.9373
val Loss: 0.7550 Acc: 0.8990
Epoch 27/150
-----
train Loss: 0.6697 Acc: 0.9402
val Loss: 0.7172 Acc: 0.9108
Epoch 28/150
-----
train Loss: 0.6532 Acc: 0.9392
val Loss: 0.7037 Acc: 0.8951
```

```
Epoch 29/150
-----
train Loss: 0.5640 Acc: 0.9520
val Loss: 0.6661 Acc: 0.9137
Epoch 30/150
-----
train Loss: 0.5610 Acc: 0.9520
val Loss: 0.6419 Acc: 0.9088
Epoch 31/150
-----
train Loss: 0.5227 Acc: 0.9578
val Loss: 0.6166 Acc: 0.9147
Epoch 32/150
-----
train Loss: 0.4997 Acc: 0.9471
val Loss: 0.6151 Acc: 0.9118
Epoch 33/150
-----
train Loss: 0.4796 Acc: 0.9539
val Loss: 0.5784 Acc: 0.9206
Epoch 34/150
-----
train Loss: 0.4104 Acc: 0.9706
val Loss: 0.5733 Acc: 0.9157
Epoch 35/150
-----
train Loss: 0.4081 Acc: 0.9667
val Loss: 0.5492 Acc: 0.9255
Epoch 36/150
-----
train Loss: 0.4181 Acc: 0.9588
val Loss: 0.5437 Acc: 0.9245
Epoch 37/150
-----
train Loss: 0.3803 Acc: 0.9667
val Loss: 0.5230 Acc: 0.9314
Epoch 38/150
-----
train Loss: 0.3691 Acc: 0.9627
val Loss: 0.5051 Acc: 0.9216
Epoch 39/150
-----
train Loss: 0.3320 Acc: 0.9637
val Loss: 0.5050 Acc: 0.9225
Epoch 40/150
-----
train Loss: 0.3500 Acc: 0.9647
val Loss: 0.4890 Acc: 0.9284
Epoch 41/150
-----
train Loss: 0.3372 Acc: 0.9637
val Loss: 0.4811 Acc: 0.9324
Epoch 42/150
-----
train Loss: 0.3242 Acc: 0.9725
val Loss: 0.4826 Acc: 0.9363
```

```
Epoch 43/150
-----
train Loss: 0.3088 Acc: 0.9676
val Loss: 0.4797 Acc: 0.9324
Epoch 44/150
-----
train Loss: 0.3400 Acc: 0.9627
val Loss: 0.4767 Acc: 0.9314
Epoch 45/150
-----
train Loss: 0.3007 Acc: 0.9725
val Loss: 0.4774 Acc: 0.9353
Epoch 46/150
-----
train Loss: 0.2995 Acc: 0.9647
val Loss: 0.4767 Acc: 0.9353
Epoch 47/150
-----
train Loss: 0.3050 Acc: 0.9696
val Loss: 0.4782 Acc: 0.9363
Epoch 48/150
-----
train Loss: 0.3159 Acc: 0.9716
val Loss: 0.4727 Acc: 0.9382
Epoch 49/150
-----
train Loss: 0.3146 Acc: 0.9735
val Loss: 0.4713 Acc: 0.9324
Epoch 50/150
-----
train Loss: 0.3294 Acc: 0.9627
val Loss: 0.4702 Acc: 0.9324
Epoch 51/150
-----
train Loss: 0.2872 Acc: 0.9716
val Loss: 0.4714 Acc: 0.9353
Epoch 52/150
-----
train Loss: 0.3069 Acc: 0.9706
val Loss: 0.4688 Acc: 0.9333
Epoch 53/150
-----
train Loss: 0.2980 Acc: 0.9716
val Loss: 0.4719 Acc: 0.9343
Epoch 54/150
-----
train Loss: 0.3308 Acc: 0.9608
val Loss: 0.4665 Acc: 0.9343
Epoch 55/150
-----
train Loss: 0.3008 Acc: 0.9725
val Loss: 0.4663 Acc: 0.9333
Epoch 56/150
-----
train Loss: 0.3141 Acc: 0.9657
val Loss: 0.4647 Acc: 0.9314
```

```
Epoch 57/150
-----
train Loss: 0.2839 Acc: 0.9725
val Loss: 0.4665 Acc: 0.9333
Epoch 58/150
-----
train Loss: 0.3029 Acc: 0.9735
val Loss: 0.4676 Acc: 0.9333
Epoch 59/150
-----
train Loss: 0.2906 Acc: 0.9765
val Loss: 0.4650 Acc: 0.9324
Epoch 60/150
-----
train Loss: 0.2794 Acc: 0.9735
val Loss: 0.4676 Acc: 0.9314
Epoch 61/150
-----
train Loss: 0.2930 Acc: 0.9696
val Loss: 0.4627 Acc: 0.9333
Epoch 62/150
-----
train Loss: 0.3233 Acc: 0.9627
val Loss: 0.4609 Acc: 0.9284
Epoch 63/150
-----
train Loss: 0.2916 Acc: 0.9784
val Loss: 0.4646 Acc: 0.9314
Epoch 64/150
-----
train Loss: 0.2741 Acc: 0.9755
val Loss: 0.4641 Acc: 0.9353
Epoch 65/150
-----
train Loss: 0.3042 Acc: 0.9775
val Loss: 0.4611 Acc: 0.9353
Epoch 66/150
-----
train Loss: 0.3038 Acc: 0.9686
val Loss: 0.4615 Acc: 0.9343
Epoch 67/150
-----
train Loss: 0.2583 Acc: 0.9784
val Loss: 0.4622 Acc: 0.9333
Epoch 68/150
-----
train Loss: 0.2633 Acc: 0.9765
val Loss: 0.4618 Acc: 0.9333
Epoch 69/150
-----
train Loss: 0.2844 Acc: 0.9676
val Loss: 0.4558 Acc: 0.9333
Epoch 70/150
-----
train Loss: 0.2646 Acc: 0.9745
val Loss: 0.4551 Acc: 0.9324
```

```
Epoch 71/150
-----
train Loss: 0.2658 Acc: 0.9735
val Loss: 0.4554 Acc: 0.9294
Epoch 72/150
-----
train Loss: 0.2568 Acc: 0.9765
val Loss: 0.4565 Acc: 0.9343
Epoch 73/150
-----
train Loss: 0.2534 Acc: 0.9775
val Loss: 0.4557 Acc: 0.9343
Epoch 74/150
-----
train Loss: 0.2582 Acc: 0.9804
val Loss: 0.4551 Acc: 0.9333
Epoch 75/150
-----
train Loss: 0.2861 Acc: 0.9725
val Loss: 0.4556 Acc: 0.9353
Epoch 76/150
-----
train Loss: 0.3036 Acc: 0.9667
val Loss: 0.4548 Acc: 0.9343
Epoch 77/150
-----
train Loss: 0.2732 Acc: 0.9775
val Loss: 0.4557 Acc: 0.9333
Epoch 78/150
-----
train Loss: 0.2559 Acc: 0.9843
val Loss: 0.4560 Acc: 0.9333
Epoch 79/150
-----
train Loss: 0.2714 Acc: 0.9745
val Loss: 0.4516 Acc: 0.9324
Epoch 80/150
-----
train Loss: 0.2603 Acc: 0.9755
val Loss: 0.4489 Acc: 0.9333
Epoch 81/150
-----
train Loss: 0.2937 Acc: 0.9686
val Loss: 0.4547 Acc: 0.9353
Epoch 82/150
-----
train Loss: 0.2684 Acc: 0.9725
val Loss: 0.4534 Acc: 0.9324
Epoch 83/150
-----
train Loss: 0.2862 Acc: 0.9716
val Loss: 0.4497 Acc: 0.9333
Epoch 84/150
-----
train Loss: 0.2773 Acc: 0.9725
val Loss: 0.4493 Acc: 0.9353
```

```
Epoch 85/150
-----
train Loss: 0.2702 Acc: 0.9775
val Loss: 0.4499 Acc: 0.9324
Epoch 86/150
-----
train Loss: 0.2663 Acc: 0.9716
val Loss: 0.4512 Acc: 0.9333
Epoch 87/150
-----
train Loss: 0.2971 Acc: 0.9647
val Loss: 0.4522 Acc: 0.9353
Epoch 88/150
-----
train Loss: 0.2754 Acc: 0.9686
val Loss: 0.4529 Acc: 0.9324
Epoch 89/150
-----
train Loss: 0.2888 Acc: 0.9667
val Loss: 0.4479 Acc: 0.9333
Epoch 90/150
-----
train Loss: 0.2527 Acc: 0.9804
val Loss: 0.4494 Acc: 0.9324
Epoch 91/150
-----
train Loss: 0.2764 Acc: 0.9725
val Loss: 0.4467 Acc: 0.9324
Epoch 92/150
-----
train Loss: 0.2668 Acc: 0.9765
val Loss: 0.4462 Acc: 0.9363
Epoch 93/150
-----
train Loss: 0.2563 Acc: 0.9804
val Loss: 0.4464 Acc: 0.9353
Epoch 94/150
-----
train Loss: 0.2402 Acc: 0.9775
val Loss: 0.4487 Acc: 0.9363
Epoch 95/150
-----
train Loss: 0.2700 Acc: 0.9745
val Loss: 0.4464 Acc: 0.9343
Epoch 96/150
-----
train Loss: 0.2727 Acc: 0.9696
val Loss: 0.4480 Acc: 0.9333
Epoch 97/150
-----
train Loss: 0.2728 Acc: 0.9716
val Loss: 0.4464 Acc: 0.9343
Epoch 98/150
-----
train Loss: 0.2683 Acc: 0.9765
val Loss: 0.4458 Acc: 0.9333
```

```
Epoch 99/150
-----
train Loss: 0.2773 Acc: 0.9725
val Loss: 0.4512 Acc: 0.9333
Epoch 100/150
-----
train Loss: 0.2986 Acc: 0.9657
val Loss: 0.4452 Acc: 0.9363
Epoch 101/150
-----
train Loss: 0.2653 Acc: 0.9755
val Loss: 0.4506 Acc: 0.9333
Epoch 102/150
-----
train Loss: 0.2841 Acc: 0.9725
val Loss: 0.4485 Acc: 0.9353
Epoch 103/150
-----
train Loss: 0.2686 Acc: 0.9804
val Loss: 0.4430 Acc: 0.9324
Epoch 104/150
-----
train Loss: 0.2795 Acc: 0.9647
val Loss: 0.4434 Acc: 0.9333
Epoch 105/150
-----
train Loss: 0.2693 Acc: 0.9686
val Loss: 0.4450 Acc: 0.9333
Epoch 106/150
-----
train Loss: 0.2564 Acc: 0.9765
val Loss: 0.4467 Acc: 0.9324
Epoch 107/150
-----
train Loss: 0.2708 Acc: 0.9745
val Loss: 0.4536 Acc: 0.9353
Epoch 108/150
-----
train Loss: 0.2840 Acc: 0.9706
val Loss: 0.4468 Acc: 0.9333
Epoch 109/150
-----
train Loss: 0.2500 Acc: 0.9794
val Loss: 0.4463 Acc: 0.9333
Epoch 110/150
-----
train Loss: 0.2733 Acc: 0.9686
val Loss: 0.4473 Acc: 0.9304
Epoch 111/150
-----
train Loss: 0.2372 Acc: 0.9784
val Loss: 0.4459 Acc: 0.9333
Epoch 112/150
-----
train Loss: 0.2867 Acc: 0.9706
val Loss: 0.4491 Acc: 0.9333
```

```
Epoch 113/150
-----
train Loss: 0.3056 Acc: 0.9627
val Loss: 0.4430 Acc: 0.9353
Epoch 114/150
-----
train Loss: 0.2800 Acc: 0.9725
val Loss: 0.4484 Acc: 0.9353
Epoch 115/150
-----
train Loss: 0.2460 Acc: 0.9804
val Loss: 0.4467 Acc: 0.9333
Epoch 116/150
-----
train Loss: 0.2707 Acc: 0.9725
val Loss: 0.4477 Acc: 0.9353
Epoch 117/150
-----
train Loss: 0.2847 Acc: 0.9696
val Loss: 0.4472 Acc: 0.9324
Epoch 118/150
-----
train Loss: 0.2711 Acc: 0.9765
val Loss: 0.4457 Acc: 0.9343
Epoch 119/150
-----
train Loss: 0.2586 Acc: 0.9706
val Loss: 0.4439 Acc: 0.9333
Epoch 120/150
-----
train Loss: 0.2566 Acc: 0.9784
val Loss: 0.4438 Acc: 0.9314
Epoch 121/150
-----
train Loss: 0.2780 Acc: 0.9686
val Loss: 0.4437 Acc: 0.9314
Epoch 122/150
-----
train Loss: 0.2886 Acc: 0.9598
val Loss: 0.4476 Acc: 0.9314
Epoch 123/150
-----
train Loss: 0.2725 Acc: 0.9735
val Loss: 0.4473 Acc: 0.9343
Epoch 124/150
-----
train Loss: 0.2618 Acc: 0.9745
val Loss: 0.4460 Acc: 0.9343
Epoch 125/150
-----
train Loss: 0.2724 Acc: 0.9686
val Loss: 0.4448 Acc: 0.9333
Epoch 126/150
-----
train Loss: 0.2838 Acc: 0.9686
val Loss: 0.4504 Acc: 0.9343
```

```
Epoch 127/150
-----
train Loss: 0.2782 Acc: 0.9745
val Loss: 0.4458 Acc: 0.9294
Epoch 128/150
-----
train Loss: 0.2642 Acc: 0.9765
val Loss: 0.4493 Acc: 0.9333
Epoch 129/150
-----
train Loss: 0.2820 Acc: 0.9716
val Loss: 0.4476 Acc: 0.9333
Epoch 130/150
-----
train Loss: 0.2730 Acc: 0.9647
val Loss: 0.4451 Acc: 0.9294
Epoch 131/150
-----
train Loss: 0.2850 Acc: 0.9716
val Loss: 0.4467 Acc: 0.9333
Epoch 132/150
-----
train Loss: 0.2561 Acc: 0.9755
val Loss: 0.4465 Acc: 0.9314
Epoch 133/150
-----
train Loss: 0.2613 Acc: 0.9745
val Loss: 0.4491 Acc: 0.9333
Epoch 134/150
-----
train Loss: 0.2701 Acc: 0.9745
val Loss: 0.4422 Acc: 0.9333
Epoch 135/150
-----
train Loss: 0.2443 Acc: 0.9814
val Loss: 0.4474 Acc: 0.9333
Epoch 136/150
-----
train Loss: 0.2534 Acc: 0.9755
val Loss: 0.4495 Acc: 0.9343
Epoch 137/150
-----
train Loss: 0.2558 Acc: 0.9775
val Loss: 0.4481 Acc: 0.9304
Epoch 138/150
-----
train Loss: 0.2612 Acc: 0.9725
val Loss: 0.4483 Acc: 0.9343
Epoch 139/150
-----
train Loss: 0.2842 Acc: 0.9716
val Loss: 0.4435 Acc: 0.9333
Epoch 140/150
-----
train Loss: 0.2614 Acc: 0.9745
val Loss: 0.4490 Acc: 0.9333
```

```
Epoch 141/150
-----
train Loss: 0.2522 Acc: 0.9716
val Loss: 0.4500 Acc: 0.9333
Epoch 142/150
-----
train Loss: 0.2724 Acc: 0.9676
val Loss: 0.4422 Acc: 0.9343
Epoch 143/150
-----
train Loss: 0.2698 Acc: 0.9667
val Loss: 0.4446 Acc: 0.9333
Epoch 144/150
-----
train Loss: 0.2714 Acc: 0.9725
val Loss: 0.4499 Acc: 0.9343
Epoch 145/150
-----
train Loss: 0.2870 Acc: 0.9657
val Loss: 0.4389 Acc: 0.9343
Epoch 146/150
-----
train Loss: 0.2452 Acc: 0.9804
val Loss: 0.4455 Acc: 0.9343
Epoch 147/150
-----
train Loss: 0.2614 Acc: 0.9833
val Loss: 0.4463 Acc: 0.9343
Epoch 148/150
-----
train Loss: 0.2612 Acc: 0.9686
val Loss: 0.4504 Acc: 0.9343
Epoch 149/150
-----
train Loss: 0.2670 Acc: 0.9716
val Loss: 0.4436 Acc: 0.9343
Epoch 150/150
-----
train Loss: 0.2594 Acc: 0.9775
val Loss: 0.4467 Acc: 0.9314
```

c. Experiment with higher learning rates

Next, keeping all the hyperparameters the same as before (including the multi-step learning rate schedule), change the learning rate to 0.01 and 0.1 uniformly for all the layers. Perform two experiments, one with a learning rate of 0.01 and one with 0.1. Again, finetune the model using the same multi-step learning rate schedule as in Part b, and report the final accuracy. Compare the accuracy with the three learning rates and discuss which learning rate gave the best accuracy on the target dataset.

(6 points)

In [7]:

```
# TODO
def experiment_higher_lr(learning_rate, root='./data', num_epochs=150):
    # Define transformations
    data_transforms = {
        'train': transforms.Compose([
            transforms.RandomResizedCrop(224),
            transforms.RandomHorizontalFlip(),
            transforms.ToTensor(),
            transforms.Normalize([0.485, 0.456, 0.406], [0.229, 0.224, 0.225])
        ]),
        'val': transforms.Compose([
            transforms.Resize(256),
            transforms.CenterCrop(224),
            transforms.ToTensor(),
            transforms.Normalize([0.485, 0.456, 0.406], [0.229, 0.224, 0.225])
        ])
    }

    # Load Flowers102 dataset
    train_dataset = torchvision.datasets.Flowers102(root=root,
                                                    split='train',
                                                    transform=data_transforms['train'],
                                                    download=True)

    val_dataset = torchvision.datasets.Flowers102(root=root,
                                                split='val',
                                                transform=data_transforms['val'],
                                                download=True)

    # Create dataloaders
    dataloaders = {
        'train': DataLoader(train_dataset, batch_size=64, shuffle=True, num_workers=4),
        'val': DataLoader(val_dataset, batch_size=64, shuffle=False, num_workers=4)
    }

    # Load pre-trained ResNet50
    model = torchvision.models.resnet50(pretrained=True)

    # Modify the final fully connected layer
    num_ftrs = model.fc.in_features
    model.fc = nn.Linear(num_ftrs, 102)

    # Define loss function and optimizer with higher learning rate
    criterion = nn.CrossEntropyLoss()
    optimizer = optim.SGD(model.parameters(), lr=learning_rate, momentum=0.9)

    # Define learning rate scheduler
    scheduler = optim.lr_scheduler.MultiStepLR(optimizer, milestones=[40, 80])

    # Train the model
    model = train_model(model, criterion, optimizer, scheduler, dataloaders, num_epochs)

    return model
```

```
In [9]: # Part 1c: Experiment with higher learning rates
print("\nFinetuning with lr=0.01")
model_01 = experiment_higher_lr(0.01)
```

```
Finetuning with lr=0.01
Epoch 1/150
-----
train Loss: 4.5126 Acc: 0.0559
val Loss: 3.6467 Acc: 0.3735
Epoch 2/150
-----
train Loss: 3.1092 Acc: 0.4441
val Loss: 1.8815 Acc: 0.5716
Epoch 3/150
-----
train Loss: 1.6507 Acc: 0.7147
val Loss: 1.1735 Acc: 0.7647
Epoch 4/150
-----
train Loss: 0.8968 Acc: 0.8343
val Loss: 0.7800 Acc: 0.8441
Epoch 5/150
-----
train Loss: 0.5598 Acc: 0.9059
val Loss: 0.5735 Acc: 0.8853
Epoch 6/150
-----
train Loss: 0.4215 Acc: 0.9206
val Loss: 0.5365 Acc: 0.8784
Epoch 7/150
-----
train Loss: 0.3144 Acc: 0.9333
val Loss: 0.4699 Acc: 0.9010
Epoch 8/150
-----
train Loss: 0.2398 Acc: 0.9529
val Loss: 0.4187 Acc: 0.9039
Epoch 9/150
-----
train Loss: 0.2090 Acc: 0.9569
val Loss: 0.3802 Acc: 0.9206
Epoch 10/150
-----
train Loss: 0.1529 Acc: 0.9765
val Loss: 0.4027 Acc: 0.9127
Epoch 11/150
-----
train Loss: 0.1714 Acc: 0.9676
val Loss: 0.4429 Acc: 0.8922
Epoch 12/150
-----
train Loss: 0.1487 Acc: 0.9716
val Loss: 0.3781 Acc: 0.9088
Epoch 13/150
-----
train Loss: 0.1395 Acc: 0.9755
val Loss: 0.3601 Acc: 0.9137
Epoch 14/150
-----
train Loss: 0.1588 Acc: 0.9676
```

```
val Loss: 0.3680 Acc: 0.9196
Epoch 15/150
-----
train Loss: 0.1469 Acc: 0.9676
val Loss: 0.3741 Acc: 0.9127
Epoch 16/150
-----
train Loss: 0.1192 Acc: 0.9755
val Loss: 0.3156 Acc: 0.9324
Epoch 17/150
-----
train Loss: 0.1333 Acc: 0.9716
val Loss: 0.3876 Acc: 0.9069
Epoch 18/150
-----
train Loss: 0.1354 Acc: 0.9735
val Loss: 0.3671 Acc: 0.9127
Epoch 19/150
-----
train Loss: 0.1335 Acc: 0.9725
val Loss: 0.3797 Acc: 0.9137
Epoch 20/150
-----
train Loss: 0.1128 Acc: 0.9755
val Loss: 0.3567 Acc: 0.9098
Epoch 21/150
-----
train Loss: 0.1337 Acc: 0.9696
val Loss: 0.3994 Acc: 0.8971
Epoch 22/150
-----
train Loss: 0.1114 Acc: 0.9755
val Loss: 0.3502 Acc: 0.9157
Epoch 23/150
-----
train Loss: 0.0751 Acc: 0.9853
val Loss: 0.3157 Acc: 0.9196
Epoch 24/150
-----
train Loss: 0.0995 Acc: 0.9784
val Loss: 0.3325 Acc: 0.9235
Epoch 25/150
-----
train Loss: 0.0878 Acc: 0.9853
val Loss: 0.3027 Acc: 0.9196
Epoch 26/150
-----
train Loss: 0.0963 Acc: 0.9775
val Loss: 0.3488 Acc: 0.9235
Epoch 27/150
-----
train Loss: 0.1009 Acc: 0.9775
val Loss: 0.3855 Acc: 0.9078
Epoch 28/150
-----
train Loss: 0.1112 Acc: 0.9696
```

```
val Loss: 0.3937 Acc: 0.9039
Epoch 29/150
-----
train Loss: 0.0755 Acc: 0.9863
val Loss: 0.3282 Acc: 0.9245
Epoch 30/150
-----
train Loss: 0.0967 Acc: 0.9804
val Loss: 0.3738 Acc: 0.9098
Epoch 31/150
-----
train Loss: 0.0830 Acc: 0.9843
val Loss: 0.3378 Acc: 0.9196
Epoch 32/150
-----
train Loss: 0.1060 Acc: 0.9755
val Loss: 0.3376 Acc: 0.9088
Epoch 33/150
-----
train Loss: 0.0719 Acc: 0.9853
val Loss: 0.3548 Acc: 0.9108
Epoch 34/150
-----
train Loss: 0.0760 Acc: 0.9853
val Loss: 0.3605 Acc: 0.9127
Epoch 35/150
-----
train Loss: 0.1065 Acc: 0.9765
val Loss: 0.3734 Acc: 0.9147
Epoch 36/150
-----
train Loss: 0.0883 Acc: 0.9814
val Loss: 0.3786 Acc: 0.9147
Epoch 37/150
-----
train Loss: 0.0696 Acc: 0.9843
val Loss: 0.3630 Acc: 0.9206
Epoch 38/150
-----
train Loss: 0.0894 Acc: 0.9814
val Loss: 0.3582 Acc: 0.9157
Epoch 39/150
-----
train Loss: 0.0747 Acc: 0.9833
val Loss: 0.3453 Acc: 0.9176
Epoch 40/150
-----
train Loss: 0.0854 Acc: 0.9794
val Loss: 0.3790 Acc: 0.9078
Epoch 41/150
-----
train Loss: 0.0824 Acc: 0.9775
val Loss: 0.3334 Acc: 0.9245
Epoch 42/150
-----
train Loss: 0.0685 Acc: 0.9882
```

```
val Loss: 0.3233 Acc: 0.9275
Epoch 43/150
-----
train Loss: 0.0550 Acc: 0.9843
val Loss: 0.3153 Acc: 0.9314
Epoch 44/150
-----
train Loss: 0.0634 Acc: 0.9882
val Loss: 0.3094 Acc: 0.9373
Epoch 45/150
-----
train Loss: 0.0522 Acc: 0.9882
val Loss: 0.3049 Acc: 0.9333
Epoch 46/150
-----
train Loss: 0.0722 Acc: 0.9833
val Loss: 0.3008 Acc: 0.9353
Epoch 47/150
-----
train Loss: 0.0467 Acc: 0.9902
val Loss: 0.2967 Acc: 0.9333
Epoch 48/150
-----
train Loss: 0.0527 Acc: 0.9892
val Loss: 0.2978 Acc: 0.9333
Epoch 49/150
-----
train Loss: 0.0301 Acc: 0.9931
val Loss: 0.2968 Acc: 0.9353
Epoch 50/150
-----
train Loss: 0.0440 Acc: 0.9922
val Loss: 0.2978 Acc: 0.9343
Epoch 51/150
-----
train Loss: 0.0461 Acc: 0.9912
val Loss: 0.2963 Acc: 0.9343
Epoch 52/150
-----
train Loss: 0.0389 Acc: 0.9912
val Loss: 0.2938 Acc: 0.9333
Epoch 53/150
-----
train Loss: 0.0302 Acc: 0.9941
val Loss: 0.2929 Acc: 0.9343
Epoch 54/150
-----
train Loss: 0.0530 Acc: 0.9882
val Loss: 0.2910 Acc: 0.9343
Epoch 55/150
-----
train Loss: 0.0373 Acc: 0.9902
val Loss: 0.2908 Acc: 0.9343
Epoch 56/150
-----
train Loss: 0.0437 Acc: 0.9902
```

```
val Loss: 0.2920 Acc: 0.9363
Epoch 57/150
-----
train Loss: 0.0430 Acc: 0.9892
val Loss: 0.2906 Acc: 0.9382
Epoch 58/150
-----
train Loss: 0.0351 Acc: 0.9931
val Loss: 0.2900 Acc: 0.9353
Epoch 59/150
-----
train Loss: 0.0449 Acc: 0.9892
val Loss: 0.2874 Acc: 0.9353
Epoch 60/150
-----
train Loss: 0.0282 Acc: 0.9922
val Loss: 0.2877 Acc: 0.9314
Epoch 61/150
-----
train Loss: 0.0587 Acc: 0.9853
val Loss: 0.2871 Acc: 0.9324
Epoch 62/150
-----
train Loss: 0.0330 Acc: 0.9922
val Loss: 0.2881 Acc: 0.9324
Epoch 63/150
-----
train Loss: 0.0241 Acc: 0.9951
val Loss: 0.2880 Acc: 0.9353
Epoch 64/150
-----
train Loss: 0.0460 Acc: 0.9902
val Loss: 0.2897 Acc: 0.9324
Epoch 65/150
-----
train Loss: 0.0263 Acc: 0.9961
val Loss: 0.2872 Acc: 0.9343
Epoch 66/150
-----
train Loss: 0.0491 Acc: 0.9902
val Loss: 0.2843 Acc: 0.9363
Epoch 67/150
-----
train Loss: 0.0282 Acc: 0.9941
val Loss: 0.2842 Acc: 0.9363
Epoch 68/150
-----
train Loss: 0.0573 Acc: 0.9853
val Loss: 0.2832 Acc: 0.9382
Epoch 69/150
-----
train Loss: 0.0348 Acc: 0.9912
val Loss: 0.2849 Acc: 0.9314
Epoch 70/150
-----
train Loss: 0.0350 Acc: 0.9922
```

```
val Loss: 0.2824 Acc: 0.9353
Epoch 71/150
-----
train Loss: 0.0449 Acc: 0.9892
val Loss: 0.2790 Acc: 0.9382
Epoch 72/150
-----
train Loss: 0.0313 Acc: 0.9961
val Loss: 0.2752 Acc: 0.9392
Epoch 73/150
-----
train Loss: 0.0332 Acc: 0.9931
val Loss: 0.2799 Acc: 0.9363
Epoch 74/150
-----
train Loss: 0.0435 Acc: 0.9892
val Loss: 0.2772 Acc: 0.9402
Epoch 75/150
-----
train Loss: 0.0302 Acc: 0.9941
val Loss: 0.2766 Acc: 0.9392
Epoch 76/150
-----
train Loss: 0.0305 Acc: 0.9951
val Loss: 0.2757 Acc: 0.9402
Epoch 77/150
-----
train Loss: 0.0268 Acc: 0.9961
val Loss: 0.2764 Acc: 0.9402
Epoch 78/150
-----
train Loss: 0.0242 Acc: 0.9941
val Loss: 0.2763 Acc: 0.9392
Epoch 79/150
-----
train Loss: 0.0401 Acc: 0.9902
val Loss: 0.2755 Acc: 0.9392
Epoch 80/150
-----
train Loss: 0.0298 Acc: 0.9922
val Loss: 0.2765 Acc: 0.9422
Epoch 81/150
-----
train Loss: 0.0275 Acc: 0.9951
val Loss: 0.2757 Acc: 0.9412
Epoch 82/150
-----
train Loss: 0.0238 Acc: 0.9951
val Loss: 0.2752 Acc: 0.9402
Epoch 83/150
-----
train Loss: 0.0339 Acc: 0.9931
val Loss: 0.2747 Acc: 0.9412
Epoch 84/150
-----
train Loss: 0.0221 Acc: 0.9951
```

```
val Loss: 0.2744 Acc: 0.9402
Epoch 85/150
-----
train Loss: 0.0349 Acc: 0.9941
val Loss: 0.2748 Acc: 0.9431
Epoch 86/150
-----
train Loss: 0.0301 Acc: 0.9931
val Loss: 0.2736 Acc: 0.9412
Epoch 87/150
-----
train Loss: 0.0331 Acc: 0.9931
val Loss: 0.2742 Acc: 0.9412
Epoch 88/150
-----
train Loss: 0.0241 Acc: 0.9941
val Loss: 0.2746 Acc: 0.9412
Epoch 89/150
-----
train Loss: 0.0260 Acc: 0.9951
val Loss: 0.2753 Acc: 0.9412
Epoch 90/150
-----
train Loss: 0.0365 Acc: 0.9912
val Loss: 0.2747 Acc: 0.9412
Epoch 91/150
-----
train Loss: 0.0311 Acc: 0.9931
val Loss: 0.2754 Acc: 0.9412
Epoch 92/150
-----
train Loss: 0.0373 Acc: 0.9892
val Loss: 0.2752 Acc: 0.9412
Epoch 93/150
-----
train Loss: 0.0302 Acc: 0.9941
val Loss: 0.2727 Acc: 0.9412
Epoch 94/150
-----
train Loss: 0.0274 Acc: 0.9951
val Loss: 0.2733 Acc: 0.9412
Epoch 95/150
-----
train Loss: 0.0305 Acc: 0.9941
val Loss: 0.2751 Acc: 0.9422
Epoch 96/150
-----
train Loss: 0.0370 Acc: 0.9922
val Loss: 0.2744 Acc: 0.9392
Epoch 97/150
-----
train Loss: 0.0271 Acc: 0.9941
val Loss: 0.2729 Acc: 0.9422
Epoch 98/150
-----
train Loss: 0.0305 Acc: 0.9922
```

```
val Loss: 0.2746 Acc: 0.9392
Epoch 99/150
-----
train Loss: 0.0122 Acc: 0.9990
val Loss: 0.2778 Acc: 0.9412
Epoch 100/150
-----
train Loss: 0.0396 Acc: 0.9892
val Loss: 0.2750 Acc: 0.9412
Epoch 101/150
-----
train Loss: 0.0314 Acc: 0.9951
val Loss: 0.2712 Acc: 0.9422
Epoch 102/150
-----
train Loss: 0.0374 Acc: 0.9902
val Loss: 0.2720 Acc: 0.9402
Epoch 103/150
-----
train Loss: 0.0384 Acc: 0.9931
val Loss: 0.2735 Acc: 0.9402
Epoch 104/150
-----
train Loss: 0.0321 Acc: 0.9941
val Loss: 0.2745 Acc: 0.9431
Epoch 105/150
-----
train Loss: 0.0443 Acc: 0.9931
val Loss: 0.2716 Acc: 0.9422
Epoch 106/150
-----
train Loss: 0.0251 Acc: 0.9961
val Loss: 0.2732 Acc: 0.9392
Epoch 107/150
-----
train Loss: 0.0432 Acc: 0.9902
val Loss: 0.2748 Acc: 0.9412
Epoch 108/150
-----
train Loss: 0.0304 Acc: 0.9961
val Loss: 0.2718 Acc: 0.9402
Epoch 109/150
-----
train Loss: 0.0240 Acc: 0.9931
val Loss: 0.2757 Acc: 0.9412
Epoch 110/150
-----
train Loss: 0.0368 Acc: 0.9922
val Loss: 0.2719 Acc: 0.9402
Epoch 111/150
-----
train Loss: 0.0263 Acc: 0.9951
val Loss: 0.2732 Acc: 0.9392
Epoch 112/150
-----
train Loss: 0.0287 Acc: 0.9912
```

```
val Loss: 0.2745 Acc: 0.9402
Epoch 113/150
-----
train Loss: 0.0201 Acc: 0.9961
val Loss: 0.2755 Acc: 0.9402
Epoch 114/150
-----
train Loss: 0.0327 Acc: 0.9922
val Loss: 0.2756 Acc: 0.9402
Epoch 115/150
-----
train Loss: 0.0233 Acc: 0.9951
val Loss: 0.2720 Acc: 0.9431
Epoch 116/150
-----
train Loss: 0.0247 Acc: 0.9922
val Loss: 0.2744 Acc: 0.9412
Epoch 117/150
-----
train Loss: 0.0273 Acc: 0.9941
val Loss: 0.2779 Acc: 0.9382
Epoch 118/150
-----
train Loss: 0.0388 Acc: 0.9892
val Loss: 0.2748 Acc: 0.9392
Epoch 119/150
-----
train Loss: 0.0382 Acc: 0.9902
val Loss: 0.2720 Acc: 0.9392
Epoch 120/150
-----
train Loss: 0.0310 Acc: 0.9951
val Loss: 0.2740 Acc: 0.9402
Epoch 121/150
-----
train Loss: 0.0358 Acc: 0.9931
val Loss: 0.2724 Acc: 0.9422
Epoch 122/150
-----
train Loss: 0.0309 Acc: 0.9961
val Loss: 0.2717 Acc: 0.9412
Epoch 123/150
-----
train Loss: 0.0174 Acc: 0.9961
val Loss: 0.2720 Acc: 0.9402
Epoch 124/150
-----
train Loss: 0.0311 Acc: 0.9951
val Loss: 0.2721 Acc: 0.9382
Epoch 125/150
-----
train Loss: 0.0305 Acc: 0.9931
val Loss: 0.2756 Acc: 0.9402
Epoch 126/150
-----
train Loss: 0.0394 Acc: 0.9902
```

```
val Loss: 0.2752 Acc: 0.9412
Epoch 127/150
-----
train Loss: 0.0406 Acc: 0.9892
val Loss: 0.2726 Acc: 0.9373
Epoch 128/150
-----
train Loss: 0.0438 Acc: 0.9902
val Loss: 0.2705 Acc: 0.9412
Epoch 129/150
-----
train Loss: 0.0354 Acc: 0.9922
val Loss: 0.2699 Acc: 0.9392
Epoch 130/150
-----
train Loss: 0.0255 Acc: 0.9941
val Loss: 0.2742 Acc: 0.9382
Epoch 131/150
-----
train Loss: 0.0281 Acc: 0.9951
val Loss: 0.2722 Acc: 0.9402
Epoch 132/150
-----
train Loss: 0.0441 Acc: 0.9892
val Loss: 0.2734 Acc: 0.9412
Epoch 133/150
-----
train Loss: 0.0332 Acc: 0.9922
val Loss: 0.2742 Acc: 0.9431
Epoch 134/150
-----
train Loss: 0.0324 Acc: 0.9922
val Loss: 0.2750 Acc: 0.9402
Epoch 135/150
-----
train Loss: 0.0319 Acc: 0.9931
val Loss: 0.2730 Acc: 0.9412
Epoch 136/150
-----
train Loss: 0.0176 Acc: 0.9961
val Loss: 0.2733 Acc: 0.9422
Epoch 137/150
-----
train Loss: 0.0312 Acc: 0.9912
val Loss: 0.2738 Acc: 0.9382
Epoch 138/150
-----
train Loss: 0.0324 Acc: 0.9941
val Loss: 0.2743 Acc: 0.9402
Epoch 139/150
-----
train Loss: 0.0363 Acc: 0.9931
val Loss: 0.2749 Acc: 0.9392
Epoch 140/150
-----
train Loss: 0.0277 Acc: 0.9931
```

```
val Loss: 0.2742 Acc: 0.9402
Epoch 141/150
-----
train Loss: 0.0473 Acc: 0.9902
val Loss: 0.2765 Acc: 0.9412
Epoch 142/150
-----
train Loss: 0.0313 Acc: 0.9922
val Loss: 0.2757 Acc: 0.9402
Epoch 143/150
-----
train Loss: 0.0332 Acc: 0.9922
val Loss: 0.2741 Acc: 0.9402
Epoch 144/150
-----
train Loss: 0.0237 Acc: 0.9941
val Loss: 0.2714 Acc: 0.9422
Epoch 145/150
-----
train Loss: 0.0295 Acc: 0.9922
val Loss: 0.2702 Acc: 0.9402
Epoch 146/150
-----
train Loss: 0.0300 Acc: 0.9941
val Loss: 0.2759 Acc: 0.9392
Epoch 147/150
-----
train Loss: 0.0504 Acc: 0.9863
val Loss: 0.2720 Acc: 0.9402
Epoch 148/150
-----
train Loss: 0.0349 Acc: 0.9902
val Loss: 0.2732 Acc: 0.9412
Epoch 149/150
-----
train Loss: 0.0256 Acc: 0.9961
val Loss: 0.2700 Acc: 0.9422
Epoch 150/150
-----
train Loss: 0.0270 Acc: 0.9951
val Loss: 0.2708 Acc: 0.9412
```

```
In [10]: print("\nFinetuning with lr=0.1")
model_1 = experiment_higher_lr(0.1)
```

```
Finetuning with lr=0.1
Epoch 1/150
-----
train Loss: 4.7264 Acc: 0.0765
val Loss: 3760.7262 Acc: 0.0167
Epoch 2/150
-----
train Loss: 5.7291 Acc: 0.0118
val Loss: 1017396.0828 Acc: 0.0098
Epoch 3/150
-----
train Loss: 4.9481 Acc: 0.0108
val Loss: 6386.8212 Acc: 0.0098
Epoch 4/150
-----
train Loss: 4.5879 Acc: 0.0255
val Loss: 27.1684 Acc: 0.0216
Epoch 5/150
-----
train Loss: 4.3518 Acc: 0.0333
val Loss: 4.6158 Acc: 0.0382
Epoch 6/150
-----
train Loss: 4.1694 Acc: 0.0441
val Loss: 4.2812 Acc: 0.0510
Epoch 7/150
-----
train Loss: 4.0406 Acc: 0.0569
val Loss: 4.1213 Acc: 0.0569
Epoch 8/150
-----
train Loss: 3.9326 Acc: 0.0696
val Loss: 3.9081 Acc: 0.0716
Epoch 9/150
-----
train Loss: 3.8186 Acc: 0.0794
val Loss: 3.9212 Acc: 0.0882
Epoch 10/150
-----
train Loss: 3.6803 Acc: 0.0892
val Loss: 3.8284 Acc: 0.0853
Epoch 11/150
-----
train Loss: 3.6034 Acc: 0.1127
val Loss: 3.8310 Acc: 0.0794
Epoch 12/150
-----
train Loss: 3.5572 Acc: 0.1216
val Loss: 4.0868 Acc: 0.0647
Epoch 13/150
-----
train Loss: 3.5735 Acc: 0.1206
val Loss: 3.7031 Acc: 0.1225
Epoch 14/150
-----
train Loss: 3.4352 Acc: 0.1206
```

```
val Loss: 3.9032 Acc: 0.1196
Epoch 15/150
-----
train Loss: 3.3575 Acc: 0.1412
val Loss: 3.6224 Acc: 0.1235
Epoch 16/150
-----
train Loss: 3.3390 Acc: 0.1500
val Loss: 3.5220 Acc: 0.1431
Epoch 17/150
-----
train Loss: 3.2617 Acc: 0.1637
val Loss: 3.4016 Acc: 0.1735
Epoch 18/150
-----
train Loss: 3.2140 Acc: 0.1608
val Loss: 3.5313 Acc: 0.1284
Epoch 19/150
-----
train Loss: 3.2296 Acc: 0.1784
val Loss: 3.3547 Acc: 0.1735
Epoch 20/150
-----
train Loss: 3.1249 Acc: 0.2000
val Loss: 3.4638 Acc: 0.1520
Epoch 21/150
-----
train Loss: 3.0882 Acc: 0.1990
val Loss: 3.9067 Acc: 0.1637
Epoch 22/150
-----
train Loss: 3.0380 Acc: 0.2137
val Loss: 3.2301 Acc: 0.1912
Epoch 23/150
-----
train Loss: 2.9749 Acc: 0.2196
val Loss: 3.2990 Acc: 0.1873
Epoch 24/150
-----
train Loss: 2.9573 Acc: 0.2255
val Loss: 3.4441 Acc: 0.1912
Epoch 25/150
-----
train Loss: 2.9136 Acc: 0.2373
val Loss: 3.6838 Acc: 0.2118
Epoch 26/150
-----
train Loss: 2.8272 Acc: 0.2569
val Loss: 3.1533 Acc: 0.2353
Epoch 27/150
-----
train Loss: 2.8089 Acc: 0.2784
val Loss: 3.7812 Acc: 0.1676
Epoch 28/150
-----
train Loss: 2.7853 Acc: 0.2735
```

```
val Loss: 3.5965 Acc: 0.2049
Epoch 29/150
-----
train Loss: 2.6653 Acc: 0.2745
val Loss: 3.5387 Acc: 0.2020
Epoch 30/150
-----
train Loss: 2.6550 Acc: 0.2951
val Loss: 3.2534 Acc: 0.2412
Epoch 31/150
-----
train Loss: 2.6079 Acc: 0.3039
val Loss: 3.2079 Acc: 0.2500
Epoch 32/150
-----
train Loss: 2.5703 Acc: 0.2941
val Loss: 3.2632 Acc: 0.2578
Epoch 33/150
-----
train Loss: 2.7205 Acc: 0.2618
val Loss: 46.2073 Acc: 0.1118
Epoch 34/150
-----
train Loss: 2.5427 Acc: 0.3147
val Loss: 2.9253 Acc: 0.2637
Epoch 35/150
-----
train Loss: 2.4884 Acc: 0.3275
val Loss: 3.5031 Acc: 0.2304
Epoch 36/150
-----
train Loss: 2.3848 Acc: 0.3402
val Loss: 3.1617 Acc: 0.2647
Epoch 37/150
-----
train Loss: 2.3211 Acc: 0.3529
val Loss: 2.9899 Acc: 0.2853
Epoch 38/150
-----
train Loss: 2.2896 Acc: 0.3637
val Loss: 3.8025 Acc: 0.2461
Epoch 39/150
-----
train Loss: 2.2393 Acc: 0.3745
val Loss: 3.4526 Acc: 0.2922
Epoch 40/150
-----
train Loss: 2.1555 Acc: 0.4108
val Loss: 3.0165 Acc: 0.2833
Epoch 41/150
-----
train Loss: 1.9729 Acc: 0.4304
val Loss: 2.7876 Acc: 0.3343
Epoch 42/150
-----
train Loss: 1.8042 Acc: 0.4775
```

```
val Loss: 2.7407 Acc: 0.3637
Epoch 43/150
-----
train Loss: 1.6619 Acc: 0.5363
val Loss: 2.7115 Acc: 0.3657
Epoch 44/150
-----
train Loss: 1.5925 Acc: 0.5578
val Loss: 2.7262 Acc: 0.3627
Epoch 45/150
-----
train Loss: 1.5984 Acc: 0.5431
val Loss: 2.7974 Acc: 0.3735
Epoch 46/150
-----
train Loss: 1.6135 Acc: 0.5431
val Loss: 2.8545 Acc: 0.3706
Epoch 47/150
-----
train Loss: 1.5727 Acc: 0.5529
val Loss: 2.8300 Acc: 0.3686
Epoch 48/150
-----
train Loss: 1.5340 Acc: 0.5637
val Loss: 2.8749 Acc: 0.3775
Epoch 49/150
-----
train Loss: 1.5383 Acc: 0.5529
val Loss: 2.8251 Acc: 0.3814
Epoch 50/150
-----
train Loss: 1.4294 Acc: 0.5882
val Loss: 2.8232 Acc: 0.3873
Epoch 51/150
-----
train Loss: 1.4639 Acc: 0.5961
val Loss: 2.9006 Acc: 0.3863
Epoch 52/150
-----
train Loss: 1.4343 Acc: 0.5912
val Loss: 2.8684 Acc: 0.3804
Epoch 53/150
-----
train Loss: 1.4496 Acc: 0.5794
val Loss: 2.9364 Acc: 0.3853
Epoch 54/150
-----
train Loss: 1.4068 Acc: 0.5833
val Loss: 2.8503 Acc: 0.3843
Epoch 55/150
-----
train Loss: 1.2993 Acc: 0.6235
val Loss: 2.9076 Acc: 0.3814
Epoch 56/150
-----
train Loss: 1.3856 Acc: 0.5931
```

```
val Loss: 2.9429 Acc: 0.3882
Epoch 57/150
-----
train Loss: 1.4159 Acc: 0.6029
val Loss: 2.9671 Acc: 0.3961
Epoch 58/150
-----
train Loss: 1.3094 Acc: 0.6235
val Loss: 2.9802 Acc: 0.3804
Epoch 59/150
-----
train Loss: 1.2983 Acc: 0.6098
val Loss: 2.9835 Acc: 0.3882
Epoch 60/150
-----
train Loss: 1.2766 Acc: 0.6284
val Loss: 2.9282 Acc: 0.3765
Epoch 61/150
-----
train Loss: 1.2516 Acc: 0.6353
val Loss: 2.9120 Acc: 0.3990
Epoch 62/150
-----
train Loss: 1.2862 Acc: 0.6137
val Loss: 3.0593 Acc: 0.3980
Epoch 63/150
-----
train Loss: 1.2061 Acc: 0.6569
val Loss: 2.9054 Acc: 0.3922
Epoch 64/150
-----
train Loss: 1.1674 Acc: 0.6637
val Loss: 2.9688 Acc: 0.4098
Epoch 65/150
-----
train Loss: 1.2802 Acc: 0.6235
val Loss: 2.9211 Acc: 0.3833
Epoch 66/150
-----
train Loss: 1.1966 Acc: 0.6608
val Loss: 2.9573 Acc: 0.4069
Epoch 67/150
-----
train Loss: 1.1196 Acc: 0.6912
val Loss: 3.0182 Acc: 0.3922
Epoch 68/150
-----
train Loss: 1.1884 Acc: 0.6441
val Loss: 2.9596 Acc: 0.4059
Epoch 69/150
-----
train Loss: 1.1607 Acc: 0.6696
val Loss: 2.9884 Acc: 0.3912
Epoch 70/150
-----
train Loss: 1.1410 Acc: 0.6735
```

```
val Loss: 3.0964 Acc: 0.3951
Epoch 71/150
-----
train Loss: 1.0936 Acc: 0.6637
val Loss: 2.8854 Acc: 0.3892
Epoch 72/150
-----
train Loss: 1.0903 Acc: 0.6745
val Loss: 2.9432 Acc: 0.4059
Epoch 73/150
-----
train Loss: 1.1310 Acc: 0.6676
val Loss: 3.0003 Acc: 0.4157
Epoch 74/150
-----
train Loss: 1.0501 Acc: 0.7049
val Loss: 2.9207 Acc: 0.4000
Epoch 75/150
-----
train Loss: 1.0280 Acc: 0.6902
val Loss: 2.8991 Acc: 0.4118
Epoch 76/150
-----
train Loss: 1.1132 Acc: 0.6735
val Loss: 2.9038 Acc: 0.4010
Epoch 77/150
-----
train Loss: 0.9816 Acc: 0.7157
val Loss: 2.9383 Acc: 0.4196
Epoch 78/150
-----
train Loss: 0.9548 Acc: 0.7206
val Loss: 2.9382 Acc: 0.4196
Epoch 79/150
-----
train Loss: 0.9368 Acc: 0.7196
val Loss: 2.9200 Acc: 0.4216
Epoch 80/150
-----
train Loss: 0.9415 Acc: 0.7392
val Loss: 2.9001 Acc: 0.4216
Epoch 81/150
-----
train Loss: 0.9798 Acc: 0.7235
val Loss: 2.8574 Acc: 0.4216
Epoch 82/150
-----
train Loss: 0.8399 Acc: 0.7745
val Loss: 2.8416 Acc: 0.4196
Epoch 83/150
-----
train Loss: 0.7577 Acc: 0.7892
val Loss: 2.8372 Acc: 0.4225
Epoch 84/150
-----
train Loss: 0.8520 Acc: 0.7657
```

```
val Loss: 2.8262 Acc: 0.4225
Epoch 85/150
-----
train Loss: 0.7890 Acc: 0.7853
val Loss: 2.8479 Acc: 0.4196
Epoch 86/150
-----
train Loss: 0.7699 Acc: 0.8029
val Loss: 2.8251 Acc: 0.4275
Epoch 87/150
-----
train Loss: 0.7879 Acc: 0.7745
val Loss: 2.8196 Acc: 0.4265
Epoch 88/150
-----
train Loss: 0.8300 Acc: 0.7676
val Loss: 2.8544 Acc: 0.4275
Epoch 89/150
-----
train Loss: 0.7774 Acc: 0.7931
val Loss: 2.8259 Acc: 0.4275
Epoch 90/150
-----
train Loss: 0.8277 Acc: 0.7598
val Loss: 2.8298 Acc: 0.4314
Epoch 91/150
-----
train Loss: 0.7822 Acc: 0.7941
val Loss: 2.8205 Acc: 0.4324
Epoch 92/150
-----
train Loss: 0.8218 Acc: 0.7686
val Loss: 2.8444 Acc: 0.4225
Epoch 93/150
-----
train Loss: 0.7704 Acc: 0.8098
val Loss: 2.8401 Acc: 0.4284
Epoch 94/150
-----
train Loss: 0.7581 Acc: 0.7980
val Loss: 2.8188 Acc: 0.4353
Epoch 95/150
-----
train Loss: 0.7574 Acc: 0.8010
val Loss: 2.8481 Acc: 0.4294
Epoch 96/150
-----
train Loss: 0.8101 Acc: 0.7814
val Loss: 2.8343 Acc: 0.4284
Epoch 97/150
-----
train Loss: 0.7904 Acc: 0.7725
val Loss: 2.8192 Acc: 0.4235
Epoch 98/150
-----
train Loss: 0.7920 Acc: 0.7745
```

```
val Loss: 2.8275 Acc: 0.4314
Epoch 99/150
-----
train Loss: 0.8010 Acc: 0.7647
val Loss: 2.8524 Acc: 0.4314
Epoch 100/150
-----
train Loss: 0.7022 Acc: 0.8176
val Loss: 2.8231 Acc: 0.4284
Epoch 101/150
-----
train Loss: 0.7544 Acc: 0.7961
val Loss: 2.8151 Acc: 0.4373
Epoch 102/150
-----
train Loss: 0.7706 Acc: 0.7922
val Loss: 2.8157 Acc: 0.4324
Epoch 103/150
-----
train Loss: 0.7427 Acc: 0.7980
val Loss: 2.8239 Acc: 0.4284
Epoch 104/150
-----
train Loss: 0.7533 Acc: 0.7863
val Loss: 2.8310 Acc: 0.4353
Epoch 105/150
-----
train Loss: 0.7084 Acc: 0.8088
val Loss: 2.8097 Acc: 0.4304
Epoch 106/150
-----
train Loss: 0.7690 Acc: 0.7863
val Loss: 2.8347 Acc: 0.4294
Epoch 107/150
-----
train Loss: 0.7322 Acc: 0.7853
val Loss: 2.8305 Acc: 0.4363
Epoch 108/150
-----
train Loss: 0.7502 Acc: 0.8000
val Loss: 2.8322 Acc: 0.4304
Epoch 109/150
-----
train Loss: 0.7035 Acc: 0.8157
val Loss: 2.8321 Acc: 0.4353
Epoch 110/150
-----
train Loss: 0.7964 Acc: 0.8010
val Loss: 2.8377 Acc: 0.4235
Epoch 111/150
-----
train Loss: 0.7037 Acc: 0.8078
val Loss: 2.8635 Acc: 0.4373
Epoch 112/150
-----
train Loss: 0.7081 Acc: 0.8039
```

```
val Loss: 2.8757 Acc: 0.4363
Epoch 113/150
-----
train Loss: 0.6967 Acc: 0.8167
val Loss: 2.8501 Acc: 0.4284
Epoch 114/150
-----
train Loss: 0.7270 Acc: 0.8108
val Loss: 2.8343 Acc: 0.4333
Epoch 115/150
-----
train Loss: 0.8137 Acc: 0.7873
val Loss: 2.8221 Acc: 0.4373
Epoch 116/150
-----
train Loss: 0.6855 Acc: 0.8098
val Loss: 2.8341 Acc: 0.4422
Epoch 117/150
-----
train Loss: 0.7411 Acc: 0.8059
val Loss: 2.8340 Acc: 0.4392
Epoch 118/150
-----
train Loss: 0.7250 Acc: 0.8020
val Loss: 2.8265 Acc: 0.4294
Epoch 119/150
-----
train Loss: 0.7816 Acc: 0.7922
val Loss: 2.8273 Acc: 0.4314
Epoch 120/150
-----
train Loss: 0.7078 Acc: 0.8186
val Loss: 2.8512 Acc: 0.4284
Epoch 121/150
-----
train Loss: 0.6963 Acc: 0.8147
val Loss: 2.8409 Acc: 0.4314
Epoch 122/150
-----
train Loss: 0.7259 Acc: 0.8020
val Loss: 2.8311 Acc: 0.4324
Epoch 123/150
-----
train Loss: 0.7594 Acc: 0.8020
val Loss: 2.8337 Acc: 0.4314
Epoch 124/150
-----
train Loss: 0.7806 Acc: 0.7882
val Loss: 2.8219 Acc: 0.4333
Epoch 125/150
-----
train Loss: 0.7553 Acc: 0.7990
val Loss: 2.8372 Acc: 0.4343
Epoch 126/150
-----
train Loss: 0.6552 Acc: 0.8461
```

```
val Loss: 2.8493 Acc: 0.4333
Epoch 127/150
-----
train Loss: 0.7684 Acc: 0.7794
val Loss: 2.8439 Acc: 0.4284
Epoch 128/150
-----
train Loss: 0.6341 Acc: 0.8373
val Loss: 2.8404 Acc: 0.4304
Epoch 129/150
-----
train Loss: 0.7083 Acc: 0.8078
val Loss: 2.8503 Acc: 0.4284
Epoch 130/150
-----
train Loss: 0.6430 Acc: 0.8235
val Loss: 2.8219 Acc: 0.4294
Epoch 131/150
-----
train Loss: 0.6774 Acc: 0.8235
val Loss: 2.8467 Acc: 0.4275
Epoch 132/150
-----
train Loss: 0.6861 Acc: 0.8225
val Loss: 2.8399 Acc: 0.4275
Epoch 133/150
-----
train Loss: 0.6685 Acc: 0.8304
val Loss: 2.8370 Acc: 0.4314
Epoch 134/150
-----
train Loss: 0.6590 Acc: 0.8441
val Loss: 2.8382 Acc: 0.4314
Epoch 135/150
-----
train Loss: 0.6737 Acc: 0.8235
val Loss: 2.8238 Acc: 0.4382
Epoch 136/150
-----
train Loss: 0.6452 Acc: 0.8392
val Loss: 2.8332 Acc: 0.4353
Epoch 137/150
-----
train Loss: 0.7460 Acc: 0.7980
val Loss: 2.8417 Acc: 0.4333
Epoch 138/150
-----
train Loss: 0.6439 Acc: 0.8206
val Loss: 2.8334 Acc: 0.4324
Epoch 139/150
-----
train Loss: 0.6551 Acc: 0.8216
val Loss: 2.8264 Acc: 0.4324
Epoch 140/150
-----
train Loss: 0.7537 Acc: 0.7980
```

```
val Loss: 2.8393 Acc: 0.4304
Epoch 141/150
-----
train Loss: 0.6478 Acc: 0.8353
val Loss: 2.8348 Acc: 0.4284
Epoch 142/150
-----
train Loss: 0.6921 Acc: 0.8098
val Loss: 2.8309 Acc: 0.4402
Epoch 143/150
-----
train Loss: 0.7110 Acc: 0.8245
val Loss: 2.8341 Acc: 0.4402
Epoch 144/150
-----
train Loss: 0.6888 Acc: 0.8206
val Loss: 2.8477 Acc: 0.4314
Epoch 145/150
-----
train Loss: 0.6814 Acc: 0.8098
val Loss: 2.8456 Acc: 0.4353
Epoch 146/150
-----
train Loss: 0.6463 Acc: 0.8275
val Loss: 2.8017 Acc: 0.4333
Epoch 147/150
-----
train Loss: 0.6881 Acc: 0.8255
val Loss: 2.8248 Acc: 0.4343
Epoch 148/150
-----
train Loss: 0.6840 Acc: 0.8245
val Loss: 2.8247 Acc: 0.4363
Epoch 149/150
-----
train Loss: 0.6735 Acc: 0.8255
val Loss: 2.8248 Acc: 0.4314
Epoch 150/150
-----
train Loss: 0.6992 Acc: 0.8186
val Loss: 2.8061 Acc: 0.4363
```

2. Using the model as a feature extractor

When using a pretrained model as a feature extractor, all the layers of the network are frozen except for the final layer. Thus, except for the last layer, none of the inner layers' gradients are updated during the backward pass with the target dataset. Since gradients do not need to be computed for most of the network, this is faster than finetuning.

a. Train only the last layer

Now train only the last layer for 0.1, 0.01, and 0.001 learning rates while keeping all the other hyperparameters and settings the same as for finetuning. Which learning rate gives you the best accuracy on the target dataset?
(5 points)

In [11]:

```
# TODO
def feature_extractor_experiment(learning_rate, root='./data', num_epochs=15):
    # Define transformations
    data_transforms = {
        'train': transforms.Compose([
            transforms.RandomResizedCrop(224),
            transforms.RandomHorizontalFlip(),
            transforms.ToTensor(),
            transforms.Normalize([0.485, 0.456, 0.406], [0.229, 0.224, 0.225])
        ]),
        'val': transforms.Compose([
            transforms.Resize(256),
            transforms.CenterCrop(224),
            transforms.ToTensor(),
            transforms.Normalize([0.485, 0.456, 0.406], [0.229, 0.224, 0.225])
        ])
    }

    # Load Flowers102 dataset
    train_dataset = torchvision.datasets.Flowers102(root=root,
                                                    split='train',
                                                    transform=data_transforms['train'],
                                                    download=True)

    val_dataset = torchvision.datasets.Flowers102(root=root,
                                                split='val',
                                                transform=data_transforms['val'],
                                                download=True)

    # Create dataloaders
    dataloaders = {
        'train': DataLoader(train_dataset, batch_size=64, shuffle=True, num_workers=4),
        'val': DataLoader(val_dataset, batch_size=64, shuffle=False, num_workers=4)
    }

    # Load pre-trained ResNet50
    model = torchvision.models.resnet50(pretrained=True)

    # Freeze all layers except the final one
    for param in model.parameters():
        param.requires_grad = False

    # Modify and train only the last layer
    num_ftrs = model.fc.in_features
    model.fc = nn.Linear(num_ftrs, 102)

    # Define loss function and optimizer (only for the last layer)
    criterion = nn.CrossEntropyLoss()
    optimizer = optim.SGD(model.fc.parameters(), lr=learning_rate, momentum=0.9)
```

```
# Define learning rate scheduler
scheduler = optim.lr_scheduler.MultiStepLR(optimizer, milestones=[40, 80])

# Train the model
model = train_model(model, criterion, optimizer, scheduler, dataloaders,
                     return model)
```

```
In [12]: print("\nFeature extraction with different learning rates")
for lr in [0.001, 0.01, 0.1]:
    print(f"\nLearning rate: {lr}")
    model_fe = feature_extractor_experiment(lr)
```

Feature extraction with different learning rates

```
Learning rate: 0.001
Epoch 1/150
-----
train Loss: 4.6756 Acc: 0.0098
val Loss: 4.6186 Acc: 0.0167
Epoch 2/150
-----
train Loss: 4.5683 Acc: 0.0284
val Loss: 4.4988 Acc: 0.0353
Epoch 3/150
-----
train Loss: 4.4556 Acc: 0.0775
val Loss: 4.3888 Acc: 0.1127
Epoch 4/150
-----
train Loss: 4.3577 Acc: 0.1588
val Loss: 4.2843 Acc: 0.1980
Epoch 5/150
-----
train Loss: 4.2564 Acc: 0.2363
val Loss: 4.1798 Acc: 0.2647
Epoch 6/150
-----
train Loss: 4.1531 Acc: 0.3510
val Loss: 4.0793 Acc: 0.3520
Epoch 7/150
-----
train Loss: 4.0624 Acc: 0.3892
val Loss: 3.9809 Acc: 0.4245
Epoch 8/150
-----
train Loss: 3.9661 Acc: 0.4824
val Loss: 3.8796 Acc: 0.4873
Epoch 9/150
-----
train Loss: 3.8676 Acc: 0.5255
val Loss: 3.7823 Acc: 0.5324
Epoch 10/150
-----
train Loss: 3.7739 Acc: 0.6049
val Loss: 3.6916 Acc: 0.5735
Epoch 11/150
-----
train Loss: 3.6873 Acc: 0.6392
val Loss: 3.5979 Acc: 0.6039
Epoch 12/150
-----
train Loss: 3.5900 Acc: 0.6539
val Loss: 3.4999 Acc: 0.6333
Epoch 13/150
-----
train Loss: 3.5188 Acc: 0.6716
val Loss: 3.4188 Acc: 0.6578
Epoch 14/150
```

```
-----  
train Loss: 3.4175 Acc: 0.7088  
val Loss: 3.3358 Acc: 0.6784  
Epoch 15/150  
-----  
train Loss: 3.3350 Acc: 0.7137  
val Loss: 3.2519 Acc: 0.6990  
Epoch 16/150  
-----  
train Loss: 3.2485 Acc: 0.7490  
val Loss: 3.1718 Acc: 0.7059  
Epoch 17/150  
-----  
train Loss: 3.1653 Acc: 0.7500  
val Loss: 3.0919 Acc: 0.7275  
Epoch 18/150  
-----  
train Loss: 3.1058 Acc: 0.7745  
val Loss: 3.0141 Acc: 0.7480  
Epoch 19/150  
-----  
train Loss: 3.0382 Acc: 0.7765  
val Loss: 2.9436 Acc: 0.7510  
Epoch 20/150  
-----  
train Loss: 2.9352 Acc: 0.7951  
val Loss: 2.8754 Acc: 0.7529  
Epoch 21/150  
-----  
train Loss: 2.8802 Acc: 0.7794  
val Loss: 2.8052 Acc: 0.7676  
Epoch 22/150  
-----  
train Loss: 2.8307 Acc: 0.8010  
val Loss: 2.7365 Acc: 0.7676  
Epoch 23/150  
-----  
train Loss: 2.7365 Acc: 0.8294  
val Loss: 2.6770 Acc: 0.7765  
Epoch 24/150  
-----  
train Loss: 2.6763 Acc: 0.8255  
val Loss: 2.6074 Acc: 0.7902  
Epoch 25/150  
-----  
train Loss: 2.6034 Acc: 0.8363  
val Loss: 2.5510 Acc: 0.7902  
Epoch 26/150  
-----  
train Loss: 2.5299 Acc: 0.8461  
val Loss: 2.4868 Acc: 0.7951  
Epoch 27/150  
-----  
train Loss: 2.4847 Acc: 0.8363  
val Loss: 2.4355 Acc: 0.8069  
Epoch 28/150
```

```
-----  
train Loss: 2.4608 Acc: 0.8314  
val Loss: 2.3735 Acc: 0.8118  
Epoch 29/150  
-----  
train Loss: 2.3654 Acc: 0.8539  
val Loss: 2.3269 Acc: 0.8108  
Epoch 30/150  
-----  
train Loss: 2.3198 Acc: 0.8520  
val Loss: 2.2820 Acc: 0.8147  
Epoch 31/150  
-----  
train Loss: 2.2539 Acc: 0.8696  
val Loss: 2.2335 Acc: 0.8206  
Epoch 32/150  
-----  
train Loss: 2.2378 Acc: 0.8676  
val Loss: 2.1827 Acc: 0.8245  
Epoch 33/150  
-----  
train Loss: 2.1679 Acc: 0.8696  
val Loss: 2.1508 Acc: 0.8216  
Epoch 34/150  
-----  
train Loss: 2.1210 Acc: 0.8490  
val Loss: 2.0967 Acc: 0.8275  
Epoch 35/150  
-----  
train Loss: 2.1102 Acc: 0.8716  
val Loss: 2.0440 Acc: 0.8294  
Epoch 36/150  
-----  
train Loss: 2.0618 Acc: 0.8588  
val Loss: 2.0088 Acc: 0.8196  
Epoch 37/150  
-----  
train Loss: 2.0042 Acc: 0.8706  
val Loss: 1.9609 Acc: 0.8353  
Epoch 38/150  
-----  
train Loss: 1.9913 Acc: 0.8696  
val Loss: 1.9379 Acc: 0.8304  
Epoch 39/150  
-----  
train Loss: 1.9082 Acc: 0.8804  
val Loss: 1.8902 Acc: 0.8304  
Epoch 40/150  
-----  
train Loss: 1.8777 Acc: 0.8706  
val Loss: 1.8640 Acc: 0.8392  
Epoch 41/150  
-----  
train Loss: 1.8537 Acc: 0.8716  
val Loss: 1.8681 Acc: 0.8402  
Epoch 42/150
```

```
-----  
train Loss: 1.8457 Acc: 0.8882  
val Loss: 1.8604 Acc: 0.8382  
Epoch 43/150  
-----  
train Loss: 1.8657 Acc: 0.8765  
val Loss: 1.8531 Acc: 0.8402  
Epoch 44/150  
-----  
train Loss: 1.8509 Acc: 0.8794  
val Loss: 1.8422 Acc: 0.8422  
Epoch 45/150  
-----  
train Loss: 1.8643 Acc: 0.8794  
val Loss: 1.8421 Acc: 0.8412  
Epoch 46/150  
-----  
train Loss: 1.8334 Acc: 0.8804  
val Loss: 1.8446 Acc: 0.8402  
Epoch 47/150  
-----  
train Loss: 1.8686 Acc: 0.8784  
val Loss: 1.8330 Acc: 0.8422  
Epoch 48/150  
-----  
train Loss: 1.7995 Acc: 0.8941  
val Loss: 1.8462 Acc: 0.8392  
Epoch 49/150  
-----  
train Loss: 1.8317 Acc: 0.8843  
val Loss: 1.8286 Acc: 0.8422  
Epoch 50/150  
-----  
train Loss: 1.7994 Acc: 0.9029  
val Loss: 1.8199 Acc: 0.8451  
Epoch 51/150  
-----  
train Loss: 1.8197 Acc: 0.8843  
val Loss: 1.8241 Acc: 0.8451  
Epoch 52/150  
-----  
train Loss: 1.8104 Acc: 0.8892  
val Loss: 1.8205 Acc: 0.8431  
Epoch 53/150  
-----  
train Loss: 1.7940 Acc: 0.8931  
val Loss: 1.8181 Acc: 0.8441  
Epoch 54/150  
-----  
train Loss: 1.8158 Acc: 0.8716  
val Loss: 1.8072 Acc: 0.8422  
Epoch 55/150  
-----  
train Loss: 1.8036 Acc: 0.8755  
val Loss: 1.8182 Acc: 0.8402  
Epoch 56/150
```

```
-----  
train Loss: 1.8038 Acc: 0.8882  
val Loss: 1.8124 Acc: 0.8422  
Epoch 57/150  
-----  
train Loss: 1.7481 Acc: 0.9059  
val Loss: 1.8125 Acc: 0.8382  
Epoch 58/150  
-----  
train Loss: 1.7879 Acc: 0.8902  
val Loss: 1.8138 Acc: 0.8392  
Epoch 59/150  
-----  
train Loss: 1.7944 Acc: 0.8833  
val Loss: 1.8011 Acc: 0.8431  
Epoch 60/150  
-----  
train Loss: 1.7926 Acc: 0.8863  
val Loss: 1.8011 Acc: 0.8441  
Epoch 61/150  
-----  
train Loss: 1.7904 Acc: 0.8853  
val Loss: 1.8041 Acc: 0.8461  
Epoch 62/150  
-----  
train Loss: 1.7750 Acc: 0.8941  
val Loss: 1.7901 Acc: 0.8451  
Epoch 63/150  
-----  
train Loss: 1.7538 Acc: 0.8951  
val Loss: 1.7867 Acc: 0.8441  
Epoch 64/150  
-----  
train Loss: 1.7575 Acc: 0.8951  
val Loss: 1.7867 Acc: 0.8431  
Epoch 65/150  
-----  
train Loss: 1.7436 Acc: 0.9000  
val Loss: 1.7791 Acc: 0.8461  
Epoch 66/150  
-----  
train Loss: 1.7398 Acc: 0.8971  
val Loss: 1.7689 Acc: 0.8480  
Epoch 67/150  
-----  
train Loss: 1.7402 Acc: 0.8863  
val Loss: 1.7640 Acc: 0.8471  
Epoch 68/150  
-----  
train Loss: 1.7487 Acc: 0.8990  
val Loss: 1.7670 Acc: 0.8480  
Epoch 69/150  
-----  
train Loss: 1.7499 Acc: 0.8951  
val Loss: 1.7725 Acc: 0.8451  
Epoch 70/150
```

```
-----  
train Loss: 1.7046 Acc: 0.9078  
val Loss: 1.7676 Acc: 0.8441  
Epoch 71/150  
-----  
train Loss: 1.7415 Acc: 0.9010  
val Loss: 1.7632 Acc: 0.8422  
Epoch 72/150  
-----  
train Loss: 1.7734 Acc: 0.8775  
val Loss: 1.7622 Acc: 0.8461  
Epoch 73/150  
-----  
train Loss: 1.7712 Acc: 0.8725  
val Loss: 1.7512 Acc: 0.8431  
Epoch 74/150  
-----  
train Loss: 1.7615 Acc: 0.8912  
val Loss: 1.7465 Acc: 0.8451  
Epoch 75/150  
-----  
train Loss: 1.7472 Acc: 0.8814  
val Loss: 1.7457 Acc: 0.8451  
Epoch 76/150  
-----  
train Loss: 1.7493 Acc: 0.8824  
val Loss: 1.7478 Acc: 0.8471  
Epoch 77/150  
-----  
train Loss: 1.7234 Acc: 0.8941  
val Loss: 1.7473 Acc: 0.8441  
Epoch 78/150  
-----  
train Loss: 1.7038 Acc: 0.9020  
val Loss: 1.7416 Acc: 0.8451  
Epoch 79/150  
-----  
train Loss: 1.6993 Acc: 0.8941  
val Loss: 1.7379 Acc: 0.8441  
Epoch 80/150  
-----  
train Loss: 1.7655 Acc: 0.8814  
val Loss: 1.7299 Acc: 0.8471  
Epoch 81/150  
-----  
train Loss: 1.7338 Acc: 0.8824  
val Loss: 1.7256 Acc: 0.8480  
Epoch 82/150  
-----  
train Loss: 1.7164 Acc: 0.8951  
val Loss: 1.7363 Acc: 0.8431  
Epoch 83/150  
-----  
train Loss: 1.7115 Acc: 0.8794  
val Loss: 1.7363 Acc: 0.8451  
Epoch 84/150
```

```
-----  
train Loss: 1.6964 Acc: 0.9000  
val Loss: 1.7292 Acc: 0.8451  
Epoch 85/150  
-----  
train Loss: 1.7149 Acc: 0.8931  
val Loss: 1.7397 Acc: 0.8451  
Epoch 86/150  
-----  
train Loss: 1.7181 Acc: 0.8931  
val Loss: 1.7398 Acc: 0.8461  
Epoch 87/150  
-----  
train Loss: 1.6772 Acc: 0.8912  
val Loss: 1.7353 Acc: 0.8480  
Epoch 88/150  
-----  
train Loss: 1.7244 Acc: 0.8853  
val Loss: 1.7259 Acc: 0.8471  
Epoch 89/150  
-----  
train Loss: 1.7205 Acc: 0.8931  
val Loss: 1.7351 Acc: 0.8490  
Epoch 90/150  
-----  
train Loss: 1.7305 Acc: 0.8794  
val Loss: 1.7308 Acc: 0.8480  
Epoch 91/150  
-----  
train Loss: 1.6930 Acc: 0.8912  
val Loss: 1.7326 Acc: 0.8461  
Epoch 92/150  
-----  
train Loss: 1.7026 Acc: 0.8853  
val Loss: 1.7356 Acc: 0.8451  
Epoch 93/150  
-----  
train Loss: 1.7133 Acc: 0.8902  
val Loss: 1.7287 Acc: 0.8451  
Epoch 94/150  
-----  
train Loss: 1.7292 Acc: 0.8794  
val Loss: 1.7347 Acc: 0.8480  
Epoch 95/150  
-----  
train Loss: 1.7357 Acc: 0.8824  
val Loss: 1.7383 Acc: 0.8480  
Epoch 96/150  
-----  
train Loss: 1.6923 Acc: 0.9000  
val Loss: 1.7189 Acc: 0.8471  
Epoch 97/150  
-----  
train Loss: 1.6803 Acc: 0.8941  
val Loss: 1.7256 Acc: 0.8451  
Epoch 98/150
```

```
-----  
train Loss: 1.6964 Acc: 0.8922  
val Loss: 1.7295 Acc: 0.8490  
Epoch 99/150  
-----  
train Loss: 1.7041 Acc: 0.8882  
val Loss: 1.7295 Acc: 0.8451  
Epoch 100/150  
-----  
train Loss: 1.7104 Acc: 0.8853  
val Loss: 1.7247 Acc: 0.8471  
Epoch 101/150  
-----  
train Loss: 1.6828 Acc: 0.8990  
val Loss: 1.7281 Acc: 0.8431  
Epoch 102/150  
-----  
train Loss: 1.7355 Acc: 0.8775  
val Loss: 1.7272 Acc: 0.8490  
Epoch 103/150  
-----  
train Loss: 1.7186 Acc: 0.8980  
val Loss: 1.7352 Acc: 0.8490  
Epoch 104/150  
-----  
train Loss: 1.7433 Acc: 0.8902  
val Loss: 1.7384 Acc: 0.8480  
Epoch 105/150  
-----  
train Loss: 1.7176 Acc: 0.8873  
val Loss: 1.7190 Acc: 0.8461  
Epoch 106/150  
-----  
train Loss: 1.7074 Acc: 0.8853  
val Loss: 1.7238 Acc: 0.8431  
Epoch 107/150  
-----  
train Loss: 1.7072 Acc: 0.8892  
val Loss: 1.7238 Acc: 0.8490  
Epoch 108/150  
-----  
train Loss: 1.7344 Acc: 0.8941  
val Loss: 1.7282 Acc: 0.8441  
Epoch 109/150  
-----  
train Loss: 1.7145 Acc: 0.8971  
val Loss: 1.7245 Acc: 0.8451  
Epoch 110/150  
-----  
train Loss: 1.6813 Acc: 0.9029  
val Loss: 1.7287 Acc: 0.8490  
Epoch 111/150  
-----  
train Loss: 1.6792 Acc: 0.8980  
val Loss: 1.7312 Acc: 0.8480  
Epoch 112/150
```

```
-----  
train Loss: 1.7268 Acc: 0.8765  
val Loss: 1.7238 Acc: 0.8490  
Epoch 113/150  
-----  
train Loss: 1.7046 Acc: 0.8951  
val Loss: 1.7177 Acc: 0.8480  
Epoch 114/150  
-----  
train Loss: 1.7115 Acc: 0.8941  
val Loss: 1.7146 Acc: 0.8480  
Epoch 115/150  
-----  
train Loss: 1.7050 Acc: 0.8902  
val Loss: 1.7156 Acc: 0.8480  
Epoch 116/150  
-----  
train Loss: 1.6740 Acc: 0.9010  
val Loss: 1.7220 Acc: 0.8471  
Epoch 117/150  
-----  
train Loss: 1.6935 Acc: 0.8892  
val Loss: 1.7152 Acc: 0.8461  
Epoch 118/150  
-----  
train Loss: 1.7026 Acc: 0.8980  
val Loss: 1.7238 Acc: 0.8480  
Epoch 119/150  
-----  
train Loss: 1.7029 Acc: 0.8990  
val Loss: 1.7289 Acc: 0.8451  
Epoch 120/150  
-----  
train Loss: 1.7207 Acc: 0.8951  
val Loss: 1.7340 Acc: 0.8441  
Epoch 121/150  
-----  
train Loss: 1.7025 Acc: 0.9118  
val Loss: 1.7332 Acc: 0.8461  
Epoch 122/150  
-----  
train Loss: 1.7000 Acc: 0.8922  
val Loss: 1.7192 Acc: 0.8461  
Epoch 123/150  
-----  
train Loss: 1.7225 Acc: 0.8892  
val Loss: 1.7226 Acc: 0.8461  
Epoch 124/150  
-----  
train Loss: 1.6983 Acc: 0.8902  
val Loss: 1.7235 Acc: 0.8480  
Epoch 125/150  
-----  
train Loss: 1.7276 Acc: 0.8804  
val Loss: 1.7187 Acc: 0.8451  
Epoch 126/150
```

```
-----  
train Loss: 1.7108 Acc: 0.9069  
val Loss: 1.7249 Acc: 0.8451  
Epoch 127/150  
-----  
train Loss: 1.6805 Acc: 0.8951  
val Loss: 1.7082 Acc: 0.8461  
Epoch 128/150  
-----  
train Loss: 1.7185 Acc: 0.8902  
val Loss: 1.7176 Acc: 0.8471  
Epoch 129/150  
-----  
train Loss: 1.7096 Acc: 0.8990  
val Loss: 1.7300 Acc: 0.8451  
Epoch 130/150  
-----  
train Loss: 1.7097 Acc: 0.8882  
val Loss: 1.7235 Acc: 0.8490  
Epoch 131/150  
-----  
train Loss: 1.7061 Acc: 0.8863  
val Loss: 1.7251 Acc: 0.8480  
Epoch 132/150  
-----  
train Loss: 1.6892 Acc: 0.8912  
val Loss: 1.7309 Acc: 0.8480  
Epoch 133/150  
-----  
train Loss: 1.7009 Acc: 0.9049  
val Loss: 1.7296 Acc: 0.8490  
Epoch 134/150  
-----  
train Loss: 1.6777 Acc: 0.9010  
val Loss: 1.7326 Acc: 0.8441  
Epoch 135/150  
-----  
train Loss: 1.7003 Acc: 0.8902  
val Loss: 1.7233 Acc: 0.8461  
Epoch 136/150  
-----  
train Loss: 1.6884 Acc: 0.9069  
val Loss: 1.7255 Acc: 0.8490  
Epoch 137/150  
-----  
train Loss: 1.7145 Acc: 0.8824  
val Loss: 1.7335 Acc: 0.8471  
Epoch 138/150  
-----  
train Loss: 1.6837 Acc: 0.9059  
val Loss: 1.7245 Acc: 0.8480  
Epoch 139/150  
-----  
train Loss: 1.6981 Acc: 0.9020  
val Loss: 1.7188 Acc: 0.8461  
Epoch 140/150
```

```
-----  
train Loss: 1.7118 Acc: 0.9020  
val Loss: 1.7146 Acc: 0.8471  
Epoch 141/150  
-----  
train Loss: 1.6824 Acc: 0.8980  
val Loss: 1.7274 Acc: 0.8461  
Epoch 142/150  
-----  
train Loss: 1.6801 Acc: 0.8990  
val Loss: 1.7293 Acc: 0.8431  
Epoch 143/150  
-----  
train Loss: 1.7053 Acc: 0.8882  
val Loss: 1.7166 Acc: 0.8461  
Epoch 144/150  
-----  
train Loss: 1.7269 Acc: 0.9029  
val Loss: 1.7185 Acc: 0.8480  
Epoch 145/150  
-----  
train Loss: 1.6885 Acc: 0.8922  
val Loss: 1.7252 Acc: 0.8480  
Epoch 146/150  
-----  
train Loss: 1.7042 Acc: 0.8892  
val Loss: 1.7246 Acc: 0.8461  
Epoch 147/150  
-----  
train Loss: 1.7077 Acc: 0.8941  
val Loss: 1.7224 Acc: 0.8471  
Epoch 148/150  
-----  
train Loss: 1.6884 Acc: 0.8863  
val Loss: 1.7225 Acc: 0.8471  
Epoch 149/150  
-----  
train Loss: 1.7209 Acc: 0.8912  
val Loss: 1.7332 Acc: 0.8471  
Epoch 150/150  
-----  
train Loss: 1.7002 Acc: 0.8784  
val Loss: 1.7113 Acc: 0.8490  
  
Learning rate: 0.01  
Epoch 1/150  
-----  
train Loss: 4.6069 Acc: 0.0363  
val Loss: 4.1248 Acc: 0.2108  
Epoch 2/150  
-----  
train Loss: 3.8968 Acc: 0.2608  
val Loss: 3.2682 Acc: 0.5373  
Epoch 3/150  
-----  
train Loss: 3.0715 Acc: 0.5755
```

```
val Loss: 2.5200 Acc: 0.6980
Epoch 4/150
-----
train Loss: 2.4332 Acc: 0.7157
val Loss: 2.0144 Acc: 0.7794
Epoch 5/150
-----
train Loss: 1.9128 Acc: 0.7902
val Loss: 1.6620 Acc: 0.8225
Epoch 6/150
-----
train Loss: 1.6116 Acc: 0.8333
val Loss: 1.4182 Acc: 0.8304
Epoch 7/150
-----
train Loss: 1.3435 Acc: 0.8569
val Loss: 1.2266 Acc: 0.8569
Epoch 8/150
-----
train Loss: 1.1899 Acc: 0.8686
val Loss: 1.1085 Acc: 0.8578
Epoch 9/150
-----
train Loss: 1.0523 Acc: 0.8941
val Loss: 1.0184 Acc: 0.8598
Epoch 10/150
-----
train Loss: 0.9323 Acc: 0.9069
val Loss: 0.9278 Acc: 0.8696
Epoch 11/150
-----
train Loss: 0.8488 Acc: 0.9098
val Loss: 0.8753 Acc: 0.8735
Epoch 12/150
-----
train Loss: 0.7684 Acc: 0.9186
val Loss: 0.8323 Acc: 0.8814
Epoch 13/150
-----
train Loss: 0.7152 Acc: 0.9353
val Loss: 0.7892 Acc: 0.8735
Epoch 14/150
-----
train Loss: 0.6632 Acc: 0.9373
val Loss: 0.7543 Acc: 0.8922
Epoch 15/150
-----
train Loss: 0.6728 Acc: 0.9176
val Loss: 0.7263 Acc: 0.8814
Epoch 16/150
-----
train Loss: 0.6002 Acc: 0.9245
val Loss: 0.7012 Acc: 0.8804
Epoch 17/150
-----
train Loss: 0.5568 Acc: 0.9343
```

```
val Loss: 0.6739 Acc: 0.8931
Epoch 18/150
-----
train Loss: 0.5524 Acc: 0.9373
val Loss: 0.6561 Acc: 0.8961
Epoch 19/150
-----
train Loss: 0.5271 Acc: 0.9363
val Loss: 0.6490 Acc: 0.8863
Epoch 20/150
-----
train Loss: 0.4714 Acc: 0.9549
val Loss: 0.6344 Acc: 0.8902
Epoch 21/150
-----
train Loss: 0.4715 Acc: 0.9392
val Loss: 0.6109 Acc: 0.8951
Epoch 22/150
-----
train Loss: 0.4308 Acc: 0.9490
val Loss: 0.5974 Acc: 0.8971
Epoch 23/150
-----
train Loss: 0.4394 Acc: 0.9510
val Loss: 0.5967 Acc: 0.8922
Epoch 24/150
-----
train Loss: 0.4062 Acc: 0.9549
val Loss: 0.5876 Acc: 0.8951
Epoch 25/150
-----
train Loss: 0.4248 Acc: 0.9461
val Loss: 0.5759 Acc: 0.9000
Epoch 26/150
-----
train Loss: 0.4283 Acc: 0.9422
val Loss: 0.5682 Acc: 0.8931
Epoch 27/150
-----
train Loss: 0.3894 Acc: 0.9549
val Loss: 0.5582 Acc: 0.8931
Epoch 28/150
-----
train Loss: 0.3614 Acc: 0.9539
val Loss: 0.5521 Acc: 0.9069
Epoch 29/150
-----
train Loss: 0.3870 Acc: 0.9569
val Loss: 0.5433 Acc: 0.8990
Epoch 30/150
-----
train Loss: 0.4100 Acc: 0.9392
val Loss: 0.5514 Acc: 0.8922
Epoch 31/150
-----
train Loss: 0.3616 Acc: 0.9569
```

```
val Loss: 0.5357 Acc: 0.8961
Epoch 32/150
-----
train Loss: 0.3659 Acc: 0.9569
val Loss: 0.5229 Acc: 0.9000
Epoch 33/150
-----
train Loss: 0.3313 Acc: 0.9598
val Loss: 0.5260 Acc: 0.9000
Epoch 34/150
-----
train Loss: 0.3339 Acc: 0.9559
val Loss: 0.5166 Acc: 0.8931
Epoch 35/150
-----
train Loss: 0.3321 Acc: 0.9676
val Loss: 0.5100 Acc: 0.9020
Epoch 36/150
-----
train Loss: 0.3136 Acc: 0.9627
val Loss: 0.4994 Acc: 0.9098
Epoch 37/150
-----
train Loss: 0.2993 Acc: 0.9588
val Loss: 0.5023 Acc: 0.9010
Epoch 38/150
-----
train Loss: 0.2979 Acc: 0.9618
val Loss: 0.4940 Acc: 0.9088
Epoch 39/150
-----
train Loss: 0.2733 Acc: 0.9676
val Loss: 0.4865 Acc: 0.9108
Epoch 40/150
-----
train Loss: 0.3253 Acc: 0.9559
val Loss: 0.4919 Acc: 0.8941
Epoch 41/150
-----
train Loss: 0.2897 Acc: 0.9618
val Loss: 0.4915 Acc: 0.8971
Epoch 42/150
-----
train Loss: 0.2826 Acc: 0.9637
val Loss: 0.4859 Acc: 0.9010
Epoch 43/150
-----
train Loss: 0.2607 Acc: 0.9578
val Loss: 0.4826 Acc: 0.9039
Epoch 44/150
-----
train Loss: 0.2809 Acc: 0.9637
val Loss: 0.4808 Acc: 0.9049
Epoch 45/150
-----
train Loss: 0.3017 Acc: 0.9559
```

```
val Loss: 0.4797 Acc: 0.9059
Epoch 46/150
-----
train Loss: 0.3142 Acc: 0.9529
val Loss: 0.4828 Acc: 0.9020
Epoch 47/150
-----
train Loss: 0.2810 Acc: 0.9608
val Loss: 0.4819 Acc: 0.9039
Epoch 48/150
-----
train Loss: 0.2635 Acc: 0.9676
val Loss: 0.4802 Acc: 0.9069
Epoch 49/150
-----
train Loss: 0.2659 Acc: 0.9676
val Loss: 0.4812 Acc: 0.9088
Epoch 50/150
-----
train Loss: 0.2735 Acc: 0.9706
val Loss: 0.4781 Acc: 0.9098
Epoch 51/150
-----
train Loss: 0.2470 Acc: 0.9755
val Loss: 0.4772 Acc: 0.9147
Epoch 52/150
-----
train Loss: 0.2591 Acc: 0.9696
val Loss: 0.4791 Acc: 0.9108
Epoch 53/150
-----
train Loss: 0.2805 Acc: 0.9598
val Loss: 0.4739 Acc: 0.9078
Epoch 54/150
-----
train Loss: 0.2908 Acc: 0.9598
val Loss: 0.4742 Acc: 0.9127
Epoch 55/150
-----
train Loss: 0.2787 Acc: 0.9608
val Loss: 0.4708 Acc: 0.9137
Epoch 56/150
-----
train Loss: 0.2712 Acc: 0.9676
val Loss: 0.4737 Acc: 0.9118
Epoch 57/150
-----
train Loss: 0.3008 Acc: 0.9588
val Loss: 0.4723 Acc: 0.9127
Epoch 58/150
-----
train Loss: 0.2728 Acc: 0.9676
val Loss: 0.4700 Acc: 0.9118
Epoch 59/150
-----
train Loss: 0.2572 Acc: 0.9657
```

```
val Loss: 0.4724 Acc: 0.9069
Epoch 60/150
-----
train Loss: 0.2985 Acc: 0.9598
val Loss: 0.4730 Acc: 0.9108
Epoch 61/150
-----
train Loss: 0.2781 Acc: 0.9647
val Loss: 0.4718 Acc: 0.9108
Epoch 62/150
-----
train Loss: 0.3001 Acc: 0.9627
val Loss: 0.4712 Acc: 0.9098
Epoch 63/150
-----
train Loss: 0.2921 Acc: 0.9598
val Loss: 0.4718 Acc: 0.9118
Epoch 64/150
-----
train Loss: 0.2681 Acc: 0.9637
val Loss: 0.4737 Acc: 0.9088
Epoch 65/150
-----
train Loss: 0.2340 Acc: 0.9745
val Loss: 0.4729 Acc: 0.9088
Epoch 66/150
-----
train Loss: 0.2991 Acc: 0.9500
val Loss: 0.4690 Acc: 0.9098
Epoch 67/150
-----
train Loss: 0.2766 Acc: 0.9637
val Loss: 0.4696 Acc: 0.9108
Epoch 68/150
-----
train Loss: 0.2344 Acc: 0.9745
val Loss: 0.4672 Acc: 0.9108
Epoch 69/150
-----
train Loss: 0.2768 Acc: 0.9637
val Loss: 0.4702 Acc: 0.9088
Epoch 70/150
-----
train Loss: 0.2497 Acc: 0.9745
val Loss: 0.4753 Acc: 0.9098
Epoch 71/150
-----
train Loss: 0.2538 Acc: 0.9706
val Loss: 0.4723 Acc: 0.9118
Epoch 72/150
-----
train Loss: 0.2427 Acc: 0.9706
val Loss: 0.4729 Acc: 0.9127
Epoch 73/150
-----
train Loss: 0.2536 Acc: 0.9765
```

```
val Loss: 0.4693 Acc: 0.9098
Epoch 74/150
-----
train Loss: 0.2428 Acc: 0.9706
val Loss: 0.4704 Acc: 0.9108
Epoch 75/150
-----
train Loss: 0.2717 Acc: 0.9627
val Loss: 0.4675 Acc: 0.9118
Epoch 76/150
-----
train Loss: 0.2406 Acc: 0.9696
val Loss: 0.4681 Acc: 0.9147
Epoch 77/150
-----
train Loss: 0.2623 Acc: 0.9627
val Loss: 0.4691 Acc: 0.9118
Epoch 78/150
-----
train Loss: 0.2664 Acc: 0.9578
val Loss: 0.4669 Acc: 0.9167
Epoch 79/150
-----
train Loss: 0.2591 Acc: 0.9667
val Loss: 0.4677 Acc: 0.9118
Epoch 80/150
-----
train Loss: 0.2677 Acc: 0.9657
val Loss: 0.4693 Acc: 0.9098
Epoch 81/150
-----
train Loss: 0.2843 Acc: 0.9598
val Loss: 0.4672 Acc: 0.9137
Epoch 82/150
-----
train Loss: 0.2603 Acc: 0.9706
val Loss: 0.4657 Acc: 0.9108
Epoch 83/150
-----
train Loss: 0.2486 Acc: 0.9716
val Loss: 0.4675 Acc: 0.9069
Epoch 84/150
-----
train Loss: 0.2548 Acc: 0.9647
val Loss: 0.4678 Acc: 0.9118
Epoch 85/150
-----
train Loss: 0.2624 Acc: 0.9686
val Loss: 0.4669 Acc: 0.9118
Epoch 86/150
-----
train Loss: 0.2667 Acc: 0.9667
val Loss: 0.4676 Acc: 0.9108
Epoch 87/150
-----
train Loss: 0.2776 Acc: 0.9676
```

```
val Loss: 0.4683 Acc: 0.9078
Epoch 88/150
-----
train Loss: 0.2452 Acc: 0.9657
val Loss: 0.4691 Acc: 0.9137
Epoch 89/150
-----
train Loss: 0.2407 Acc: 0.9706
val Loss: 0.4666 Acc: 0.9118
Epoch 90/150
-----
train Loss: 0.2366 Acc: 0.9745
val Loss: 0.4691 Acc: 0.9108
Epoch 91/150
-----
train Loss: 0.2437 Acc: 0.9775
val Loss: 0.4734 Acc: 0.9127
Epoch 92/150
-----
train Loss: 0.2828 Acc: 0.9627
val Loss: 0.4685 Acc: 0.9118
Epoch 93/150
-----
train Loss: 0.2916 Acc: 0.9578
val Loss: 0.4667 Acc: 0.9137
Epoch 94/150
-----
train Loss: 0.2816 Acc: 0.9569
val Loss: 0.4668 Acc: 0.9118
Epoch 95/150
-----
train Loss: 0.2517 Acc: 0.9627
val Loss: 0.4696 Acc: 0.9127
Epoch 96/150
-----
train Loss: 0.2603 Acc: 0.9657
val Loss: 0.4696 Acc: 0.9127
Epoch 97/150
-----
train Loss: 0.2670 Acc: 0.9637
val Loss: 0.4683 Acc: 0.9118
Epoch 98/150
-----
train Loss: 0.2712 Acc: 0.9588
val Loss: 0.4679 Acc: 0.9088
Epoch 99/150
-----
train Loss: 0.2578 Acc: 0.9725
val Loss: 0.4662 Acc: 0.9108
Epoch 100/150
-----
train Loss: 0.2707 Acc: 0.9608
val Loss: 0.4668 Acc: 0.9088
Epoch 101/150
-----
train Loss: 0.2716 Acc: 0.9588
```

```
val Loss: 0.4646 Acc: 0.9098
Epoch 102/150
-----
train Loss: 0.2698 Acc: 0.9647
val Loss: 0.4676 Acc: 0.9098
Epoch 103/150
-----
train Loss: 0.2598 Acc: 0.9716
val Loss: 0.4653 Acc: 0.9108
Epoch 104/150
-----
train Loss: 0.2335 Acc: 0.9765
val Loss: 0.4678 Acc: 0.9118
Epoch 105/150
-----
train Loss: 0.2755 Acc: 0.9647
val Loss: 0.4685 Acc: 0.9118
Epoch 106/150
-----
train Loss: 0.2576 Acc: 0.9647
val Loss: 0.4679 Acc: 0.9098
Epoch 107/150
-----
train Loss: 0.2777 Acc: 0.9676
val Loss: 0.4669 Acc: 0.9078
Epoch 108/150
-----
train Loss: 0.2714 Acc: 0.9657
val Loss: 0.4688 Acc: 0.9098
Epoch 109/150
-----
train Loss: 0.2582 Acc: 0.9686
val Loss: 0.4682 Acc: 0.9118
Epoch 110/150
-----
train Loss: 0.2505 Acc: 0.9716
val Loss: 0.4650 Acc: 0.9088
Epoch 111/150
-----
train Loss: 0.2732 Acc: 0.9598
val Loss: 0.4674 Acc: 0.9088
Epoch 112/150
-----
train Loss: 0.2363 Acc: 0.9725
val Loss: 0.4668 Acc: 0.9069
Epoch 113/150
-----
train Loss: 0.2720 Acc: 0.9627
val Loss: 0.4647 Acc: 0.9098
Epoch 114/150
-----
train Loss: 0.2749 Acc: 0.9588
val Loss: 0.4664 Acc: 0.9137
Epoch 115/150
-----
train Loss: 0.2976 Acc: 0.9539
```

```
val Loss: 0.4688 Acc: 0.9108
Epoch 116/150
-----
train Loss: 0.2546 Acc: 0.9627
val Loss: 0.4696 Acc: 0.9069
Epoch 117/150
-----
train Loss: 0.2449 Acc: 0.9735
val Loss: 0.4641 Acc: 0.9108
Epoch 118/150
-----
train Loss: 0.2766 Acc: 0.9618
val Loss: 0.4664 Acc: 0.9098
Epoch 119/150
-----
train Loss: 0.2861 Acc: 0.9559
val Loss: 0.4657 Acc: 0.9108
Epoch 120/150
-----
train Loss: 0.2720 Acc: 0.9657
val Loss: 0.4641 Acc: 0.9098
Epoch 121/150
-----
train Loss: 0.2605 Acc: 0.9667
val Loss: 0.4654 Acc: 0.9088
Epoch 122/150
-----
train Loss: 0.2473 Acc: 0.9696
val Loss: 0.4661 Acc: 0.9098
Epoch 123/150
-----
train Loss: 0.2674 Acc: 0.9686
val Loss: 0.4676 Acc: 0.9108
Epoch 124/150
-----
train Loss: 0.2535 Acc: 0.9725
val Loss: 0.4686 Acc: 0.9098
Epoch 125/150
-----
train Loss: 0.2758 Acc: 0.9676
val Loss: 0.4655 Acc: 0.9088
Epoch 126/150
-----
train Loss: 0.2704 Acc: 0.9627
val Loss: 0.4656 Acc: 0.9127
Epoch 127/150
-----
train Loss: 0.2674 Acc: 0.9657
val Loss: 0.4640 Acc: 0.9127
Epoch 128/150
-----
train Loss: 0.2277 Acc: 0.9725
val Loss: 0.4703 Acc: 0.9098
Epoch 129/150
-----
train Loss: 0.2410 Acc: 0.9765
```

```
val Loss: 0.4678 Acc: 0.9098
Epoch 130/150
-----
train Loss: 0.2500 Acc: 0.9755
val Loss: 0.4660 Acc: 0.9088
Epoch 131/150
-----
train Loss: 0.2553 Acc: 0.9745
val Loss: 0.4686 Acc: 0.9127
Epoch 132/150
-----
train Loss: 0.2662 Acc: 0.9647
val Loss: 0.4648 Acc: 0.9088
Epoch 133/150
-----
train Loss: 0.2405 Acc: 0.9755
val Loss: 0.4662 Acc: 0.9127
Epoch 134/150
-----
train Loss: 0.2746 Acc: 0.9647
val Loss: 0.4654 Acc: 0.9108
Epoch 135/150
-----
train Loss: 0.2429 Acc: 0.9667
val Loss: 0.4627 Acc: 0.9147
Epoch 136/150
-----
train Loss: 0.2723 Acc: 0.9598
val Loss: 0.4631 Acc: 0.9118
Epoch 137/150
-----
train Loss: 0.2561 Acc: 0.9696
val Loss: 0.4639 Acc: 0.9118
Epoch 138/150
-----
train Loss: 0.2698 Acc: 0.9647
val Loss: 0.4664 Acc: 0.9078
Epoch 139/150
-----
train Loss: 0.2541 Acc: 0.9706
val Loss: 0.4648 Acc: 0.9108
Epoch 140/150
-----
train Loss: 0.2715 Acc: 0.9598
val Loss: 0.4635 Acc: 0.9088
Epoch 141/150
-----
train Loss: 0.2683 Acc: 0.9637
val Loss: 0.4632 Acc: 0.9108
Epoch 142/150
-----
train Loss: 0.2518 Acc: 0.9686
val Loss: 0.4673 Acc: 0.9118
Epoch 143/150
-----
train Loss: 0.2721 Acc: 0.9578
```

```
val Loss: 0.4616 Acc: 0.9098
Epoch 144/150
-----
train Loss: 0.2482 Acc: 0.9716
val Loss: 0.4691 Acc: 0.9108
Epoch 145/150
-----
train Loss: 0.2493 Acc: 0.9735
val Loss: 0.4671 Acc: 0.9098
Epoch 146/150
-----
train Loss: 0.2510 Acc: 0.9686
val Loss: 0.4663 Acc: 0.9108
Epoch 147/150
-----
train Loss: 0.2622 Acc: 0.9706
val Loss: 0.4666 Acc: 0.9118
Epoch 148/150
-----
train Loss: 0.2725 Acc: 0.9696
val Loss: 0.4657 Acc: 0.9078
Epoch 149/150
-----
train Loss: 0.2667 Acc: 0.9627
val Loss: 0.4663 Acc: 0.9118
Epoch 150/150
-----
train Loss: 0.2233 Acc: 0.9755
val Loss: 0.4678 Acc: 0.9137

Learning rate: 0.1
Epoch 1/150
-----
train Loss: 4.8377 Acc: 0.1098
val Loss: 4.2301 Acc: 0.4461
Epoch 2/150
-----
train Loss: 3.4977 Acc: 0.5137
val Loss: 2.8727 Acc: 0.5461
Epoch 3/150
-----
train Loss: 1.7516 Acc: 0.6784
val Loss: 1.2519 Acc: 0.7147
Epoch 4/150
-----
train Loss: 0.9799 Acc: 0.7588
val Loss: 1.0534 Acc: 0.7608
Epoch 5/150
-----
train Loss: 0.6264 Acc: 0.8324
val Loss: 0.9116 Acc: 0.7755
Epoch 6/150
-----
train Loss: 0.5186 Acc: 0.8667
val Loss: 1.0276 Acc: 0.7735
Epoch 7/150
```

```
-----  
train Loss: 0.5671 Acc: 0.8608  
val Loss: 0.8117 Acc: 0.8069  
Epoch 8/150  
-----  
train Loss: 0.4751 Acc: 0.8765  
val Loss: 0.9815 Acc: 0.7588  
Epoch 9/150  
-----  
train Loss: 0.5232 Acc: 0.8676  
val Loss: 0.7465 Acc: 0.8216  
Epoch 10/150  
-----  
train Loss: 0.4322 Acc: 0.8892  
val Loss: 0.6641 Acc: 0.8480  
Epoch 11/150  
-----  
train Loss: 0.3004 Acc: 0.9186  
val Loss: 0.6853 Acc: 0.8275  
Epoch 12/150  
-----  
train Loss: 0.2959 Acc: 0.9137  
val Loss: 0.6324 Acc: 0.8392  
Epoch 13/150  
-----  
train Loss: 0.2824 Acc: 0.9304  
val Loss: 0.7364 Acc: 0.8265  
Epoch 14/150  
-----  
train Loss: 0.3027 Acc: 0.9255  
val Loss: 0.7224 Acc: 0.8314  
Epoch 15/150  
-----  
train Loss: 0.2438 Acc: 0.9314  
val Loss: 0.7018 Acc: 0.8255  
Epoch 16/150  
-----  
train Loss: 0.2554 Acc: 0.9275  
val Loss: 0.5622 Acc: 0.8696  
Epoch 17/150  
-----  
train Loss: 0.2770 Acc: 0.9294  
val Loss: 0.6255 Acc: 0.8402  
Epoch 18/150  
-----  
train Loss: 0.2269 Acc: 0.9412  
val Loss: 0.5702 Acc: 0.8735  
Epoch 19/150  
-----  
train Loss: 0.2246 Acc: 0.9451  
val Loss: 0.7790 Acc: 0.8245  
Epoch 20/150  
-----  
train Loss: 0.3026 Acc: 0.9294  
val Loss: 0.6336 Acc: 0.8500  
Epoch 21/150
```

```
-----  
train Loss: 0.2536 Acc: 0.9363  
val Loss: 0.5982 Acc: 0.8480  
Epoch 22/150  
-----  
train Loss: 0.2034 Acc: 0.9353  
val Loss: 0.5397 Acc: 0.8618  
Epoch 23/150  
-----  
train Loss: 0.1885 Acc: 0.9539  
val Loss: 0.6306 Acc: 0.8500  
Epoch 24/150  
-----  
train Loss: 0.2021 Acc: 0.9480  
val Loss: 0.6045 Acc: 0.8549  
Epoch 25/150  
-----  
train Loss: 0.1699 Acc: 0.9500  
val Loss: 0.6743 Acc: 0.8490  
Epoch 26/150  
-----  
train Loss: 0.2281 Acc: 0.9431  
val Loss: 0.7591 Acc: 0.8235  
Epoch 27/150  
-----  
train Loss: 0.2086 Acc: 0.9373  
val Loss: 0.6890 Acc: 0.8461  
Epoch 28/150  
-----  
train Loss: 0.1731 Acc: 0.9451  
val Loss: 0.5987 Acc: 0.8618  
Epoch 29/150  
-----  
train Loss: 0.1956 Acc: 0.9402  
val Loss: 0.7153 Acc: 0.8363  
Epoch 30/150  
-----  
train Loss: 0.2186 Acc: 0.9441  
val Loss: 0.7242 Acc: 0.8471  
Epoch 31/150  
-----  
train Loss: 0.1897 Acc: 0.9490  
val Loss: 0.5807 Acc: 0.8510  
Epoch 32/150  
-----  
train Loss: 0.1626 Acc: 0.9539  
val Loss: 0.6358 Acc: 0.8618  
Epoch 33/150  
-----  
train Loss: 0.1752 Acc: 0.9480  
val Loss: 0.6391 Acc: 0.8529  
Epoch 34/150  
-----  
train Loss: 0.2028 Acc: 0.9461  
val Loss: 0.6249 Acc: 0.8510  
Epoch 35/150
```

```
-----  
train Loss: 0.1796 Acc: 0.9549  
val Loss: 0.6463 Acc: 0.8510  
Epoch 36/150  
-----  
train Loss: 0.1792 Acc: 0.9510  
val Loss: 0.6252 Acc: 0.8510  
Epoch 37/150  
-----  
train Loss: 0.1482 Acc: 0.9627  
val Loss: 0.5935 Acc: 0.8618  
Epoch 38/150  
-----  
train Loss: 0.1914 Acc: 0.9520  
val Loss: 0.7232 Acc: 0.8539  
Epoch 39/150  
-----  
train Loss: 0.1452 Acc: 0.9627  
val Loss: 0.6787 Acc: 0.8441  
Epoch 40/150  
-----  
train Loss: 0.1682 Acc: 0.9490  
val Loss: 0.6730 Acc: 0.8480  
Epoch 41/150  
-----  
train Loss: 0.1637 Acc: 0.9549  
val Loss: 0.5625 Acc: 0.8706  
Epoch 42/150  
-----  
train Loss: 0.1264 Acc: 0.9637  
val Loss: 0.4970 Acc: 0.8941  
Epoch 43/150  
-----  
train Loss: 0.1199 Acc: 0.9735  
val Loss: 0.4750 Acc: 0.8922  
Epoch 44/150  
-----  
train Loss: 0.0916 Acc: 0.9755  
val Loss: 0.4598 Acc: 0.8990  
Epoch 45/150  
-----  
train Loss: 0.0791 Acc: 0.9804  
val Loss: 0.4495 Acc: 0.8990  
Epoch 46/150  
-----  
train Loss: 0.0931 Acc: 0.9725  
val Loss: 0.4389 Acc: 0.9029  
Epoch 47/150  
-----  
train Loss: 0.0790 Acc: 0.9784  
val Loss: 0.4471 Acc: 0.8990  
Epoch 48/150  
-----  
train Loss: 0.0888 Acc: 0.9804  
val Loss: 0.4482 Acc: 0.9020  
Epoch 49/150
```

```
-----  
train Loss: 0.1036 Acc: 0.9725  
val Loss: 0.4505 Acc: 0.9029  
Epoch 50/150  
-----  
train Loss: 0.0891 Acc: 0.9755  
val Loss: 0.4490 Acc: 0.9049  
Epoch 51/150  
-----  
train Loss: 0.0895 Acc: 0.9784  
val Loss: 0.4462 Acc: 0.9039  
Epoch 52/150  
-----  
train Loss: 0.0643 Acc: 0.9843  
val Loss: 0.4464 Acc: 0.9049  
Epoch 53/150  
-----  
train Loss: 0.0603 Acc: 0.9843  
val Loss: 0.4439 Acc: 0.9059  
Epoch 54/150  
-----  
train Loss: 0.0604 Acc: 0.9833  
val Loss: 0.4502 Acc: 0.9020  
Epoch 55/150  
-----  
train Loss: 0.0714 Acc: 0.9833  
val Loss: 0.4621 Acc: 0.9000  
Epoch 56/150  
-----  
train Loss: 0.0624 Acc: 0.9804  
val Loss: 0.4688 Acc: 0.8961  
Epoch 57/150  
-----  
train Loss: 0.0947 Acc: 0.9745  
val Loss: 0.4585 Acc: 0.8971  
Epoch 58/150  
-----  
train Loss: 0.1120 Acc: 0.9745  
val Loss: 0.4505 Acc: 0.9029  
Epoch 59/150  
-----  
train Loss: 0.0613 Acc: 0.9833  
val Loss: 0.4515 Acc: 0.9039  
Epoch 60/150  
-----  
train Loss: 0.0558 Acc: 0.9804  
val Loss: 0.4504 Acc: 0.9069  
Epoch 61/150  
-----  
train Loss: 0.0861 Acc: 0.9755  
val Loss: 0.4515 Acc: 0.8990  
Epoch 62/150  
-----  
train Loss: 0.1168 Acc: 0.9667  
val Loss: 0.4497 Acc: 0.9010  
Epoch 63/150
```

```
-----  
train Loss: 0.0633 Acc: 0.9843  
val Loss: 0.4413 Acc: 0.9020  
Epoch 64/150  
-----  
train Loss: 0.0947 Acc: 0.9765  
val Loss: 0.4417 Acc: 0.9039  
Epoch 65/150  
-----  
train Loss: 0.0669 Acc: 0.9843  
val Loss: 0.4401 Acc: 0.9000  
Epoch 66/150  
-----  
train Loss: 0.0723 Acc: 0.9853  
val Loss: 0.4415 Acc: 0.9000  
Epoch 67/150  
-----  
train Loss: 0.0744 Acc: 0.9814  
val Loss: 0.4390 Acc: 0.8990  
Epoch 68/150  
-----  
train Loss: 0.0752 Acc: 0.9814  
val Loss: 0.4398 Acc: 0.9010  
Epoch 69/150  
-----  
train Loss: 0.0861 Acc: 0.9804  
val Loss: 0.4453 Acc: 0.9059  
Epoch 70/150  
-----  
train Loss: 0.0882 Acc: 0.9794  
val Loss: 0.4560 Acc: 0.9039  
Epoch 71/150  
-----  
train Loss: 0.0679 Acc: 0.9833  
val Loss: 0.4655 Acc: 0.9059  
Epoch 72/150  
-----  
train Loss: 0.0743 Acc: 0.9833  
val Loss: 0.4610 Acc: 0.9000  
Epoch 73/150  
-----  
train Loss: 0.0762 Acc: 0.9833  
val Loss: 0.4538 Acc: 0.8971  
Epoch 74/150  
-----  
train Loss: 0.0834 Acc: 0.9765  
val Loss: 0.4476 Acc: 0.9039  
Epoch 75/150  
-----  
train Loss: 0.0822 Acc: 0.9755  
val Loss: 0.4551 Acc: 0.9010  
Epoch 76/150  
-----  
train Loss: 0.0969 Acc: 0.9725  
val Loss: 0.4595 Acc: 0.8980  
Epoch 77/150
```

```
-----  
train Loss: 0.0501 Acc: 0.9863  
val Loss: 0.4569 Acc: 0.8990  
Epoch 78/150  
-----  
train Loss: 0.0794 Acc: 0.9794  
val Loss: 0.4570 Acc: 0.8980  
Epoch 79/150  
-----  
train Loss: 0.0796 Acc: 0.9804  
val Loss: 0.4554 Acc: 0.9000  
Epoch 80/150  
-----  
train Loss: 0.0946 Acc: 0.9745  
val Loss: 0.4519 Acc: 0.9000  
Epoch 81/150  
-----  
train Loss: 0.0571 Acc: 0.9853  
val Loss: 0.4506 Acc: 0.9000  
Epoch 82/150  
-----  
train Loss: 0.0748 Acc: 0.9794  
val Loss: 0.4491 Acc: 0.8971  
Epoch 83/150  
-----  
train Loss: 0.0754 Acc: 0.9843  
val Loss: 0.4482 Acc: 0.9010  
Epoch 84/150  
-----  
train Loss: 0.0912 Acc: 0.9765  
val Loss: 0.4497 Acc: 0.8980  
Epoch 85/150  
-----  
train Loss: 0.0802 Acc: 0.9784  
val Loss: 0.4540 Acc: 0.8980  
Epoch 86/150  
-----  
train Loss: 0.0687 Acc: 0.9814  
val Loss: 0.4494 Acc: 0.8980  
Epoch 87/150  
-----  
train Loss: 0.0840 Acc: 0.9755  
val Loss: 0.4485 Acc: 0.9010  
Epoch 88/150  
-----  
train Loss: 0.0609 Acc: 0.9863  
val Loss: 0.4469 Acc: 0.9000  
Epoch 89/150  
-----  
train Loss: 0.1054 Acc: 0.9775  
val Loss: 0.4492 Acc: 0.9010  
Epoch 90/150  
-----  
train Loss: 0.0682 Acc: 0.9794  
val Loss: 0.4480 Acc: 0.9010  
Epoch 91/150
```

```
-----  
train Loss: 0.0896 Acc: 0.9784  
val Loss: 0.4443 Acc: 0.9020  
Epoch 92/150  
-----  
train Loss: 0.0795 Acc: 0.9784  
val Loss: 0.4485 Acc: 0.9020  
Epoch 93/150  
-----  
train Loss: 0.0768 Acc: 0.9824  
val Loss: 0.4489 Acc: 0.9020  
Epoch 94/150  
-----  
train Loss: 0.0754 Acc: 0.9804  
val Loss: 0.4460 Acc: 0.9029  
Epoch 95/150  
-----  
train Loss: 0.0753 Acc: 0.9745  
val Loss: 0.4479 Acc: 0.8990  
Epoch 96/150  
-----  
train Loss: 0.0725 Acc: 0.9843  
val Loss: 0.4430 Acc: 0.9000  
Epoch 97/150  
-----  
train Loss: 0.0729 Acc: 0.9824  
val Loss: 0.4460 Acc: 0.9059  
Epoch 98/150  
-----  
train Loss: 0.1027 Acc: 0.9755  
val Loss: 0.4465 Acc: 0.9010  
Epoch 99/150  
-----  
train Loss: 0.0568 Acc: 0.9833  
val Loss: 0.4463 Acc: 0.9020  
Epoch 100/150  
-----  
train Loss: 0.0649 Acc: 0.9814  
val Loss: 0.4464 Acc: 0.9010  
Epoch 101/150  
-----  
train Loss: 0.0995 Acc: 0.9725  
val Loss: 0.4439 Acc: 0.9010  
Epoch 102/150  
-----  
train Loss: 0.0719 Acc: 0.9833  
val Loss: 0.4457 Acc: 0.9010  
Epoch 103/150  
-----  
train Loss: 0.0907 Acc: 0.9735  
val Loss: 0.4431 Acc: 0.9010  
Epoch 104/150  
-----  
train Loss: 0.0634 Acc: 0.9804  
val Loss: 0.4446 Acc: 0.9039  
Epoch 105/150
```

```
-----  
train Loss: 0.0768 Acc: 0.9833  
val Loss: 0.4409 Acc: 0.9020  
Epoch 106/150  
-----  
train Loss: 0.0782 Acc: 0.9755  
val Loss: 0.4421 Acc: 0.9039  
Epoch 107/150  
-----  
train Loss: 0.0585 Acc: 0.9824  
val Loss: 0.4436 Acc: 0.9039  
Epoch 108/150  
-----  
train Loss: 0.0877 Acc: 0.9794  
val Loss: 0.4427 Acc: 0.9029  
Epoch 109/150  
-----  
train Loss: 0.0639 Acc: 0.9824  
val Loss: 0.4456 Acc: 0.9010  
Epoch 110/150  
-----  
train Loss: 0.0687 Acc: 0.9853  
val Loss: 0.4423 Acc: 0.9049  
Epoch 111/150  
-----  
train Loss: 0.0622 Acc: 0.9804  
val Loss: 0.4442 Acc: 0.9059  
Epoch 112/150  
-----  
train Loss: 0.0929 Acc: 0.9735  
val Loss: 0.4431 Acc: 0.9029  
Epoch 113/150  
-----  
train Loss: 0.0605 Acc: 0.9824  
val Loss: 0.4440 Acc: 0.9010  
Epoch 114/150  
-----  
train Loss: 0.0617 Acc: 0.9814  
val Loss: 0.4439 Acc: 0.9029  
Epoch 115/150  
-----  
train Loss: 0.0825 Acc: 0.9794  
val Loss: 0.4416 Acc: 0.9049  
Epoch 116/150  
-----  
train Loss: 0.0943 Acc: 0.9784  
val Loss: 0.4399 Acc: 0.9029  
Epoch 117/150  
-----  
train Loss: 0.0966 Acc: 0.9725  
val Loss: 0.4431 Acc: 0.9010  
Epoch 118/150  
-----  
train Loss: 0.0915 Acc: 0.9765  
val Loss: 0.4437 Acc: 0.9010  
Epoch 119/150
```

```
-----  
train Loss: 0.0715 Acc: 0.9794  
val Loss: 0.4412 Acc: 0.9039  
Epoch 120/150  
-----  
train Loss: 0.0907 Acc: 0.9725  
val Loss: 0.4470 Acc: 0.9029  
Epoch 121/150  
-----  
train Loss: 0.0539 Acc: 0.9873  
val Loss: 0.4422 Acc: 0.9029  
Epoch 122/150  
-----  
train Loss: 0.0700 Acc: 0.9824  
val Loss: 0.4433 Acc: 0.9010  
Epoch 123/150  
-----  
train Loss: 0.0518 Acc: 0.9853  
val Loss: 0.4468 Acc: 0.9049  
Epoch 124/150  
-----  
train Loss: 0.0797 Acc: 0.9775  
val Loss: 0.4477 Acc: 0.9059  
Epoch 125/150  
-----  
train Loss: 0.0714 Acc: 0.9833  
val Loss: 0.4479 Acc: 0.9039  
Epoch 126/150  
-----  
train Loss: 0.0803 Acc: 0.9794  
val Loss: 0.4437 Acc: 0.9029  
Epoch 127/150  
-----  
train Loss: 0.0553 Acc: 0.9873  
val Loss: 0.4459 Acc: 0.9020  
Epoch 128/150  
-----  
train Loss: 0.0724 Acc: 0.9863  
val Loss: 0.4435 Acc: 0.9000  
Epoch 129/150  
-----  
train Loss: 0.0846 Acc: 0.9784  
val Loss: 0.4444 Acc: 0.9039  
Epoch 130/150  
-----  
train Loss: 0.0682 Acc: 0.9794  
val Loss: 0.4443 Acc: 0.9039  
Epoch 131/150  
-----  
train Loss: 0.0721 Acc: 0.9833  
val Loss: 0.4418 Acc: 0.9020  
Epoch 132/150  
-----  
train Loss: 0.0675 Acc: 0.9814  
val Loss: 0.4428 Acc: 0.9039  
Epoch 133/150
```

```
-----  
train Loss: 0.0889 Acc: 0.9765  
val Loss: 0.4432 Acc: 0.9029  
Epoch 134/150  
-----  
train Loss: 0.0798 Acc: 0.9824  
val Loss: 0.4417 Acc: 0.9039  
Epoch 135/150  
-----  
train Loss: 0.0998 Acc: 0.9794  
val Loss: 0.4443 Acc: 0.9029  
Epoch 136/150  
-----  
train Loss: 0.0693 Acc: 0.9824  
val Loss: 0.4421 Acc: 0.9000  
Epoch 137/150  
-----  
train Loss: 0.0714 Acc: 0.9804  
val Loss: 0.4460 Acc: 0.9039  
Epoch 138/150  
-----  
train Loss: 0.0725 Acc: 0.9863  
val Loss: 0.4431 Acc: 0.9029  
Epoch 139/150  
-----  
train Loss: 0.0630 Acc: 0.9824  
val Loss: 0.4441 Acc: 0.9039  
Epoch 140/150  
-----  
train Loss: 0.0728 Acc: 0.9804  
val Loss: 0.4399 Acc: 0.9039  
Epoch 141/150  
-----  
train Loss: 0.1225 Acc: 0.9686  
val Loss: 0.4430 Acc: 0.9029  
Epoch 142/150  
-----  
train Loss: 0.0816 Acc: 0.9765  
val Loss: 0.4437 Acc: 0.9059  
Epoch 143/150  
-----  
train Loss: 0.0896 Acc: 0.9745  
val Loss: 0.4455 Acc: 0.9069  
Epoch 144/150  
-----  
train Loss: 0.0988 Acc: 0.9755  
val Loss: 0.4411 Acc: 0.9059  
Epoch 145/150  
-----  
train Loss: 0.0799 Acc: 0.9814  
val Loss: 0.4468 Acc: 0.9029  
Epoch 146/150  
-----  
train Loss: 0.0638 Acc: 0.9804  
val Loss: 0.4435 Acc: 0.9020  
Epoch 147/150
```

```
-----  
train Loss: 0.0791 Acc: 0.9765  
val Loss: 0.4420 Acc: 0.9049  
Epoch 148/150  
-----  
train Loss: 0.0679 Acc: 0.9833  
val Loss: 0.4388 Acc: 0.9049  
Epoch 149/150  
-----  
train Loss: 0.0740 Acc: 0.9804  
val Loss: 0.4426 Acc: 0.9039  
Epoch 150/150  
-----  
train Loss: 0.1016 Acc: 0.9735  
val Loss: 0.4428 Acc: 0.9039
```

b. Final comparison of approaches

Find the best final accuracy (across all the learning rates) from the two transfer learning approaches. Which approach and learning rate was the winner? Provide a plausible explanation to support your observation.

(4 points)

After looking through all our experiments, I found some really interesting patterns! Let me break down what I discovered:

For finetuning (where we trained all layers):

- With a learning rate of 0.001: The model did pretty well, hitting 93.14% accuracy, but wasn't quite our best
- With 0.01: This was our star performer! It reached an impressive 94.12% accuracy
- With 0.1: This was way too aggressive - the model struggled and only managed 43.63% accuracy

For feature extraction (where we only trained the last layer):

- With 0.001: Got to 84.90% accuracy, but clearly not optimal
- With 0.01: Better results at 91.37% accuracy
- With 0.1: a bit lower performance at 90.39% accuracy

The clear winner here was finetuning with a learning rate of 0.01, reaching that 94.12% sweet spot. Let me explain why I think this worked so well:

First off, that 0.01 learning rate turned out to be just right - kind of like the "Goldilocks zone" for our model. The 0.001 rate was playing it too safe (though it still did well), while 0.1 was just too aggressive and made the model unstable. 0.01 hit that perfect balance where it could learn effectively without going off the rails.

What's particularly interesting is how finetuning beat feature extraction. While feature extraction is often a safe bet (and did give us decent results), finetuning all layers really let the model adapt to the specific features of flowers. Think about it - while ImageNet gave us great general features to start with, flowers have their own unique patterns and details. By carefully tuning all layers with that 0.01 learning rate, we let the model refine both basic features (like edge detection) and complex features (like petal arrangements) specifically for flower recognition.

I think what really made this work was our learning rate schedule - dropping the rate at specific points (epochs 40, 80, and 120) helped the model fine-tune those features with increasing precision. It's like starting with broad brush strokes and gradually moving to finer and finer details.

The numbers really tell the story - that 94.12% accuracy shows that letting the model adapt all its layers, while keeping the learning controlled with the right rate and schedule, was definitely the way to go!

References

- Pytorch blog. Transfer Learning for Computer Vision Tutorial by S. Chilamkurthy
Available at
https://pytorch.org/tutorials/beginner/transfer_learning_tutorial.html
- Notes on Transfer Learning. CS231n Convolutional Neural Networks for Visual Recognition
Available at <https://cs231n.github.io/transfer-learning/>
- [Visual Domain Decathlon](#)

Problem 3 - *Data Parallelism in Pytorch* (20 points)

We are going to experiment with PyTorch's DataParallel Module, which is PyTorch's Synchronous SGD implementation across a number of GPUs on the same server. In particular, we will train ResNet-18 implementation from <https://github.com/kuangliu/pytorch-cifar> with `num_workers=2`, running up to 4 GPUs with the DataParallel (DP) Module. Use SGD optimizers with 0.1 as the learning rate, momentum 0.9, and weight decay 5e-4. For this question, you need to experiment with multiple GPUs on the same server. You may need to execute this on the NYU Greene Cluster.

Create a PyTorch program with a DataLoader that loads the images and the related labels from the torchvision CIFAR10 dataset. Import the CIFAR10 dataset from the torchvision package, with the following sequence of transformations:

- Random cropping, with size 32x32 and padding 4
- Random horizontal flipping with a probability of 0.5
- Normalize each image's RGB channel with mean(0.4914, 0.4822, 0.4465) and variance (0.2023, 0.1994, 0.2010)

The DataLoader for the training set uses a minibatch size of 128 and 3 I/O processes (i.e., `num_workers=2`). The DataLoader for the testing set uses a minibatch size of 100 and 3 I/O processes (i.e., `num_workers=2`). Create a main function that creates the DataLoaders for the training set and the neural network.

1. Measure training time for different batch sizes on a single GPU

Measure how long it takes to complete 1 epoch of training using different batch sizes on a single GPU. Start with a batch size of 32 and increase by 4-fold for each measurement (i.e., 32, 128, 512, etc.) until the single GPU memory cannot hold the batch size. For each run, perform 2 epochs; the first epoch is used to warm up CPU/GPU cache, and you should report the training time (excluding data I/O but including data movement from CPU to GPU, gradient calculation, and weight updates) based on the 2nd epoch training.

(5 points)

In [16]:

```
import torch
import torchvision
import torchvision.transforms as transforms
import torch.nn as nn
import torch.optim as optim
import time
import numpy as np
from torch.utils.data import DataLoader
from torch.nn.parallel import DataParallel

# Data loading setup
transform_train = transforms.Compose([
    transforms.RandomCrop(32, padding=4),
    transforms.RandomHorizontalFlip(),
    transforms.ToTensor(),
    transforms.Normalize((0.4914, 0.4822, 0.4465), (0.2023, 0.1994, 0.2010))
])

transform_test = transforms.Compose([
    transforms.ToTensor(),
```

```
        transforms.Normalize((0.4914, 0.4822, 0.4465), (0.2023, 0.1994, 0.2010))
])
```

```
In [17]: def measure_single_gpu_time(batch_size):
    # Load CIFAR10 dataset
    trainset = torchvision.datasets.CIFAR10(root='./data', train=True,
                                             download=True, transform=transform)
    trainloader = DataLoader(trainset, batch_size=batch_size,
                             shuffle=True, num_workers=2)

    # Initialize model
    device = torch.device("cuda:0" if torch.cuda.is_available() else "cpu")
    model = torchvision.models.resnet18(pretrained=False)
    model = model.to(device)

    criterion = nn.CrossEntropyLoss()
    optimizer = optim.SGD(model.parameters(), lr=0.1, momentum=0.9, weight_c

    # Warmup epoch
    model.train()
    for inputs, targets in trainloader:
        inputs, targets = inputs.to(device), targets.to(device)
        optimizer.zero_grad()
        outputs = model(inputs)
        loss = criterion(outputs, targets)
        loss.backward()
        optimizer.step()

    # Measure second epoch
    model.train()
    start_time = time.time()

    for inputs, targets in trainloader:
        inputs, targets = inputs.to(device), targets.to(device)
        optimizer.zero_grad()
        outputs = model(inputs)
        loss = criterion(outputs, targets)
        loss.backward()
        optimizer.step()

    epoch_time = time.time() - start_time
    return epoch_time

    # Test different batch sizes
batch_sizes = [32, 128, 512]
for bs in batch_sizes:
    time_taken = measure_single_gpu_time(bs)
    print(f"Batch size {bs}: {time_taken:.2f} seconds")
```

```
Files already downloaded and verified
Batch size 32: 12.13 seconds
Files already downloaded and verified
Batch size 128: 7.53 seconds
Files already downloaded and verified
Batch size 512: 7.48 seconds
```

2. Measure running time and speedup on multiple GPUs

Measure running time with batch sizes used in part 1 (i.e., 32, 128, etc.) on 2 GPUs and 4 GPUs, and calculate speedup for each setup. Again, for each setup, perform 2 epochs, and only measure the 2nd epoch. When measuring speedup, include all the training components (e.g., data loading, CPU-GPU time, compute time).

(5 points)

Expected Answer: Table 1 records the training time and speedup for different batch sizes up to 4 GPUs. Comment on which type of scaling you are measuring: weak-scaling or strong-scaling? Comment on whether, if the other type of scaling were used, the speedup numbers would be better or worse than what you are measuring.

```
In [18]: def measure_multi_gpu_time(batch_size, num_gpus):
    if torch.cuda.device_count() < num_gpus:
        return None

    # Load dataset
    trainset = torchvision.datasets.CIFAR10(root='./data', train=True,
                                            download=True, transform=transform_
    trainloader = DataLoader(trainset, batch_size=batch_size,
                            shuffle=True, num_workers=2)

    # Initialize model with DataParallel
    device = torch.device("cuda:0")
    model = torchvision.models.resnet18(pretrained=False)
    model = DataParallel(model, device_ids=list(range(num_gpus))).to(device)

    criterion = nn.CrossEntropyLoss()
    optimizer = optim.SGD(model.parameters(), lr=0.1, momentum=0.9, weight_c

    # Warmup epoch
    model.train()
    for inputs, targets in trainloader:
        inputs, targets = inputs.to(device), targets.to(device)
        optimizer.zero_grad()
        outputs = model(inputs)
        loss = criterion(outputs, targets)
        loss.backward()
        optimizer.step()

    # Measure second epoch
    model.train()
    start_time = time.time()

    compute_time = 0
    communication_time = 0
```

```

for inputs, targets in trainloader:
    inputs, targets = inputs.to(device), targets.to(device)

    # Compute phase
    optimizer.zero_grad()
    comp_start = time.time()
    outputs = model(inputs)
    loss = criterion(outputs, targets)
    loss.backward()
    torch.cuda.synchronize()
    compute_time += time.time() - comp_start

    # Communication phase
    comm_start = time.time()
    optimizer.step()
    torch.cuda.synchronize()
    communication_time += time.time() - comm_start

epoch_time = time.time() - start_time
return epoch_time, compute_time, communication_time

# Test with different configurations
batch_sizes = [32, 128, 512]
num_gpus_list = [2, 4]

results = {}
for bs in batch_sizes:
    results[bs] = {}
    # Get single GPU time for speedup calculation
    single_time = measure_single_gpu_time(bs)
    results[bs][1] = single_time

# Test multi-GPU configurations
for num_gpus in num_gpus_list:
    timing = measure_multi_gpu_time(bs, num_gpus)
    if timing:
        epoch_time, compute_time, comm_time = timing
        speedup = single_time / epoch_time
        results[bs][num_gpus] = {
            'time': epoch_time,
            'speedup': speedup,
            'compute_time': compute_time,
            'comm_time': comm_time
        }
        print(f"Batch size {bs}, {num_gpus} GPUs:")
        print(f"Time: {epoch_time:.2f}s, Speedup: {speedup:.2f}x")

```

Files already downloaded and verified

Files already downloaded and verified

Files already downloaded and verified

Batch-size 32 per GPU	Batch-size 128 per GPU	Batch-size 512 per GPU			
Time (sec)	Speedup	Time (sec)	Speedup	Time (sec)	Speedup

	Batch-size 32 per GPU		Batch-size 128 per GPU		Batch-size 512 per GPU	
1-GPU	14.62	1	8.50	1	8.86	1
2-GPU	38.23	0.38	10.59	0.80	7.71	1.15
4-GPU	49.22	0.30	12.37	0.69	7.66	1.16

Table 1: Speedup Measurement for different batch sizes.

3. Report computation and communication time for 2-GPU and 4-GPU setups

For each batch size per GPU (i.e., 32, 128, 512, etc.), report how much time is spent in computation (including CPU-GPU transferring and calculation) and how much time is spent in communication in the 2-GPU and 4-GPU case for one epoch.

Expected Answer: First, describe how you calculate the compute and communication time in each setup. Second, list compute and communication time in Table 2.

(5 points)

```
In [ ]: # TODO
# Print detailed timing breakdown from previous results
print("\nComputation and Communication Time Breakdown:")
for bs in batch_sizes:
    for num_gpus in [2, 4]:
        if num_gpus in results[bs]:
            res = results[bs][num_gpus]
            print(f"\nBatch size {bs}, {num_gpus} GPUs:")
            print(f"Compute time: {res['compute_time']:.2f}s")
            print(f"Communication time: {res['comm_time']:.2f}s")
```

	Batch-size 32 per GPU		Batch-size 128 per GPU		Batch-size 512 per GPU	
	Compute (sec)	Comm (sec)	Compute (sec)	Comm (sec)	Compute (sec)	Comm (sec)
2-GPU	33.85	2.36	9.10	0.60	3.38	0.16
4-GPU	44.81	2.39	10.89	0.60	3.53	0.16

Table 2: Compute and Communication time for different batch sizes.

4. Calculate communication bandwidth utilization

Assume PyTorch DP implements the all-reduce algorithm as discussed in class (reference below). Calculate the communication bandwidth utilization for each multi-GPU/batch-size-per-GPU setup.

Expected Answer: First, list the formula to calculate how long it takes to finish an all-reduce. Second, list the formula to calculate the bandwidth utilization. Third, list the calculated results in Table 3.
(5 points)

```
In [ ]: def calculate_bandwidth(comm_time, num_gpus):
    # Get model size
    model = torchvision.models.resnet18(pretrained=False)
    total_params = sum(p.numel() for p in model.parameters())
    bytes_per_param = 4 # float32
    total_bytes = total_params * bytes_per_param

    # Calculate theoretical bandwidth using the all-reduce algorithm
    # For ring all-reduce:  $2(n-1)/n \times \text{size}$  where  $n$  is number of GPUs
    scaled_size = 2 * (num_gpus - 1) / num_gpus * total_bytes

    # Convert to GB/s
    bandwidth = (scaled_size / comm_time) / (1024**3)
    return bandwidth

# Calculate and print bandwidth utilization
print("\nBandwidth Utilization:")
for bs in batch_sizes:
    for num_gpus in [2, 4]:
        if num_gpus in results[bs]:
            bandwidth = calculate_bandwidth(
                results[bs][num_gpus]['comm_time'],
                num_gpus
            )
            print(f"Batch size {bs}, {num_gpus} GPUs: {bandwidth:.2f} GB/s")
```

	Batch-size-per-GPU 32	Batch-size-per-GPU 128	Batch-size-per-GPU 512
	Bandwidth Utilization (GB/s)	Bandwidth Utilization (GB/s)	Bandwidth Utilization (GB/s)
2-GPU	0.59	9.35	143.54
4-GPU	0.88	14.01	208.67

Table 3: Communication Bandwidth Utilization.

References

- PyTorch Data Parallel, Available at
[https://pytorch.org/docs/stable/_modules/torch/nn/parallel/data_parallel.html]
(https://pytorch.org/docs/stable/_modules/torch/nn/parallel/data_parallel.html)
- [Bringing HPC Techniques to Deep Learning](#)

Problem 4 - Math Problem Solving with Large Language Models: Exploring Prompting Techniques (20 points)

Set up Environment

Before we begin, we need to install the necessary libraries. We'll use:

- `transformers` for our language model
- `torch` as the underlying deep learning framework
- `datasets` to load and manage our dataset

Run the following cell to install these packages:

```
In [ ]: !pip install datasets  
!pip install openai==0.28
```

Load Dataset and Select a Random Problem

In this section, we import the necessary libraries and load the GSM8K (Grade School Math 8K) dataset. The GSM8K dataset contains a variety of math word problems, making it ideal for testing problem-solving capabilities.

We use the Hugging Face `datasets` library to easily access this dataset. After loading the dataset, we randomly select a math problem from the test set. This approach allows us to test our model on a problem it hasn't seen before, simulating a real-world scenario.

The code below performs the following steps:

1. Imports required libraries
2. Loads the GSM8K dataset
3. Selects a random problem from the test set
4. Prints the selected problem and its correct solution

This setup provides us with a randomly chosen math problem that we'll use to evaluate our model's performance.

In [113...]

```
import random
import torch
from datasets import load_dataset
from huggingface_hub import InferenceClient

# Load the GSM8K dataset
dataset = load_dataset("gsm8k", "main")

# Randomly choose a math question from the test set
random_index = random.randint(0, len(dataset['test']) - 1)
problem_to_solve = dataset['test'][random_index]['question']
correct_answer = dataset['test'][random_index]['answer']

# Print results
print("Problem:")
print(problem_to_solve)
print("\nCorrect Solution:")
print(correct_answer)
```

Problem:

Martin's weight is 55 kg. Carl's weight is 16 kg more than Martin's weight. Christian's weight is 8 kg more than Carl's weight. Harry is 5 kg less than Christian's weight. What is the weight of Harry, in kg?

Correct Solution:

Carl's weight is $55 + 16 = 71$ kg.
Christian's weight is $71 + 8 = 79$ kg.
So, Harry's weight is $79 - 5 = 74$ kg.
74

In [114...]

```
# Validate Hugging Face API token
from huggingface_hub import HfApi, InferenceClient
import os

def validate_hf_token(token):
    try:
        api = HfApi(token=token)
        # Try to get user info to validate token
        user_info = api.whoami()

        # Check if we can access the basic user info
        if user_info and isinstance(user_info, dict):
            print("Token is valid!")
            print(f"Connected as: {user_info.get('name', 'Unknown user')}")
            return True
        return False
    except Exception as e:
        print("Token validation failed!")
        print(f"Error: {str(e)}")
        return False

# Your token (replace with your actual token)
```

```
==== Basic Prompt Results ====
```

Problem: Martin's weight is 55 kg. Carl's weight is 16 kg more than Martin's weight. Christian's weight is 8 kg more than Carl's weight. Harry is 5 kg less than Christian's weight. What is the weight of Harry, in kg?

Prompt used: Solve this math problem:

Generated Solution: Christian's weight is $55 + 16 = 64$ kg. Christian's weight is $64 + 8 = 82$ kg. Harry's weight is $82 - 5 = 58$ kg.

Correct Answer: Carl's weight is $55 + 16 = <<55+16=71>>71$ kg.

Christian's weight is $71 + 8 = <<71+8=79>>79$ kg.

So, Harry's weight is $79 - 5 = <<79-5=74>>74$ kg.

74

One-Shot and Two-Shot Prompting with Numerical Answers

In this section, we'll explore one-shot and two-shot prompting techniques, but with a focus on concise, numerical answers rather than detailed explanations. This approach aims to:

1. Provide the model with examples of correct problem-solving without exposing it to step-by-step reasoning.
2. Test whether the model can infer the problem-solving process from seeing only the question and the final numerical answer.
3. Contrast this method with the more detailed Chain-of-Thought prompting we'll explore later.

We'll create functions for both techniques, re-using our existing `generate_solution` function. Then, we'll test these approaches on our randomly selected problem and compare the results to our previous outputs.

This experiment will help us understand how providing examples with just numerical answers affects the model's problem-solving performance and output style, setting the stage for comparison with the more detailed Chain-of-Thought approach.

```
In [118]: def extract_numeric_answer(answer):  
    # Extract just the numeric answer from the full solution  
    return answer.split('####')[-1].strip()
```

```
In [119]: def one_shot_prompting_numeric(problem_to_solve):  
    one_shot_example = dataset['train'][0]  
    numeric_answer = extract_numeric_answer(one_shot_example['answer'])  
  
    prompt = f"""Solve the following math problem and provide only the numer  
  
Problem: {one_shot_example['question']}
```

```
Answer: {numeric_answer}
```

```
Now, solve this problem and provide only numerical answer:"""
```

```
solution = generate_solution_with_hf(prompt, problem_to_solve)
return prompt, solution
```

```
In [120...]: def two_shot_prompting_numeric(problem_to_solve):
    example1 = dataset['train'][0]
    example2 = dataset['train'][1]
    numeric_answer1 = extract_numeric_answer(example1['answer'])
    numeric_answer2 = extract_numeric_answer(example2['answer'])

    prompt = f"""Solve the following math problems and provide only the numerical answers.

Problem 1: {example1['question']}
Answer 1: {numeric_answer1}

Problem 2: {example2['question']}
Answer 2: {numeric_answer2}

Now, solve this problem and provide only numerical answer:"""

    solution = generate_solution_with_hf(prompt, problem_to_solve)
    return prompt, solution
```

```
In [121...]: # Generate solutions using one-shot and two-shot prompting with numeric answers
# Run one-shot prompting
print("\n==== One-Shot Prompting Results ===")
one_shot_prompt, one_shot_solution = one_shot_prompting_numeric(problem_to_solve)
print("Problem:", problem_to_solve)
print("\nPrompt used:", one_shot_prompt)
print("\nGenerated Solution:", one_shot_solution)
print("\nCorrect Answer:", correct_answer)

# Run two-shot prompting
print("\n==== Two-Shot Prompting Results ===")
two_shot_prompt, two_shot_solution = two_shot_prompting_numeric(problem_to_solve)
print("Problem:", problem_to_solve)
print("\nPrompt used:", two_shot_prompt)
print("\nGenerated Solution:", two_shot_solution)
print("\nCorrect Answer:", correct_answer)
```

==== One-Shot Prompting Results ====

Problem: Martin's weight is 55 kg. Carl's weight is 16 kg more than Martin's weight. Christian's weight is 8 kg more than Carl's weight. Harry is 5 kg less than Christian's weight. What is the weight of Harry, in kg?

Prompt used: Solve the following math problem and provide only the numeric answer:

Problem: Natalia sold clips to 48 of her friends in April, and then she sold half as many clips in May. How many clips did Natalia sell altogether in April and May?

Answer: 72

Now, solve this problem and provide only numerical answer:

Generated Solution: 58

Correct Answer: Carl's weight is $55 + 16 = <<55+16=71>>71$ kg.

Christian's weight is $71 + 8 = <<71+8=79>>79$ kg.

So, Harry's weight is $79 - 5 = <<79-5=74>>74$ kg.

74

==== Two-Shot Prompting Results ====

Problem: Martin's weight is 55 kg. Carl's weight is 16 kg more than Martin's weight. Christian's weight is 8 kg more than Carl's weight. Harry is 5 kg less than Christian's weight. What is the weight of Harry, in kg?

Prompt used: Solve the following math problems and provide only the numeric answers:

Problem 1: Natalia sold clips to 48 of her friends in April, and then she sold half as many clips in May. How many clips did Natalia sell altogether in April and May?

Answer 1: 72

Problem 2: Weng earns \$12 an hour for babysitting. Yesterday, she just did 50 minutes of babysitting. How much did she earn?

Answer 2: 10

Now, solve this problem and provide only numerical answer:

Generated Solution: 71.

Correct Answer: Carl's weight is $55 + 16 = <<55+16=71>>71$ kg.

Christian's weight is $71 + 8 = <<71+8=79>>79$ kg.

So, Harry's weight is $79 - 5 = <<79-5=74>>74$ kg.

74

Chain-of-Thought (CoT) Prompting with Two-Shot Examples

In this section, we'll implement Chain-of-Thought (CoT) prompting, an advanced technique that encourages the model to articulate its problem-solving process step by step. CoT prompting is particularly effective for complex reasoning tasks,

as it allows the model to break down problems into manageable steps and show its work.

Key aspects of CoT prompting:

1. It provides examples of detailed, step-by-step reasoning.
2. It encourages the model to explain its thought process, not just give an answer.
3. It can lead to more accurate results, especially for multi-step problems.
4. It offers transparency into the model's problem-solving approach.

We'll create a two-shot CoT prompting function that demonstrates step-by-step reasoning for two example problems before presenting the target problem. This approach aims to guide the model in providing similarly detailed solutions.

By comparing the CoT results with our previous numeric-only prompting techniques, we can evaluate whether this more detailed approach improves the model's problem-solving capabilities and output quality on our math task.

```
In [122...]  
def select_random_examples(dataset, n=2):  
    """  
        Randomly select n examples from the dataset.  
  
    Args:  
        dataset: Dataset to choose from (could be a dict or any iterable).  
        n (int): Number of examples to select.  
  
    Returns:  
        list: n randomly selected examples.  
    """  
    if isinstance(dataset, dict):  
        # If it's a dictionary, select from its values  
        examples = list(dataset.values())  
    elif hasattr(dataset, '__iter__'):  
        # If it's any other iterable, convert to list  
        examples = list(dataset)  
    else:  
        raise TypeError("Dataset must be a dictionary or iterable.")  
  
    return random.sample(examples, min(n, len(examples)))
```

```
In [123...]  
def two_shot_cot_prompts(problem_to_solve):  
    # Randomly select two examples  
    random_examples = select_random_examples(dataset['train'], n=2)  
  
    prompt = "Let's solve these math problems step by step:\n\n"  
  
    for i, example in enumerate(random_examples, 1):  
        prompt += f"Problem {i}: {example['question']}\n"  
        prompt += "Solution:\n"  
        prompt += "1) Let's identify the important information.\n"  
        prompt += "2) We'll determine the necessary calculations.\n"
```

```
prompt += "3) We'll perform each calculation step by step.\n"
prompt += "4) Finally, we'll state our answer clearly.\n"
prompt += f"{example['answer']}\\n\\n"

prompt += f"Problem: \\n {problem_to_solve}\\n"
prompt += "Now, let's solve this problem using the same step-by-step app"
prompt += "Solution:\\n) "

solution = generate_solution_with_hf(prompt, problem_to_solve)
return prompt, solution
```

In [124...]: # Generate solution using Chain-of-Thought prompting

```
# Print the results
print("\n==== Chain-of-Thought Prompting Results ===")
cot_prompt, cot_solution = two_shot_cot_prompts(problem_to_solve)
print("Problem:", problem_to_solve)
print("\nPrompt used:", cot_prompt)
print("\nGenerated Solution:", cot_solution)
print("\nCorrect Answer:", correct_answer)
```

==== Chain-of-Thought Prompting Results ====

Problem: Martin's weight is 55 kg. Carl's weight is 16 kg more than Martin's weight. Christian's weight is 8 kg more than Carl's weight. Harry is 5 kg less than Christian's weight. What is the weight of Harry, in kg?

Prompt used: Let's solve these math problems step by step:

Problem 1: A leaf is being blown down a sidewalk by swirling gusts of wind. For every five feet that a gust blows it forward, the wind swirls and blows it back two feet. How many feet has it traveled down the sidewalk after 11 gusts of wind?

Solution:

- 1) Let's identify the important information.
- 2) We'll determine the necessary calculations.
- 3) We'll perform each calculation step by step.
- 4) Finally, we'll state our answer clearly.

Each gust blows the leaf forward 5 feet, so 11 gusts will blow it forward $5 * 11 = <<5*11=55>>55$ feet.

Each swirl after a gust blows it back 2 feet, so 11 swirls will blow it back $2 * 11 = <<11*2=22>>22$ feet.

After 11 gusts, the leaf has traveled $55 - 22 = <<55-22=33>>33$ feet down the sidewalk.

33

Problem 2: Emma got \$2000 from the bank. She bought \$400 of furniture and gave $\frac{3}{4}$ of the rest to her friend Anna. How much is left with Emma?

Solution:

- 1) Let's identify the important information.
- 2) We'll determine the necessary calculations.
- 3) We'll perform each calculation step by step.
- 4) Finally, we'll state our answer clearly.

Emma had $\$2000 - \$400 = \$<<2000-400=1600>>1600$ left after buying furniture
Emma gave Anna $\frac{3}{4} * \$1600 = \$<<3/4*1600=1200>>1200$

So, Emma is left with $\$1600 - \$1200 = \$<<1600-1200=400>>400$

400

Problem:

Martin's weight is 55 kg. Carl's weight is 16 kg more than Martin's weight. Christian's weight is 8 kg more than Carl's weight. Harry is 5 kg less than Christian's weight. What is the weight of Harry, in kg?

Now, let's solve this problem using the same step-by-step approach:

Solution:

)

Generated Solution: Carl's weight is $16 + 16 = 32$ kg. Christian's weight is $8 + 8 = 16$ kg. Harry's weight is $32 + 5 = 44$ kg. The answer: 44.

Correct Answer: Carl's weight is $55 + 16 = <<55+16=71>>71$ kg.

Christian's weight is $71 + 8 = <<71+8=79>>79$ kg.

So, Harry's weight is $79 - 5 = <<79-5=74>>74$ kg.

74

```
In [125...]: # Compare accuracy
def extract_final_number(solution):
    try:
        import re
```

```

"""Generate solution using Gemini model"""
model = genai.GenerativeModel('gemini-pro')
full_prompt = f"{prompt}\n{problem}"
response = model.generate_content(full_prompt)
return response.text

def basic_prompt(problem_to_solve):
    prompt = """Solve this math problem step by step. Show your work and end with an equals sign. Problem: {problem_to_solve}"""
    solution = generate_solution_with_gemini(prompt, problem_to_solve)
    return prompt, solution

def one_shot_prompting_numeric(problem_to_solve):
    example = dataset['train'][0]
    example_answer = extract_numeric_answer(example['answer'])

    prompt = f"""Let me show you how to solve a math problem:
Example Problem: {example['question']}
Step-by-step solution:
{example['answer']}
Final answer: {example_answer}

Now solve this new problem using the same approach. Show your work and end with an equals sign. Problem: {problem_to_solve}"""

    solution = generate_solution_with_gemini(prompt, problem_to_solve)
    return prompt, solution

def two_shot_cot_prompting(problem_to_solve):
    example1 = dataset['train'][0]
    example2 = dataset['train'][1]

    prompt = f"""Here are two example math problems solved step by step:
Problem 1: {example1['question']}
Solution 1: {example1['answer']}

Problem 2: {example2['question']}
Solution 2: {example2['answer']}"""

    solution = generate_solution_with_gemini(prompt, problem_to_solve)
    return prompt, solution

# Run and display results
print("\n==== Basic Prompt Results ===")
basic_prompt, basic_solution = basic_prompt(problem_to_solve)
print("Problem:", problem_to_solve)
print("Prompt used:", basic_prompt)
print("Generated Solution:", basic_solution)
print("Correct Answer:", correct_answer)

print("\n==== One-Shot Results ===")
one_shot_prompt, one_shot_solution = one_shot_prompting_numeric(problem_to_solve)
print("Generated Solution:", one_shot_solution)

```

```

print("\n==== Chain-of-Thought Results ===")
cot_prompt, cot_solution = two_shot_cot_prompting(problem_to_solve)
print("\nGenerated Solution:", cot_solution)

# Compare accuracy
def extract_final_number(solution):
    try:
        # Look for number after #####
        if "#####" in str(solution):
            answer_part = str(solution).split("#####")[-1].strip()
            import re
            numbers = re.findall(r'\d+', answer_part)
            return int(numbers[0]) if numbers else None
        # Fallback to last number in text
        numbers = re.findall(r'\d+', str(solution))
        return int(numbers[-1]) if numbers else None
    except:
        return None

print("\n==== Accuracy Comparison ===")
correct_num = extract_final_number(correct_answer)
methods = {
    "Basic": extract_final_number(basic_solution),
    "One-Shot": extract_final_number(one_shot_solution),
    "Chain-of-Thought": extract_final_number(cot_solution)
}

print(f"Correct answer: {correct_num}")
for method, result in methods.items():
    is_correct = result == correct_num if result is not None else False
    print(f"{method}: {result} ({'Correct' if is_correct else 'Incorrect'})")

```

Selected Problem:

An apple orchard sells apples in bags of 10. The orchard sold a total of 2000 apples one day. How much did an orchard earn for selling this at \$5 per bag?

Correct Answer:

There were $2000/10 = <<2000/10=200>>200$ bags of apples sold.

Therefore, the orchard earned $200 \times \$5 = \$<<200*5=1000>>1000$.

1000

== Basic Prompt Results ==

Problem: An apple orchard sells apples in bags of 10. The orchard sold a total of 2000 apples one day. How much did an orchard earn for selling this at \$5 per bag?

Prompt used: Solve this math problem step by step. Show your work and end with a line containing '####' followed by just the final numerical answer:

Generated Solution: 1. The number of bags of apples sold = $2000/10 = 200$ bags.

2. The total amount earned = $200 \times \$5 = \1000 .

\$1000

Correct Answer: There were $2000/10 = <<2000/10=200>>200$ bags of apples sold.

Therefore, the orchard earned $200 \times \$5 = \$<<200*5=1000>>1000$.

1000

== One-Shot Results ==

Generated Solution: The orchard sold $2000/10 = <<2000/10=200>>200$ bags of apples.

The orchard earned $200 \times \$5 = <<200 * \$5=1000>>1000$ dollars.

1000

== Chain-of-Thought Results ==

Generated Solution: The orchard sold $2000/10 = <<2000/10=200>>200$ bags of apples.

The orchard earned $5 \times 200 = \$<<5*200=1000>>1000$.

1000

== Accuracy Comparison ==

Correct answer: 1000

Basic: 1000 (Correct)

One-Shot: 1000 (Correct)

Chain-of-Thought: 1000 (Correct)

Question 5: Implement ReAct Agent with Multiple Tools (20 points)

Implement a ReAct (Reasoning and Acting) agent as described by Yao et al. [1], incorporating three main tools: search, compare, and analyze. This agent should

be able to handle complex queries by reasoning about which tool to use and when.

a) (4 points) Implement the search tool using the SerpAPI integration from previous questions. Ensure it can be easily used by the ReAct agent.

- Proper integration with SerpAPI
- Formatting the search results for use by the ReAct agent

b) (5 points) Create a custom comparison tool using LangChain's `Tool` class.

The tool should accept multiple items and a category as input and return a comparison result.

- Implementing the comparison logic
- Creating an appropriate prompt template for the comparison
- Proper error handling for invalid inputs

c) (5 points) Implement an analysis tool that can summarize and extract key information from search results or comparisons. This tool should use the OpenAI model to generate insightful analyses.

- Implementing the analysis logic
- Creating an appropriate prompt template for the analysis
- Ensuring the analysis output is concise and relevant

d) (6 points) Integrate these tools with a ReAct agent using LangChain. Your implementation should:

- Use LangChain's `initialize_agent` function with the `AgentType.ZERO_SHOT_REACT_DESCRIPTION` agent type
- Include all three tools (search, compare, analyze) as available actions for the agent
- Implement proper error handling and fallback strategies
- Ensure smooth transitions between tools in the agent's reasoning process

e) **Bonus (5 points)** Implement a simple Streamlit user interface for your ReAct agent. Your implementation should include:

- A text input field for users to enter their queries
- A button to submit the query and trigger the ReAct agent
- A display area for showing the final results
- A section to display the step-by-step reasoning process of the ReAct agent

```
In [ ]: # Install required packages
!pip install langchain_openai google-search-results langchain
!pip install -U langchain-community
```

```

        return "\n\n".join([
            f"Title: {r.get('title', '')}\n",
            f"Summary: {r.get('snippet', '')}\n",
            f"Source: {r.get('link', '')}"
            for r in formatted_results[:3] # Limit to top 3 results
        ])
    return results
except Exception as e:
    return f"Search error: {str(e)}"

search_tool = Tool(
    name="Search",
    func=enhanced_search,
    description="Search the internet for current information on a topic. Ret"
)

```

b) Create a custom comparison tool

```

In [4]: # TODO
# comparison_prompt = PromptTemplate(input_variables=, template="")
# comparison_chain = LLMChain(llm=, prompt=)
# def compare_items(query: str) -> str:
#     return comparison_chain.run(items=items, category=category)

# Create comparison tool
comparison_prompt = PromptTemplate(
    input_variables=["items", "category"],
    template="""Compare the following items in the category of {category}:

Items to compare:
{items}

Provide a detailed comparison addressing:
1. Key features and specifications
2. Strengths and weaknesses
3. Notable differences
4. Best use cases
5. Price comparison (if applicable)

Format the response with clear headings and bullet points for easy reading.
"""
)
comparison_chain = LLMChain(llm=llm, prompt=comparison_prompt)

def compare_items(query: str) -> str:
    """
    Compare multiple items based on the query
    """
    try:
        # Parse the comparison query
        if " in terms of " not in query.lower():
            return "Please specify what to compare using format: 'Compare [i

```

```

        items_part, category = query.lower().split(" in terms of ", 1)
        items = items_part.replace("compare", "").strip()
        category = category.strip()

    return comparison_chain.run(items=items, category=category)
except Exception as e:
    return f"Comparison error: {str(e)}"

comparison_tool = Tool(
    name="Compare",
    func=compare_items,
    description="Compare multiple items in a specific category. Use format:")
)

# Create analysis tool
analysis_prompt = PromptTemplate(
    input_variables=["results", "query"],
    template="""
Analyze the following information based on this query:

Query: {query}

Information to analyze:
{results}

Provide:
1. Key findings and insights
2. Patterns or trends
3. Critical evaluation
4. Main takeaways
5. Actionable recommendations

Keep the analysis focused and practical.
"""
)
analysis_chain = LLMChain(llm=llm, prompt=analysis_prompt)

```

```

/state/partition1/job-52933127/ipykernel_701116/1986465123.py:26: LangChainD
eprecationWarning: The class `LLMChain` was deprecated in LangChain 0.1.17 a
nd will be removed in 1.0. Use :meth:`~RunnableSequence`, e.g., `prompt | llm
`` instead.
    comparison_chain = LLMChain(llm=llm, prompt=comparison_prompt)

```

c) Implement an analysis tool

```

In [6]: # TODO
# New function to analyze search results and perform comparisons
# def analyze_results(results: str, query: str) -> str:
# If query is empty or doesn't request comparison, provide a sample query
# analysis_prompt = PromptTemplate(input_variables=, template="")
# analysis_chain = LLMChain(llm= , prompt= )
# return analysis_chain.run(results=results, query=query)
def analyze_results(results: str, query: str) -> str:
    """

```

```
Analyze results from searches or comparisons
"""

try:
    if not results or not query:
        return "Both results and query are required for analysis."

    return analysis_chain.run(results=results, query=query)
except Exception as e:
    return f"Analysis error: {str(e)}"

analysis_tool = Tool(
    name="Analyze",
    func=analyze_results,
    description="Analyze search results or comparisons to provide insights"
)
```

d) Integrate tools with a ReAct agent

```
In [ ]: # TODO
# Integrate tools with ReAct agent
# comparison_tool = Tool(name="", func=, description="")

# Modify the AnalyzeSearchResults tool definition
# analysis_tool = Tool(name="", func=, # Directly use the function
# description="Analyze search results and perform comparisons if needed.")

# Assuming search_tool is a list, and you want to include each item in the list
# tools = [item for item in ] + []

# Initialize the ReAct agent with all tools
tools = [search_tool, comparison_tool, analysis_tool]

agent = initialize_agent(
    tools,
    llm,
    agent=AgentType.ZERO_SHOT_REACT_DESCRIPTION,
    verbose=True,
    max_iterations=10
)
```

```
In [7]: def process_query(query: str, max_steps: int = 100) -> str:
    try:
        return agent({"input": query, "max_iterations": max_steps})["output"]
    except RecursionError:
        return "The query was too complex and exceeded the maximum number of iterations."
    except Exception as e:
        return f"An error occurred: {str(e)}"
```

Test Your Implementation

Use the cell below to test your implementation with a sample query.

```
In [10]: def test_agent():
    """
    Test the ReAct agent with a sample query
    """
    test_query = "What are the top 3 electric vehicles in 2024, and how do they compare in terms of range and charging speed?"
    print(f"Testing query: {test_query}\n")
    print("Processing...\n")
    result = process_query(test_query)
    print("Result:")
    print(result)

if __name__ == "__main__":
    test_agent()
```

Testing query: What are the top 3 electric vehicles in 2024, and how do they compare in terms of range and charging speed?

Processing...

```
> Entering new AgentExecutor chain...
I should first search for the top 3 electric vehicles in 2024.
Action: Search
Action Input: Top 3 electric vehicles in 2024
Observation:
Thought:I should now compare the top 3 electric vehicles in terms of range and charging speed.
Action: Compare
Action Input: Compare Tesla Model 3, Ford Mustang Mach-E, and Chevrolet Bolt in terms of range and charging speed
```

```
/state/partition1/job-52933127/ipykernel_701116/1986465123.py:41: LangChainDeprecationWarning: The method `Chain.run` was deprecated in langchain 0.1.0 and will be removed in 1.0. Use :meth:`~invoke` instead.
```

```
    return comparison_chain.run(items=items, category=category)
```

Observation: ## Range and Charging Speed Comparison

Key Features and Specifications

Feature	Tesla Model 3	Ford Mustang Mach-E	Chevrolet Bolt
Range (EPA estimated)	272-353 miles	247-305 miles	259 miles
Battery capacity	50-82 kWh	68-98.8 kWh	65 kWh
Charging speed (DC fast charging)	Up to 250 kW	Up to 150 kW	Up to 55 kW
Charging time (0-80%)	30-45 minutes	45-60 minutes	60-90 minutes

Strengths and Weaknesses

Tesla Model 3

- * **Strengths:**
 - * Longest range
 - * Fastest charging speed
 - * Supercharger network
- * **Weaknesses:**
 - * Higher price
 - * Limited cargo space

Ford Mustang Mach-E

- * **Strengths:**
 - * Spacious interior
 - * Good range
 - * Comfortable ride
- * **Weaknesses:**
 - * Slower charging speed than Tesla
 - * Limited availability of DC fast chargers

Chevrolet Bolt

- * **Strengths:**
 - * Affordable price
 - * Good range for its price
 - * Compact size
- * **Weaknesses:**
 - * Slowest charging speed
 - * Limited cargo space

Notable Differences

- * **Range:** The Tesla Model 3 has the longest range, followed by the Ford Mustang Mach-E and Chevrolet Bolt.
- * **Charging speed:** The Tesla Model 3 has the fastest charging speed, followed by the Ford Mustang Mach-E and Chevrolet Bolt.
- * **Charging network:** Tesla has the most extensive Supercharger network, which gives it an advantage in terms of charging convenience.
- * **Price:** The Chevrolet Bolt is the most affordable option, followed by the Ford Mustang Mach-E and Tesla Model 3.

Best Use Cases

```
* **Tesla Model 3:** Long-distance travel, daily commuting, performance driving  
* **Ford Mustang Mach-E:** Family transportation, road trips, outdoor adventures  
* **Chevrolet Bolt:** City driving, short-distance commuting, budget-conscious buyers
```

Price Comparison

Model	Starting Price
Tesla Model 3	\$46,990
Ford Mustang Mach-E	\$43,895
Chevrolet Bolt	\$25,600

Thought: I have now compared the top 3 electric vehicles in terms of range and charging speed.

Final Answer: The top 3 electric vehicles in 2024 are the Tesla Model 3, Ford Mustang Mach-E, and Chevrolet Bolt. The Tesla Model 3 has the longest range and fastest charging speed, while the Chevrolet Bolt is the most affordable option.

> Finished chain.

Result:

The top 3 electric vehicles in 2024 are the Tesla Model 3, Ford Mustang Mach-E, and Chevrolet Bolt. The Tesla Model 3 has the longest range and fastest charging speed, while the Chevrolet Bolt is the most affordable option.

Submission Requirements

Please submit the following items as part of your solution:

1. Your complete code implementation for the ReAct agent and its tools.
2. A sample question that you used to test your tool (make it complex enough to demonstrate the use of multiple tools).
3. The final answer provided by your ReAct agent for the sample question.
4. The complete history traces of the ReAct agent for your sample question, showing its thought process, actions, and observations. Your traces should follow a format similar to this example:

Thought: I need to find information about top smartphones first

Action: Search[top smartphones 2023]

Observation: [Search results about top smartphones]

Thought: Now I should compare the top two options

Action: Compare[iPhone 14 Pro, Samsung Galaxy S23 Ultra, smartphones]

Observation: [Comparison result]

Thought: I should analyze this comparison for the user

Action: Analyze[comparison result]

Observation: [Analysis of the comparison]

Final Answer: [Your agent's final response to the user's query]

Ensure that your submission clearly demonstrates the agent's ability to reason about which tool to use and how to interpret the results from each tool. Your history traces should show a logical flow of thoughts, actions, and observations, culminating in a final answer that addresses the initial query.

Note: Ensure that your ReAct agent can seamlessly switch between these tools based on the task at hand. The agent should be able to reason about which tool to use next and how to interpret the results from each tool.

References

- [1] Yao, S., Zhao, J., Yu, D., Du, N., Shafran, I., Narasimhan, K., & Cao, Y. (2022). ReAct: Synergizing reasoning and acting in language models. arXiv preprint arXiv:2210.03629. <https://arxiv.org/pdf/2210.03629>

This notebook was converted with [convert.ploomber.io](#)